

O'REILLY®

Third
Edition

Learning SQL

Generate, Manipulate, and Retrieve Data



Alan Beaulieu

Learning SQL

As data floods into your company, you need to put it to work right away—and SQL is the best tool for the job. With the latest edition of this introductory guide, author Alan Beaulieu helps developers get up to speed with SQL fundamentals for writing database applications, performing administrative tasks, and generating reports. You'll find new chapters on analytic functions, strategies for working with large databases, and SQL and big data.

Each chapter presents a self-contained lesson on a key SQL concept or technique using numerous illustrations and annotated examples. Exercises let you practice the skills you learn. Knowledge of SQL is a must for interacting with data. With *Learning SQL*, you'll quickly discover how to put the power and flexibility of this language to work.

- Move quickly through SQL basics and several advanced features
- Use SQL data statements to generate, manipulate, and retrieve data
- Create database objects such as tables, indexes, and constraints with SQL schema statements
- Learn how data sets interact with queries; understand the importance of subqueries
- Convert and manipulate data with SQL's built-in functions and use conditional logic in data statements

"From the basics of SQL to advanced topics such as analytical functions and working with large databases, this third edition of *Learning SQL* provides everything you need to know about SQL in today's modern database world."

—Mark Richards

Author of *Fundamentals of Software Architecture* (O'Reilly)

Alan Beaulieu has been designing, building, and implementing custom database applications for over 25 years. He's the coauthor of *Mastering Oracle SQL* (O'Reilly) and has written an online course on SQL for the University of California. Alan runs his own consulting company that specializes in database design and development for financial services and telecommunications.

DATABASES

US \$59.99 CAN \$79.99

ISBN: 978-1-492-05761-1



Twitter: @oreillymedia
facebook.com/oreilly

THIRD EDITION

Learning SQL

Generate, Manipulate, and Retrieve Data

Alan Beaulieu

Beijing • Boston • Farnham • Sebastopol • Tokyo

O'REILLY®

Learning SQL

by Alan Beaulieu

Copyright © 2020 Alan Beaulieu. All rights reserved.

Printed in the United States of America.

Published by O'Reilly Media, Inc., 1005 Gravenstein Highway North, Sebastopol, CA 95472.

O'Reilly books may be purchased for educational, business, or sales promotional use. Online editions are also available for most titles (<http://oreilly.com>). For more information, contact our corporate/institutional sales department: 800-998-9938 or corporate@oreilly.com.

Acquisitions Editor: Jessica Haberman

Indexer: Angela Howard

Development Editor: Jeff Bleiel

Interior Designer: David Futato

Production Editor: Deborah Baker

Cover Designer: Karen Montgomery

Copyeditor: Charles Roumeliots

Illustrator: Rebecca Demarest

Proofreader: Chris Morris

August 2005: First Edition

April 2009: Second Edition

April 2020: Third Edition

Revision History for the Third Edition

2020-03-04: First Release

2020-09-04: Second Release

See <http://oreilly.com/catalog/errata.csp?isbn=9781492057611> for release details.

The O'Reilly logo is a registered trademark of O'Reilly Media, Inc. *Learning SQL*, the cover image, and related trade dress are trademarks of O'Reilly Media, Inc.

The views expressed in this work are those of the author, and do not represent the publisher's views. While the publisher and the author have used good faith efforts to ensure that the information and instructions contained in this work are accurate, the publisher and the author disclaim all responsibility for errors or omissions, including without limitation responsibility for damages resulting from the use of or reliance on this work. Use of the information and instructions contained in this work is at your own risk. If any code samples or other technology this work contains or describes is subject to open source licenses or the intellectual property rights of others, it is your responsibility to ensure that your use thereof complies with such licenses and/or rights.

978-1-492-05761-1

[MBP]

Table of Contents

Preface.....	xi
1. A Little Background.....	1
Introduction to Databases	1
Nonrelational Database Systems	2
The Relational Model	5
Some Terminology	7
What Is SQL?	8
SQL Statement Classes	9
SQL: A Nonprocedural Language	10
SQL Examples	11
What Is MySQL?	13
SQL Unplugged	14
What's in Store	15
2. Creating and Populating a Database.....	17
Creating a MySQL Database	17
Using the mysql Command-Line Tool	18
MySQL Data Types	20
Character Data	20
Numeric Data	23
Temporal Data	25
Table Creation	27
Step 1: Design	27
Step 2: Refinement	28
Step 3: Building SQL Schema Statements	30
Populating and Modifying Tables	33
Inserting Data	33

Updating Data	38
Deleting Data	39
When Good Statements Go Bad	39
Nonunique Primary Key	39
Nonexistent Foreign Key	40
Column Value Violations	40
Invalid Date Conversions	40
The Sakila Database	41
3. Query Primer.....	45
Query Mechanics	45
Query Clauses	47
The select Clause	48
Column Aliases	50
Removing Duplicates	51
The from Clause	53
Tables	53
Table Links	56
Defining Table Aliases	57
The where Clause	58
The group by and having Clauses	60
The order by Clause	61
Ascending Versus Descending Sort Order	63
Sorting via Numeric Placeholders	64
Test Your Knowledge	65
Exercise 3-1	65
Exercise 3-2	65
Exercise 3-3	65
Exercise 3-4	65
4. Filtering.....	67
Condition Evaluation	67
Using Parentheses	68
Using the not Operator	69
Building a Condition	70
Condition Types	71
Equality Conditions	71
Range Conditions	73
Membership Conditions	77
Matching Conditions	79
Null: That Four-Letter Word	82
Test Your Knowledge	85

Exercise 4-1	86
Exercise 4-2	86
Exercise 4-3	86
Exercise 4-4	86
5. Querying Multiple Tables.....	87
What Is a Join?	87
Cartesian Product	88
Inner Joins	89
The ANSI Join Syntax	91
Joining Three or More Tables	93
Using Subqueries as Tables	95
Using the Same Table Twice	96
Self-Joins	98
Test Your Knowledge	99
Exercise 5-1	99
Exercise 5-2	99
Exercise 5-3	100
6. Working with Sets.....	101
Set Theory Primer	101
Set Theory in Practice	104
Set Operators	105
The union Operator	106
The intersect Operator	108
The except Operator	109
Set Operation Rules	111
Sorting Compound Query Results	111
Set Operation Precedence	112
Test Your Knowledge	114
Exercise 6-1	114
Exercise 6-2	114
Exercise 6-3	114
7. Data Generation, Manipulation, and Conversion.....	115
Working with String Data	115
String Generation	116
String Manipulation	121
Working with Numeric Data	129
Performing Arithmetic Functions	129
Controlling Number Precision	131
Handling Signed Data	133

Working with Temporal Data	134
Dealing with Time Zones	134
Generating Temporal Data	136
Manipulating Temporal Data	140
Conversion Functions	144
Test Your Knowledge	145
Exercise 7-1	145
Exercise 7-2	145
Exercise 7-3	145
8. Grouping and Aggregates.....	147
Grouping Concepts	147
Aggregate Functions	150
Implicit Versus Explicit Groups	151
Counting Distinct Values	152
Using Expressions	153
How Nulls Are Handled	153
Generating Groups	155
Single-Column Grouping	155
Multicolumn Grouping	156
Grouping via Expressions	157
Generating Rollups	157
Group Filter Conditions	159
Test Your Knowledge	160
Exercise 8-1	160
Exercise 8-2	160
Exercise 8-3	160
9. Subqueries.....	161
What Is a Subquery?	161
Subquery Types	163
Noncorrelated Subqueries	163
Multiple-Row, Single-Column Subqueries	164
Multicolumn Subqueries	169
Correlated Subqueries	171
The exists Operator	173
Data Manipulation Using Correlated Subqueries	174
When to Use Subqueries	175
Subqueries as Data Sources	176
Subqueries as Expression Generators	182
Subquery Wrap-Up	184
Test Your Knowledge	185

Exercise 9-1	185
Exercise 9-2	185
Exercise 9-3	185
10. Joins Revisited.....	187
Outer Joins	187
Left Versus Right Outer Joins	190
Three-Way Outer Joins	191
Cross Joins	192
Natural Joins	198
Test Your Knowledge	199
Exercise 10-1	200
Exercise 10-2	200
Exercise 10-3 (Extra Credit)	200
11. Conditional Logic.....	201
What Is Conditional Logic?	201
The case Expression	202
Searched case Expressions	202
Simple case Expressions	204
Examples of case Expressions	205
Result Set Transformations	205
Checking for Existence	206
Division-by-Zero Errors	208
Conditional Updates	209
Handling Null Values	210
Test Your Knowledge	211
Exercise 11-1	211
Exercise 11-2	211
12. Transactions.....	213
Multiuser Databases	213
Locking	214
Lock Granularities	214
What Is a Transaction?	215
Starting a Transaction	217
Ending a Transaction	218
Transaction Savepoints	219
Test Your Knowledge	222
Exercise 12-1	222

13. Indexes and Constraints.....	223
Indexes	223
Index Creation	224
Types of Indexes	229
How Indexes Are Used	231
The Downside of Indexes	232
Constraints	233
Constraint Creation	234
Test Your Knowledge	237
Exercise 13-1	237
Exercise 13-2	237
14. Views.....	239
What Are Views?	239
Why Use Views?	242
Data Security	242
Data Aggregation	243
Hiding Complexity	244
Joining Partitioned Data	244
Updatable Views	245
Updating Simple Views	246
Updating Complex Views	247
Test Your Knowledge	249
Exercise 14-1	249
Exercise 14-2	250
15. Metadata.....	251
Data About Data	251
information_schema	252
Working with Metadata	257
Schema Generation Scripts	257
Deployment Verification	260
Dynamic SQL Generation	261
Test Your Knowledge	265
Exercise 15-1	265
Exercise 15-2	265
16. Analytic Functions.....	267
Analytic Function Concepts	267
Data Windows	268
Localized Sorting	269
Ranking	270

Ranking Functions	271
Generating Multiple Rankings	274
Reporting Functions	277
Window Frames	279
Lag and Lead	281
Column Value Concatenation	283
Test Your Knowledge	284
Exercise 16-1	284
Exercise 16-2	285
Exercise 16-3	285
17. Working with Large Databases.....	287
Partitioning	287
Partitioning Concepts	288
Table Partitioning	288
Index Partitioning	289
Partitioning Methods	289
Partitioning Benefits	297
Clustering	297
Sharding	298
Big Data	299
Hadoop	299
NoSQL and Document Databases	300
Cloud Computing	300
Conclusion	301
18. SQL and Big Data.....	303
Introduction to Apache Drill	303
Querying Files Using Drill	304
Querying MySQL Using Drill	306
Querying MongoDB Using Drill	309
Drill with Multiple Data Sources	315
Future of SQL	317
A. ER Diagram for Example Database.....	319
B. Solutions to Exercises.....	321
Index.....	349

Preface

Programming languages come and go constantly, and very few languages in use today have roots going back more than a decade or so. Some examples are COBOL, which is still used quite heavily in mainframe environments; Java, which was born in the mid-1990s and has become one of the most popular programming languages; and C, which is still quite popular for operating systems and server development and for embedded systems. In the database arena, we have SQL, whose roots go all the way back to the 1970s.

SQL was initially created to be the language for generating, manipulating, and retrieving data from relational databases, which have been around for more than 40 years. Over the past decade or so, however, other data platforms such as Hadoop, Spark, and NoSQL have gained a great deal of traction, eating away at the relational database market. As will be discussed in the last few chapters of this book, however, the SQL language has been evolving to facilitate the retrieval of data from various platforms, regardless of whether the data is stored in tables, documents, or flat files.

Why Learn SQL?

Whether you will be using a relational database or not, if you are working in data science, business intelligence, or some other facet of data analysis, you will likely need to know SQL, along with other languages/platforms such as Python and R. Data is everywhere, in huge quantities, and arriving at a rapid pace, and people who can extract meaningful information from all this data are in big demand.

Why Use This Book to Do It?

There are plenty of books out there that treat you like a dummy, idiot, or some other flavor of simpleton, but these books tend to just skim the surface. At the other end of the spectrum are reference guides that detail every permutation of every statement in a language, which can be useful if you already have a good idea of what you want to

do but just need the syntax. This book strives to find the middle ground, starting with some background of the SQL language, moving through the basics, and then progressing into some of the more advanced features that will allow you to really shine. Additionally, this book ends with a chapter showing how to query data in nonrelational databases, which is a topic rarely covered in introductory books.

Structure of This Book

This book is divided into 18 chapters and 2 appendixes:

Chapter 1, A Little Background

Explores the history of computerized databases, including the rise of the relational model and the SQL language.

Chapter 2, Creating and Populating a Database

Demonstrates how to create a MySQL database, create the tables used for the examples in this book, and populate the tables with data.

Chapter 3, Query Primer

Introduces the `select` statement and further demonstrates the most common clauses (`select`, `from`, `where`).

Chapter 4, Filtering

Demonstrates the different types of conditions that can be used in the `where` clause of a `select`, `update`, or `delete` statement.

Chapter 5, Querying Multiple Tables

Shows how queries can utilize multiple tables via table joins.

Chapter 6, Working with Sets

This chapter is all about data sets and how they can interact within queries.

Chapter 7, Data Generation, Manipulation, and Conversion

Demonstrates several built-in functions used for manipulating or converting data.

Chapter 8, Grouping and Aggregates

Shows how data can be aggregated.

Chapter 9, Subqueries

Introduces subqueries (a personal favorite) and shows how and where they can be utilized.

Chapter 10, Joins Revisited

Further explores the various types of table joins.

Chapter 11, Conditional Logic

Explores how conditional logic (i.e., if-then-else) can be utilized in `select`, `insert`, `update`, and `delete` statements.

Chapter 12, Transactions

Introduces transactions and shows how to use them.

Chapter 13, Indexes and Constraints

Explores indexes and constraints.

Chapter 14, Views

Shows how to build an interface to shield users from data complexities.

Chapter 15, Metadata

Demonstrates the utility of the data dictionary.

Chapter 16, Analytic Functions

Covers functionality used to generate rankings, subtotals, and other values used heavily in reporting and analysis.

Chapter 17, Working with Large Databases

Demonstrates techniques for making very large databases easier to manage and traverse.

Chapter 18, SQL and Big Data

Explores the transformation of the SQL language to allow retrieval of data from nonrelational data platforms.

Appendix A, ER Diagram for Example Database

Shows the database schema used for all examples in the book.

Appendix B, Solutions to Exercises

Shows solutions to the chapter exercises.

Conventions Used in This Book

The following typographical conventions are used in this book:

Italic

Indicates new terms, URLs, email addresses, filenames, and file extensions.

Constant width

Used for program listings, as well as within paragraphs to refer to program elements such as variable or function names, databases, data types, environment variables, statements, and keywords.

Constant width italic

Shows text that should be replaced with user-supplied values or by values determined by context.

Constant width bold

Shows commands or other text that should be typed literally by the user.



Indicates a tip, suggestion, or general note. For example, I use notes to point you to useful new features in Oracle9*i*.



Indicates a warning or caution. For example, I'll tell you if a certain SQL clause might have unintended consequences if not used carefully.

Using the Examples in This Book

To experiment with the data used for the examples in this book, you have two options:

- Download and install the MySQL server version 8.0 (or later) and load the Sakila example database from <https://dev.mysql.com/doc/index-other.html>.
- Go to <https://www.katacoda.com/mysql-db-sandbox/scenarios/mysql-sandbox> to access the MySQL Sandbox, which has the Sakila sample database loaded in a MySQL instance. You'll have to set up a (free) Katacoda account. Then, click the Start Scenario button.

If you choose the second option, once you start the scenario, a MySQL server is installed and started, and then the Sakila schema and data are loaded. When it's ready, a standard `mysql>` prompt appears, and you can then start querying the sample database. This is certainly the easiest option, and I anticipate that most readers will choose this option; if this sounds good to you, feel free to skip ahead to the next section.

If you prefer to have your own copy of the data and want any changes you have made to be permanent, or if you are just interested in installing the MySQL server on your own machine, you may prefer the first option. You may also opt to use a MySQL server hosted in an environment such as Amazon Web Services or Google Cloud. In either case, you will need to perform the installation/configuration yourself, as it is beyond the scope of this book. Once your database is available, you will need to follow a few steps to load the Sakila sample database.

First, you will need to launch the `mysql` command-line client and provide a password, and then perform the following steps:

1. Go to <https://dev.mysql.com/doc/index-other.html> and download the files for “sakila database” under the Example Databases section.
2. Put the files in a local directory such as `C:\temp\sakila-db` (used for the next two steps, but overwrite with your directory path).
3. Type `source c:\temp\sakila-db\sakila-schema.sql`; and press Enter.
4. Type `source c:\temp\sakila-db\sakila-data.sql`; and press Enter.

You should now have a working database populated with all the data needed for the examples in this book.

O'Reilly Online Learning



For more than 40 years, *O'Reilly Media* has provided technology and business training, knowledge, and insight to help companies succeed.

Our unique network of experts and innovators share their knowledge and expertise through books, articles, and our online learning platform. O'Reilly's online learning platform gives you on-demand access to live training courses, in-depth learning paths, interactive coding environments, and a vast collection of text and video from O'Reilly and 200+ other publishers. For more information, please visit <http://oreilly.com>.

How to Contact Us

Please address comments and questions concerning this book to the publisher:

O'Reilly Media, Inc.
1005 Gravenstein Highway North
Sebastopol, CA 95472
800-998-9938 (in the United States or Canada)
707-829-0515 (international or local)
707-829-0104 (fax)

We have a web page for this book, where we list errata and any additional information. You can access this page at https://oreil.ly/Learning_SQL3.

Email bookquestions@oreilly.com to comment or ask technical questions about this book.

For news and more information about our books and courses, see our website at <http://www.oreilly.com>.

Find us on Facebook: <http://facebook.com/oreilly>

Follow us on Twitter: <http://twitter.com/oreillymedia>

Watch us on YouTube: <http://www.youtube.com/oreillymedia>

Acknowledgments

I would like to thank my editor, Jeff Bleiel, for helping to make this third edition a reality, along with Thomas Nield, Ann White-Watkins, and Charles Givre, who were kind enough to review the book for me. Thanks also go to Deb Baker, Jess Haberman, and all the other folks at O'Reilly Media who were involved. Lastly, I thank my wife, Nancy, and my daughters, Michelle and Nicole, for their encouragement and inspiration.

A Little Background

Before we roll up our sleeves and get to work, it would be helpful to survey the history of database technology in order to better understand how relational databases and the SQL language evolved. Therefore, I'd like to start by introducing some basic database concepts and looking at the history of computerized data storage and retrieval.



For those readers anxious to start writing queries, feel free to skip ahead to [Chapter 3](#), but I recommend returning later to the first two chapters in order to better understand the history and utility of the SQL language.

Introduction to Databases

A *database* is nothing more than a set of related information. A telephone book, for example, is a database of the names, phone numbers, and addresses of all people living in a particular region. While a telephone book is certainly a ubiquitous and frequently used database, it suffers from the following:

- Finding a person's telephone number can be time consuming, especially if the telephone book contains a large number of entries.
- A telephone book is indexed only by last/first names, so finding the names of the people living at a particular address, while possible in theory, is not a practical use for this database.
- From the moment the telephone book is printed, the information becomes less and less accurate as people move into or out of a region, change their telephone numbers, or move to another location within the same region.

The same drawbacks attributed to telephone books can also apply to any manual data storage system, such as patient records stored in a filing cabinet. Because of the cumbersome nature of paper databases, some of the first computer applications developed were *database systems*, which are computerized data storage and retrieval mechanisms. Because a database system stores data electronically rather than on paper, a database system is able to retrieve data more quickly, index data in multiple ways, and deliver up-to-the-minute information to its user community.

Early database systems managed data stored on magnetic tapes. Because there were generally far more tapes than tape readers, technicians were tasked with loading and unloading tapes as specific data was requested. Because the computers of that era had very little memory, multiple requests for the same data generally required the data to be read from the tape multiple times. While these database systems were a significant improvement over paper databases, they are a far cry from what is possible with today's technology. (Modern database systems can manage petabytes of data, accessed by clusters of servers each caching tens of gigabytes of that data in high-speed memory, but I'm getting a bit ahead of myself.)

Nonrelational Database Systems



This section contains some background information about pre-relational database systems. For those readers eager to dive into SQL, feel free to skip ahead a couple of pages to the next section.

Over the first several decades of computerized database systems, data was stored and represented to users in various ways. In a *hierarchical database system*, for example, data is represented as one or more tree structures. [Figure 1-1](#) shows how data relating to George Blake's and Sue Smith's bank accounts might be represented via tree structures.

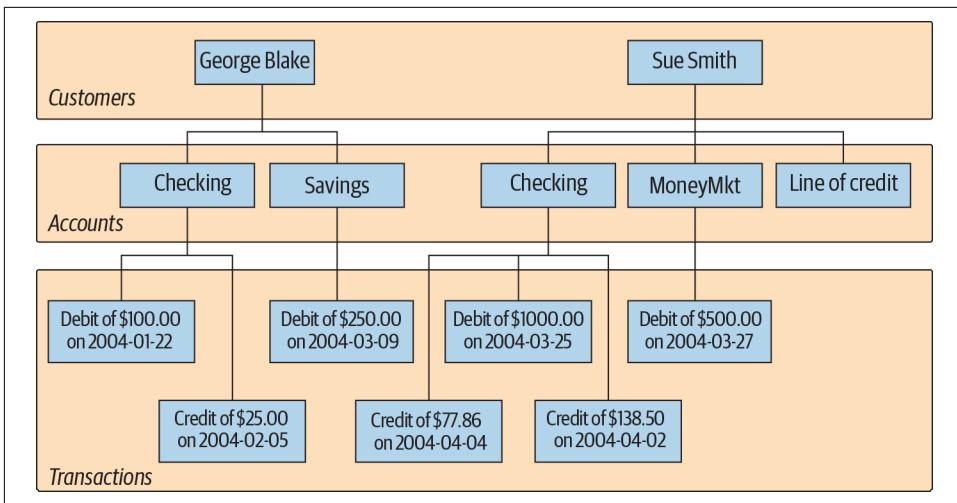


Figure 1-1. Hierarchical view of account data

George and Sue each have their own tree containing their accounts and the transactions on those accounts. The hierarchical database system provides tools for locating a particular customer's tree and then traversing the tree to find the desired accounts and/or transactions. Each node in the tree may have either zero or one parent and zero, one, or many children. This configuration is known as a *single-parent hierarchy*.

Another common approach, called the *network database system*, exposes sets of records and sets of links that define relationships between different records. **Figure 1-2** shows how George's and Sue's same accounts might look in such a system.

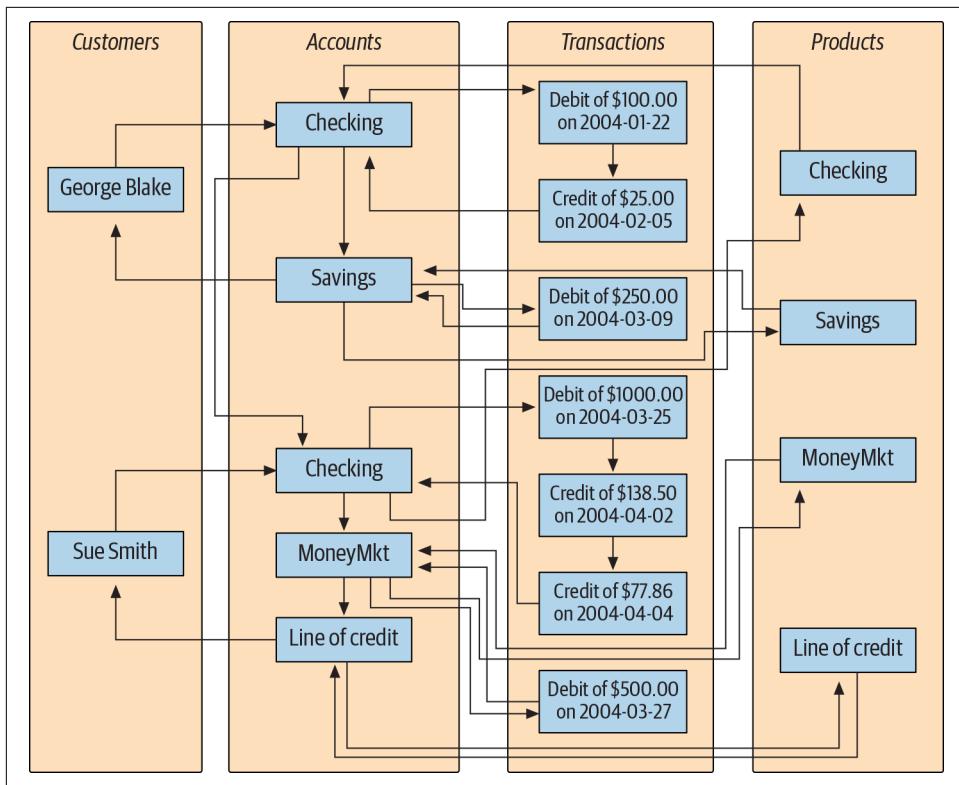


Figure 1-2. Network view of account data

In order to find the transactions posted to Sue's money market account, you would need to perform the following steps:

1. Find the customer record for Sue Smith.
2. Follow the link from Sue Smith's customer record to her list of accounts.
3. Traverse the chain of accounts until you find the money market account.
4. Follow the link from the money market record to its list of transactions.

One interesting feature of network database systems is demonstrated by the set of product records on the far right of [Figure 1-2](#). Notice that each **product** record (Checking, Savings, etc.) points to a list of account records that are of that product type. Account records, therefore, can be accessed from multiple places (both customer records and product records), allowing a network database to act as a *multi-parent hierarchy*.

Both hierarchical and network database systems are alive and well today, although generally in the mainframe world. Additionally, hierarchical database systems have

enjoyed a rebirth in the directory services realm, such as Microsoft's Active Directory and the open source Apache Directory Server. Beginning in the 1970s, however, a new way of representing data began to take root, one that was more rigorous yet easy to understand and implement.

The Relational Model

In 1970, Dr. E. F. Codd of IBM's research laboratory published a paper titled "A Relational Model of Data for Large Shared Data Banks" that proposed that data be represented as sets of *tables*. Rather than using pointers to navigate between related entities, redundant data is used to link records in different tables. [Figure 1-3](#) shows how George's and Sue's account information would appear in this context.

Customer			Account			
cust_id	fname	lname	account_id	product_cd	cust_id	balance
1	George	Blake	103	CHK	1	\$75.00
2	Sue	Smith	104	SAV	1	\$250.00
			105	CHK	2	\$783.64
			106	MM	2	\$500.00
			107	LOC	2	0

Product		Transaction			
product_cd	name	txn_id	txn_type_cd	account_id	amount
CHK	Checking	978	DBT	103	\$100.00
SAV	Savings	979	CDT	103	\$25.00
MM	Money market	980	DBT	104	\$250.00
LOC	Line of credit	981	DBT	105	\$1000.00
		982	CDT	105	\$138.50
		983	CDT	105	\$77.86
		984	DBT	106	\$500.00

Figure 1-3. Relational view of account data

The four tables in [Figure 1-3](#) represent the four entities discussed so far: `customer`, `product`, `account`, and `transaction`. Looking across the top of the `customer` table in [Figure 1-3](#), you can see three *columns*: `cust_id` (which contains the customer's ID number), `fname` (which contains the customer's first name), and `lname` (which contains the customer's last name). Looking down the side of the `customer` table, you can see two *rows*, one containing George Blake's data and the other containing Sue Smith's data. The number of columns that a table may contain differs from server to server, but it is generally large enough not to be an issue (Microsoft SQL Server, for example, allows up to 1,024 columns per table). The number of rows that a table may contain is more a matter of physical limits (i.e., how much disk drive space is available) and maintainability (i.e., how large a table can get before it becomes difficult to work with) than of database server limitations.

Each table in a relational database includes information that uniquely identifies a row in that table (known as the *primary key*), along with additional information needed to describe the entity completely. Looking again at the `customer` table, the `cust_id` column holds a different number for each customer; George Blake, for example, can be uniquely identified by customer ID 1. No other customer will ever be assigned that identifier, and no other information is needed to locate George Blake's data in the `customer` table.



Every database server provides a mechanism for generating unique sets of numbers to use as primary key values, so you won't need to worry about keeping track of what numbers have been assigned.

While I might have chosen to use the combination of the `fname` and `lname` columns as the primary key (a primary key consisting of two or more columns is known as a *compound key*), there could easily be two or more people with the same first and last names who have accounts at the bank. Therefore, I chose to include the `cust_id` column in the `customer` table specifically for use as a primary key column.



In this example, choosing `fname/lname` as the primary key would be referred to as a *natural key*, whereas the choice of `cust_id` would be referred to as a *surrogate key*. The decision whether to employ natural or surrogate keys is up to the database designer, but in this particular case the choice is clear, since a person's last name may change (such as when a person adopts a spouse's last name), and primary key columns should never be allowed to change once a value has been assigned.

Some of the tables also include information used to navigate to another table; this is where the “redundant data” mentioned earlier comes in. For example, the `account` table includes a column called `cust_id`, which contains the unique identifier of the customer who opened the account, along with a column called `product_cd`, which contains the unique identifier of the product to which the account will conform. These columns are known as *foreign keys*, and they serve the same purpose as the lines that connect the entities in the hierarchical and network versions of the account information. If you are looking at a particular account record and want to know more information about the customer who opened the account, you would take the value of the `cust_id` column and use it to find the appropriate row in the `customer` table (this process is known, in relational database lingo, as a *join*; joins are introduced in [Chapter 3](#) and probed deeply in Chapters [5](#) and [10](#)).

It might seem wasteful to store the same data many times, but the relational model is quite clear on what redundant data may be stored. For example, it is proper for the `account` table to include a column for the unique identifier of the customer who opened the account, but it is not proper to include the customer’s first and last names in the `account` table as well. If a customer were to change her name, for example, you want to make sure that there is only one place in the database that holds the customer’s name; otherwise, the data might be changed in one place but not another, causing the data in the database to be unreliable. The proper place for this data is the `customer` table, and only the `cust_id` values should be included in other tables. It is also not proper for a single column to contain multiple pieces of information, such as a `name` column that contains both a person’s first and last names, or an `address` column that contains street, city, state, and zip code information. The process of refining a database design to ensure that each independent piece of information is in only one place (except for foreign keys) is known as *normalization*.

Getting back to the four tables in [Figure 1-3](#), you may wonder how you would use these tables to find George Blake’s transactions against his checking account. First, you would find George Blake’s unique identifier in the `customer` table. Then, you would find the row in the `account` table whose `cust_id` column contains George’s unique identifier and whose `product_cd` column matches the row in the `product` table whose `name` column equals “Checking.” Finally, you would locate the rows in the `transaction` table whose `account_id` column matches the unique identifier from the `account` table. This might sound complicated, but you can do it in a single command, using the SQL language, as you will see shortly.

Some Terminology

I introduced some new terminology in the previous sections, so maybe it’s time for some formal definitions. [Table 1-1](#) shows the terms we use for the remainder of the book along with their definitions.

Table 1-1. Terms and definitions

Term	Definition
Entity	Something of interest to the database user community. Examples include customers, parts, geographic locations, etc.
Column	An individual piece of data stored in a table.
Row	A set of columns that together completely describe an entity or some action on an entity. Also called a record.
Table	A set of rows, held either in memory (nonpersistent) or on permanent storage (persistent).
Result set	Another name for a nonpersistent table, generally the result of an SQL query.
Primary key	One or more columns that can be used as a unique identifier for each row in a table.
Foreign key	One or more columns that can be used together to identify a single row in another table.

What Is SQL?

Along with Codd’s definition of the relational model, he proposed a language called DSL/Alpha for manipulating the data in relational tables. Shortly after Codd’s paper was released, IBM commissioned a group to build a prototype based on Codd’s ideas. This group created a simplified version of DSL/Alpha that they called SQUARE. Refinements to SQUARE led to a language called SEQUEL, which was, finally, shortened to SQL. While SQL began as a language used to manipulate data in relational databases, it has evolved (as you will see toward the end of this book) to be a language for manipulating data across various database technologies.

SQL is now more than 40 years old, and it has undergone a great deal of change along the way. In the mid-1980s, the American National Standards Institute (ANSI) began working on the first standard for the SQL language, which was published in 1986. Subsequent refinements led to new releases of the SQL standard in 1989, 1992, 1999, 2003, 2006, 2008, 2011, and 2016. Along with refinements to the core language, new features have been added to the SQL language to incorporate object-oriented functionality, among other things. The later standards focus on the integration of related technologies, such as extensible markup language (XML) and JavaScript object notation (JSON).

SQL goes hand in hand with the relational model because the result of an SQL query is a table (also called, in this context, a *result set*). Thus, a new permanent table can be created in a relational database simply by storing the result set of a query. Similarly, a query can use both permanent tables and the result sets from other queries as inputs (we explore this in detail in [Chapter 9](#)).

One final note: SQL is not an acronym for anything (although many people will insist it stands for “Structured Query Language”). When referring to the language, it is equally acceptable to say the letters individually (i.e., S. Q. L.) or to use the word *sequel*.

SQL Statement Classes

The SQL language is divided into several distinct parts: the parts that we explore in this book include *SQL schema statements*, which are used to define the data structures stored in the database; *SQL data statements*, which are used to manipulate the data structures previously defined using SQL schema statements; and *SQL transaction statements*, which are used to begin, end, and roll back transactions (concepts covered in [Chapter 12](#)). For example, to create a new table in your database, you would use the SQL schema statement `create table`, whereas the process of populating your new table with data would require the SQL data statement `insert`.

To give you a taste of what these statements look like, here's an SQL schema statement that creates a table called `corporation`:

```
CREATE TABLE corporation
  (corp_id SMALLINT,
   name VARCHAR(30),
   CONSTRAINT pk_corporation PRIMARY KEY (corp_id)
);
```

This statement creates a table with two columns, `corp_id` and `name`, with the `corp_id` column identified as the primary key for the table. We probe the finer details of this statement, such as the different data types available with MySQL, in [Chapter 2](#). Next, here's an SQL data statement that inserts a row into the `corporation` table for Acme Paper Corporation:

```
INSERT INTO corporation (corp_id, name)
VALUES (27, 'Acme Paper Corporation');
```

This statement adds a row to the `corporation` table with a value of 27 for the `corp_id` column and a value of `Acme Paper Corporation` for the `name` column.

Finally, here's a simple `select` statement to retrieve the data that was just created:

```
mysql> SELECT name
    -> FROM corporation
    -> WHERE corp_id = 27;
+-----+
| name           |
+-----+
| Acme Paper Corporation |
+-----+
```

All database elements created via SQL schema statements are stored in a special set of tables called the *data dictionary*. This “data about the database” is known collectively as *metadata* and is explored in [Chapter 15](#). Just like tables that you create yourself, data dictionary tables can be queried via a `select` statement, thereby allowing you to discover the current data structures deployed in the database at runtime. For example, if you are asked to write a report showing the new accounts created last month,

you could either hardcode the names of the columns in the account table that were known to you when you wrote the report, or query the data dictionary to determine the current set of columns and dynamically generate the report each time it is executed.

Most of this book is concerned with the data portion of the SQL language, which consists of the `select`, `update`, `insert`, and `delete` commands. SQL schema statements are demonstrated in [Chapter 2](#), which will lead you through the design and creation of some simple tables. In general, SQL schema statements do not require much discussion apart from their syntax, whereas SQL data statements, while few in number, offer numerous opportunities for detailed study. Therefore, while I try to introduce you to many of the SQL schema statements, most chapters in this book concentrate on the SQL data statements.

SQL: A Nonprocedural Language

If you have worked with programming languages in the past, you are used to defining variables and data structures, using conditional logic (i.e., if-then-else) and looping constructs (i.e., do while ... end), and breaking your code into small, reusable pieces (i.e., objects, functions, procedures). Your code is handed to a compiler, and the executable that results does exactly (well, not always *exactly*) what you programmed it to do. Whether you work with Java, Python, Scala, or some other *procedural* language, you are in complete control of what the program does.



A procedural language defines both the desired results and the mechanism, or process, by which the results are generated. Non-procedural languages also define the desired results, but the process by which the results are generated is left to an external agent.

With SQL, however, you will need to give up some of the control you are used to, because SQL statements define the necessary inputs and outputs, but the manner in which a statement is executed is left to a component of your database engine known as the *optimizer*. The optimizer's job is to look at your SQL statements and, taking into account how your tables are configured and what indexes are available, decide the most efficient execution path (well, not always the *most* efficient). Most database engines will allow you to influence the optimizer's decisions by specifying *optimizer hints*, such as suggesting that a particular index be used; most SQL users, however, will never get to this level of sophistication and will leave such tweaking to their database administrator or performance expert.

Therefore, with SQL, you will not be able to write complete applications. Unless you are writing a simple script to manipulate certain data, you will need to integrate SQL with your favorite programming language. Some database vendors have done this for

you, such as Oracle's PL/SQL language, MySQL's stored procedure language, and Microsoft's Transact-SQL language. With these languages, the SQL data statements are part of the language's grammar, allowing you to seamlessly integrate database queries with procedural commands. If you are using a non-database-specific language such as Java or Python, however, you will need to use a toolkit/API to execute SQL statements from your code. Some of these toolkits are provided by your database vendor, whereas others have been created by third-party vendors or by open source providers. [Table 1-2](#) shows some of the available options for integrating SQL into a specific language.

Table 1-2. SQL integration toolkits

Language	Toolkit
Java	JDBC (Java Database Connectivity)
C#	ADO.NET (Microsoft)
Ruby	Ruby DBI
Python	Python DB
Go	Package database/sql

If you only need to execute SQL commands interactively, every database vendor provides at least a simple command-line tool for submitting SQL commands to the database engine and inspecting the results. Most vendors provide a graphical tool as well that includes one window showing your SQL commands and another window showing the results from your SQL commands. Additionally, there are third-party tools such as SQuirrel, which will connect via a JDBC connection to many different database servers. Since the examples in this book are executed against a MySQL database, I use the `mysql` command-line tool that is included as part of the MySQL installation to run the examples and format the results.

SQL Examples

Earlier in this chapter, I promised to show you an SQL statement that would return all the transactions against George Blake's checking account. Without further ado, here it is:

```
SELECT t.txn_id, t.txn_type_cd, t.txn_date, t.amount
FROM individual i
    INNER JOIN account a ON i.cust_id = a.cust_id
    INNER JOIN product p ON p.product_cd = a.product_cd
    INNER JOIN transaction t ON t.account_id = a.account_id
WHERE i.fname = 'George' AND i.lname = 'Blake'
    AND p.name = 'checking account';

+-----+-----+-----+-----+
| txn_id | txn_type_cd | txn_date           | amount |
+-----+-----+-----+-----+
```

```
+-----+-----+-----+-----+
| 11 | DBT      | 2008-01-05 00:00:00 | 100.00 |
+-----+-----+-----+
1 row in set (0.00 sec)
```

Without going into too much detail at this point, this query identifies the row in the `individual` table for George Blake and the row in the `product` table for the “checking” product, finds the row in the `account` table for this individual/product combination, and returns four columns from the `transaction` table for all transactions posted to this account. If you happen to know that George Blake’s customer ID is 8 and that checking accounts are designated by the code ‘CHK’, then you can simply find George Blake’s checking account in the `account` table based on the customer ID and use the account ID to find the appropriate transactions:

```
SELECT t.txn_id, t.txn_type_cd, t.txn_date, t.amount
FROM account a
INNER JOIN transaction t ON t.account_id = a.account_id
WHERE a.cust_id = 8 AND a.product_cd = 'CHK';
```

I cover all of the concepts in these queries (plus a lot more) in the following chapters, but I wanted to at least show what they would look like.

The previous queries contain three different *clauses*: `select`, `from`, and `where`. Almost every query that you encounter will include at least these three clauses, although there are several more that can be used for more specialized purposes. The role of each of these three clauses is demonstrated by the following:

```
SELECT /* one or more things */ ...
FROM /* one or more places */ ...
WHERE /* one or more conditions apply */ ...
```



Most SQL implementations treat any text between the `/*` and `*/` tags as comments.

When constructing your query, your first task is generally to determine which table or tables will be needed and then add them to your `from` clause. Next, you will need to add conditions to your `where` clause to filter out the data from these tables that you aren’t interested in. Finally, you will decide which columns from the different tables need to be retrieved and add them to your `select` clause. Here’s a simple example that shows how you would find all customers with the last name “Smith”:

```
SELECT cust_id, fname
FROM individual
WHERE lname = 'Smith';
```

This query searches the `individual` table for all rows whose `lname` column matches the string '`Smith`' and returns the `cust_id` and `fname` columns from those rows.

Along with querying your database, you will most likely be involved with populating and modifying the data in your database. Here's a simple example of how you would insert a new row into the `product` table:

```
INSERT INTO product (product_cd, name)
VALUES ('CD', 'Certificate of Depysit')
```

Whoops, looks like you misspelled "Deposit." No problem. You can clean that up with an `update` statement:

```
UPDATE product
SET name = 'Certificate of Deposit'
WHERE product_cd = 'CD';
```

Notice that the `update` statement also contains a `where` clause, just like the `select` statement. This is because an `update` statement must identify the rows to be modified; in this case, you are specifying that only those rows whose `product_cd` column matches the string '`CD`' should be modified. Since the `product_cd` column is the primary key for the `product` table, you should expect your `update` statement to modify exactly one row (or zero, if the value doesn't exist in the table). Whenever you execute an SQL data statement, you will receive feedback from the database engine as to how many rows were affected by your statement. If you are using an interactive tool such as the `mysql` command-line tool mentioned earlier, then you will receive feedback concerning how many rows were either:

- Returned by your `select` statement
- Created by your `insert` statement
- Modified by your `update` statement
- Removed by your `delete` statement

If you are using a procedural language with one of the toolkits mentioned earlier, the toolkit will include a call to ask for this information after your SQL data statement has executed. In general, it's a good idea to check this info to make sure your statement didn't do something unexpected (like when you forgot to put a `where` clause on your `delete` statement and delete every row in the table!).

What Is MySQL?

Relational databases have been available commercially for more than three decades. Some of the most mature and popular commercial products include:

- Oracle Database from Oracle Corporation
- SQL Server from Microsoft
- DB2 Universal Database from IBM

All these database servers do approximately the same thing, although some are better equipped to run very large or very high throughput databases. Others are better at handling objects or very large files or XML documents, and so on. Additionally, all these servers do a pretty good job of complying with the latest ANSI SQL standard. This is a good thing, and I make it a point to show you how to write SQL statements that will run on any of these platforms with little or no modification.

Along with the commercial database servers, there has been quite a bit of activity in the open source community in the past two decades with the goal of creating a viable alternative. Two of the most commonly used open source database servers are PostgreSQL and MySQL. The MySQL server is available for free, and I have found it to be extremely simple to download and install. For these reasons, I have decided that all examples for this book be run against a MySQL (version 8.0) database, and that the `mysql` command-line tool be used to format query results. Even if you are already using another server and never plan to use MySQL, I urge you to install the latest MySQL server, load the sample schema and data, and experiment with the data and examples in this book.

However, keep in mind the following caveat:

This is not a book about MySQL's SQL implementation.

Rather, this book is designed to teach you how to craft SQL statements that will run on MySQL with no modifications, and will run on recent releases of Oracle Database, DB2, and SQL Server with few or no modifications.

SQL Unplugged

A great deal has happened in the database world during the decade between the second and third editions of this book. While relational databases are still heavily used and will continue to be for some time, new database technologies have emerged to meet the needs of companies like Amazon and Google. These technologies include Hadoop, Spark, NoSQL, and NewSQL, which are distributed, scalable systems typically deployed on clusters of commodity servers. While it is beyond the scope of this book to explore these technologies in detail, they do all share something in common with relational databases: SQL.

Since organizations frequently store data using multiple technologies, there is a need to unplug SQL from a particular database server and provide a service that can span multiple databases. For example, a report may need to bring together data stored in

Oracle, Hadoop, JSON files, CSV files, and Unix log files. A new generation of tools have been built to meet this type of challenge, and one of the most promising is Apache Drill, which is an open source query engine that allows users to write queries that can access data stored in most any database or filesystem. We will explore Apache Drill in [Chapter 18](#).

What's in Store

The overall goal of the next four chapters is to introduce the SQL data statements, with a special emphasis on the three main clauses of the `select` statement. Additionally, you will see many examples that use the Sakila schema (introduced in the next chapter), which will be used for all examples in the book. It is my hope that familiarity with a single database will allow you to get to the crux of an example without having to stop and examine the tables being used each time. If it becomes a bit tedious working with the same set of tables, feel free to augment the sample database with additional tables or to invent your own database with which to experiment.

After you have a solid grasp on the basics, the remaining chapters will drill deep into additional concepts, most of which are independent of each other. Thus, if you find yourself getting confused, you can always move ahead and come back later to revisit a chapter. When you have finished the book and worked through all of the examples, you will be well on your way to becoming a seasoned SQL practitioner.

For readers interested in learning more about relational databases, the history of computerized database systems, or the SQL language than was covered in this short introduction, here are a few resources worth checking out:

- *Database in Depth: Relational Theory for Practitioners* by C. J. Date (O'Reilly)
- *An Introduction to Database Systems*, Eighth Edition, by C. J. Date (Addison-Wesley)
- *The Database Relational Model: A Retrospective Review and Analysis*, by C. J. Date (Addison-Wesley)
- [Wikipedia subarticle on definition of “Database Management System”](#)

Creating and Populating a Database

This chapter provides you with the information you need to create your first database and to create the tables and associated data used for the examples in this book. You will also learn about various data types and see how to create tables using them. Because the examples in this book are executed against a MySQL database, this chapter is somewhat skewed toward MySQL's features and syntax, but most concepts are applicable to any server.

Creating a MySQL Database

If you want the ability to experiment with the data used for the examples in this book, you have two options:

- Download and install the MySQL server version 8.0 (or later) and load the Sakila example database from <https://dev.mysql.com/doc/index-other.html>.
- Go to <https://www.katacoda.com/mysql-db-sandbox/scenarios/mysql-sandbox> to access the MySQL Sandbox, which has the Sakila sample database loaded in a MySQL instance. You'll have to set up a (free) Katacoda account. Then, click the Start Scenario button.

If you choose the second option, once you start the scenario, a MySQL server is installed and started, and then the Sakila schema and data are loaded. When it's ready, a standard `mysql>` prompt appears, and you can then start querying the sample database. This is certainly the easiest option, and I anticipate that most readers will choose this option; if this sounds good to you, feel free to skip ahead to the next section.

If you prefer to have your own copy of the data and want any changes you have made to be permanent, or if you are just interested in installing the MySQL server on your

own machine, you may prefer the first option. You may also opt to use a MySQL server hosted in an environment such as Amazon Web Services or Google Cloud. In either case, you will need to perform the installation/configuration yourself, as it is beyond the scope of this book. Once your database is available, you will need to follow a few steps to load the Sakila sample database.

First, you will need to launch the `mysql` command-line client and provide a password, and then perform the following steps:

1. Go to <https://dev.mysql.com/doc/index-other.html> and download the files for “sakila database” under the Example Databases section.
2. Put the files in a local directory such as `C:\temp\sakila-db` (used for the next two steps, but overwrite with your directory path).
3. Type `source c:\temp\sakila-db\sakila-schema.sql`; and press Enter.
4. Type `source c:\temp\sakila-db\sakila-data.sql`; and press Enter.

You should now have a working database populated with all the data needed for the examples in this book.



The Sakila sample database is made available by MySQL and is licensed via the New BSD license. Sakila contains data for a fictitious movie rental company, and includes tables such as `store`, `inventory`, `film`, `customer`, and `payment`. While actual movie rental stores are largely a thing of the past, with a little imagination we could rebrand it as a movie-streaming company by ignoring the `staff` and `address` tables and renaming `store` to `streaming_service`. However, the examples in this book will stick to the original script (pun intended).

Using the `mysql` Command-Line Tool

Unless you are using a temporary database session (the second option in the previous section), you will need to start the `mysql` command-line tool in order to interact with the database. To do so, you will need to open a Windows or Unix shell and execute the `mysql` utility. For example, if you are logging in using the root account, you would do the following:

```
mysql -u root -p;
```

You will then be asked for your password, after which you will see the `mysql>` prompt. To see all of the available databases, you can use the following command:

```
mysql> show databases;
+-----+
| Database      |
+-----+
| information_schema |
| mysql          |
| performance_schema |
| sakila         |
| sys            |
+-----+
5 rows in set (0.01 sec)
```

Since you will be using the Sakila database, you will need to specify the database you want to work with via the `use` command:

```
mysql> use sakila;
Database changed
```

Whenever you invoke the `mysql` command-line tool, you can specify both the user-name and database to use, as in the following:

```
mysql -u root -p sakila;
```

This will save you from having to type `use sakila;` every time you start up the tool. Now that you have established a session and specified the database, you will be able to issue SQL statements and view the results. For example, if you want to know the current date and time, you could issue the following query:

```
mysql> SELECT now();
+-----+
| now()      |
+-----+
| 2019-04-04 20:44:26 |
+-----+
1 row in set (0.01 sec)
```

The `now()` function is a built-in MySQL function that returns the current date and time. As you can see, the `mysql` command-line tool formats the results of your queries within a rectangle bounded by +, -, and | characters. After the results have been exhausted (in this case, there is only a single row of results), the `mysql` command-line tool shows how many rows were returned, along with how long the SQL statement took to execute.

About Missing from Clauses

With some database servers, you won't be able to issue a query without a `from` clause that names at least one table. Oracle Database is a commonly used server for which this is true. For cases when you only need to call a function, Oracle provides a table called `dual`, which consists of a single column called `dummy` that contains a single row of data. In order to be compatible with Oracle Database, MySQL also provides a `dual`

table. The previous query to determine the current date and time could therefore be written as:

```
mysql> SELECT now()
      FROM dual;
+-----+
| now()          |
+-----+
| 2019-04-04 20:44:26 |
+-----+
1 row in set (0.01 sec)
```

If you are not using Oracle and have no need to be compatible with Oracle, you can ignore the `dual` table altogether and use just a `select` clause without a `from` clause.

When you are done with the `mysql` command-line tool, simply type `quit;` or `exit;` to return to the Unix or Windows command shell.

MySQL Data Types

In general, all the popular database servers have the capacity to store the same types of data, such as strings, dates, and numbers. Where they typically differ is in the specialty data types, such as XML and JSON documents or spatial data. Since this is an introductory book on SQL and since 98% of the columns you encounter will be simple data types, this chapter covers only the character, date (a.k.a. temporal), and numeric data types. The use of SQL to query JSON documents will be explored in [Chapter 18](#).

Character Data

Character data can be stored as either fixed-length or variable-length strings; the difference is that fixed-length strings are right-padded with spaces and always consume the same number of bytes, and variable-length strings are not right-padded with spaces and don't always consume the same number of bytes. When defining a character column, you must specify the maximum size of any string to be stored in the column. For example, if you want to store strings up to 20 characters in length, you could use either of the following definitions:

```
char(20) /* fixed-length */
varchar(20) /* variable-length */
```

The maximum length for `char` columns is currently 255 bytes, whereas `varchar` columns can be up to 65,535 bytes. If you need to store longer strings (such as emails, XML documents, etc.), then you will want to use one of the text types (`mediumtext` and `longtext`), which I cover later in this section. In general, you should use the `char` type when all strings to be stored in the column are of the same length, such as state

abbreviations, and the `varchar` type when strings to be stored in the column are of varying lengths. Both `char` and `varchar` are used in a similar fashion in all the major database servers.



An exception is made in the use of `varchar` for Oracle Database. Oracle users should use the `varchar2` type when defining variable-length character columns.

Character sets

For languages that use the Latin alphabet, such as English, there is a sufficiently small number of characters such that only a single byte is needed to store each character. Other languages, such as Japanese and Korean, contain large numbers of characters, thus requiring multiple bytes of storage for each character. Such character sets are therefore called *multibyte character sets*.

MySQL can store data using various character sets, both single- and multibyte. To view the supported character sets in your server, you can use the `SHOW` command, as shown in the following example:

```
mysql> SHOW CHARACTER SET;
```

Charset	Description	Default collation	Maxlen
armSCII8	ARMSCII-8 Armenian	armSCII8_general_ci	1
ascii	US ASCII	ascii_general_ci	1
big5	Big5 Traditional Chinese	big5_chinese_ci	2
binary	Binary pseudo charset	binary	1
cp1250	Windows Central European	cp1250_general_ci	1
cp1251	Windows Cyrillic	cp1251_general_ci	1
cp1256	Windows Arabic	cp1256_general_ci	1
cp1257	Windows Baltic	cp1257_general_ci	1
cp850	DOS West European	cp850_general_ci	1
cp852	DOS Central European	cp852_general_ci	1
cp866	DOS Russian	cp866_general_ci	1
cp932	SJIS for Windows Japanese	cp932_japanese_ci	2
dec8	DEC West European	dec8_swedish_ci	1
eucjpm	UJIS for Windows Japanese	eucjpm_s_japanese_ci	3
euckr	EUC-KR Korean	euckr_korean_ci	2
gb18030	China National Standard GB18030	gb18030_chinese_ci	4
gb2312	GB2312 Simplified Chinese	gb2312_chinese_ci	2
gbk	GBK Simplified Chinese	gbk_chinese_ci	2
geostd8	GEOSTD8 Georgian	geostd8_general_ci	1
greek	ISO 8859-7 Greek	greek_general_ci	1
hebrew	ISO 8859-8 Hebrew	hebrew_general_ci	1
hp8	HP West European	hp8_english_ci	1
keybcs2	DOS Kamenicky Czech-Slovak	keybcs2_general_ci	1
koi8r	KOI8-R Relcom Russian	koi8r_general_ci	1

koi8u	KOI8-U Ukrainian	koi8u_general_ci	1
latin1	cp1252 West European	latin1_swedish_ci	1
latin2	ISO 8859-2 Central European	latin2_general_ci	1
latin5	ISO 8859-9 Turkish	latin5_turkish_ci	1
latin7	ISO 8859-13 Baltic	latin7_general_ci	1
macce	Mac Central European	macce_general_ci	1
macroman	Mac West European	macroman_general_ci	1
sjis	Shift-JIS Japanese	sjis_japanese_ci	2
swe7	7bit Swedish	swe7_swedish_ci	1
tis620	TIS620 Thai	tis620_thai_ci	1
ucs2	UCS-2 Unicode	ucs2_general_ci	2
ujis	EUC-JP Japanese	ujis_japanese_ci	3
utf16	UTF-16 Unicode	utf16_general_ci	4
utf16le	UTF-16LE Unicode	utf16le_general_ci	4
utf32	UTF-32 Unicode	utf32_general_ci	4
utf8	UTF-8 Unicode	utf8_general_ci	3
utf8mb4	UTF-8 Unicode	utf8mb4_0900_ai_ci	4

41 rows in set (0.04 sec)

If the value in the fourth column, `maxlen`, is greater than 1, then the character set is a multibyte character set.

In prior versions of the MySQL server, the `latin1` character set was automatically chosen as the default character set, but version 8 defaults to `utf8mb4`. However, you may choose to use a different character set for each character column in your database, and you can even store different character sets within the same table. To choose a character set other than the default when defining a column, simply name one of the supported character sets after the type definition, as in:

```
varchar(20) character set latin1
```

With MySQL, you may also set the default character set for your entire database:

```
create database european_sales character set latin1;
```

While this is as much information regarding character sets as is appropriate for an introductory book, there is a great deal more to the topic of internationalization than what is shown here. If you plan to deal with multiple or unfamiliar character sets, you may want to pick up a book such as Jukka Korpela's *Unicode Explained: Internationalize Documents, Programs, and Web Sites* (O'Reilly).

Text data

If you need to store data that might exceed the 64 KB limit for `varchar` columns, you will need to use one of the text types.

Table 2-1 shows the available text types and their maximum sizes.

Table 2-1. MySQL text types

Text type	Maximum number of bytes
tinytext	255
text	65,535
mediumtext	16,777,215
longtext	4,294,967,295

When choosing to use one of the text types, you should be aware of the following:

- If the data being loaded into a text column exceeds the maximum size for that type, the data will be truncated.
- Trailing spaces will not be removed when data is loaded into the column.
- When using `text` columns for sorting or grouping, only the first 1,024 bytes are used, although this limit may be increased if necessary.
- The different text types are unique to MySQL. SQL Server has a single `text` type for large character data, whereas DB2 and Oracle use a data type called `clob`, for Character Large Object.
- Now that MySQL allows up to 65,535 bytes for `varchar` columns (it was limited to 255 bytes in version 4), there isn't any particular need to use the `tinytext` or `text` type.

If you are creating a column for free-form data entry, such as a `notes` column to hold data about customer interactions with your company's customer service department, then `varchar` will probably be adequate. If you are storing documents, however, you should choose either the `mediumtext` or `longtext` type.



Oracle Database allows up to 2,000 bytes for `char` columns and 4,000 bytes for `varchar2` columns. For larger documents you may use the `clob` type. SQL Server can handle up to 8,000 bytes for both `char` and `varchar` data, but you can store up to 2 GB of data in a column defined as `varchar(max)`.

Numeric Data

Although it might seem reasonable to have a single numeric data type called "numeric," there are actually several different numeric data types that reflect the various ways in which numbers are used, as illustrated here:

A column indicating whether a customer order has been shipped

This type of column, referred to as a *Boolean*, would contain a `0` to indicate `false` and a `1` to indicate `true`.

A system-generated primary key for a transaction table

This data would generally start at 1 and increase in increments of one up to a potentially very large number.

An item number for a customer's electronic shopping basket

The values for this type of column would be positive whole numbers between 1 and, perhaps, 200 (for shopaholics).

Positional data for a circuit board drill machine

High-precision scientific or manufacturing data often requires accuracy to eight decimal points.

To handle these types of data (and more), MySQL has several different numeric data types. The most commonly used numeric types are those used to store whole numbers, or *integers*. When specifying one of these types, you may also specify that the data is *unsigned*, which tells the server that all data stored in the column will be greater than or equal to zero. **Table 2-2** shows the five different data types used to store whole-number integers.

Table 2-2. MySQL integer types

Type	Signed range	Unsigned range
tinyint	-128 to 127	0 to 255
smallint	-32,768 to 32,767	0 to 65,535
mediumint	-8,388,608 to 8,388,607	0 to 16,777,215
int	-2,147,483,648 to 2,147,483,647	0 to 4,294,967,295
bigint	-2 ⁶³ to 2 ⁶³ - 1	0 to 2 ⁶⁴ - 1

When you create a column using one of the integer types, MySQL will allocate an appropriate amount of space to store the data, which ranges from one byte for a *tinyint* to eight bytes for a *bigint*. Therefore, you should try to choose a type that will be large enough to hold the biggest number you can envision being stored in the column without needlessly wasting storage space.

For floating-point numbers (such as 3.1415927), you may choose from the numeric types shown in **Table 2-3**.

Table 2-3. MySQL floating-point types

Type	Numeric range
float(p , s)	-3.402823466E+38 to -1.175494351E-38 and 1.175494351E-38 to 3.402823466E+38
double(p , s)	-1.7976931348623157E+308 to -2.2250738585072014E-308 and 2.2250738585072014E-308 to 1.7976931348623157E+308

When using a floating-point type, you can specify a *precision* (the total number of allowable digits both to the left and to the right of the decimal point) and a *scale* (the number of allowable digits to the right of the decimal point), but they are not required. These values are represented in [Table 2-3](#) as *p* and *s*. If you specify a precision and scale for your floating-point column, remember that the data stored in the column will be rounded if the number of digits exceeds the scale and/or precision of the column. For example, a column defined as `float(4,2)` will store a total of four digits, two to the left of the decimal and two to the right of the decimal. Therefore, such a column would handle the numbers 27.44 and 8.19 just fine, but the number 17.8675 would be rounded to 17.87, and attempting to store the number 178.375 in your `float(4,2)` column would generate an error.

Like the integer types, floating-point columns can be defined as `unsigned`, but this designation only prevents negative numbers from being stored in the column rather than altering the range of data that may be stored in the column.

Temporal Data

Along with strings and numbers, you will almost certainly be working with information about dates and/or times. This type of data is referred to as *temporal*, and some examples of temporal data in a database include:

- The future date that a particular event is expected to happen, such as shipping a customer's order
- The date that a customer's order was shipped
- The date and time that a user modified a particular row in a table
- An employee's birth date
- The year corresponding to a row in a `yearly_sales` fact table in a data warehouse
- The elapsed time needed to complete a wiring harness on an automobile assembly line

MySQL includes data types to handle all of these situations. [Table 2-4](#) shows the temporal data types supported by MySQL.

Table 2-4. MySQL temporal types

Type	Default format	Allowable values
<code>date</code>	<code>YYYY-MM-DD</code>	<code>1000-01-01</code> to <code>9999-12-31</code>
<code>datetime</code>	<code>YYYY-MM-DD HH:MI:SS</code>	<code>1000-01-01 00:00:00.000000</code> to <code>9999-12-31 23:59:59.999999</code>

Type	Default format	Allowable values
timestamp	YYYY-MM-DD HH:MI:SS	1970-01-01 00:00:00.000000 to 2038-01-18 22:14:07.999999
year	YYYY	1901 to 2155
time	HHH:MI:SS	-838:59:59.000000 to 838:59:59.000000

While database servers store temporal data in various ways, the purpose of a format string (second column of [Table 2-4](#)) is to show how the data will be represented when retrieved, along with how a date string should be constructed when inserting or updating a temporal column. Thus, if you wanted to insert the date March 23, 2020, into a `date` column using the default format `YYYY-MM-DD`, you would use the string '`2020-03-23`'. [Chapter 7](#) fully explores how temporal data is constructed and displayed.

The `datetime`, `timestamp`, and `time` types also allow fractional seconds of up to 6 decimal places (microseconds). When defining columns using one of these data types, you may supply a value from 0 to 6; for example, specifying `datetime(2)` would allow your time values to include hundredths of a second.



Each database server allows a different range of dates for temporal columns. Oracle Database accepts dates ranging from 4712 BC to 9999 AD, while SQL Server only handles dates ranging from 1753 AD to 9999 AD (unless you are using SQL Server 2008's `datetime2` data type, which allows for dates ranging from 1 AD to 9999 AD). MySQL falls in between Oracle and SQL Server and can store dates from 1000 AD to 9999 AD. Although this might not make any difference for most systems that track current and future events, it is important to keep in mind if you are storing historical dates.

[Table 2-5](#) describes the various components of the date formats shown in [Table 2-4](#).

Table 2-5. Date format components

Component	Definition	Range
YYYY	Year, including century	1000 to 9999
MM	Month	01 (January) to 12 (December)
DD	Day	01 to 31
HH	Hour	00 to 23
HHH	Hours (elapsed)	-838 to 838
MI	Minute	00 to 59
SS	Second	00 to 59

Here's how the various temporal types would be used to implement the examples shown earlier:

- Columns to hold the expected future shipping date of a customer order and an employee's birth date would use the `date` type, since it is unrealistic to schedule a future shipment down to the second and unnecessary to know at what time a person was born.
- A column to hold information about when a customer order was actually shipped would use the `datetime` type, since it is important to track not only the date that the shipment occurred but the time as well.
- A column that tracks when a user last modified a particular row in a table would use the `timestamp` type. The `timestamp` type holds the same information as the `datetime` type (year, month, day, hour, minute, second), but a `timestamp` column will automatically be populated with the current date/time by the MySQL server when a row is added to a table or when a row is later modified.
- A column holding just year data would use the `year` type.
- Columns that hold data regarding the length of time needed to complete a task would use the `time` type. For this type of data, it would be unnecessary and confusing to store a date component, since you are interested only in the number of hours/minutes/seconds needed to complete the task. This information could be derived using two `datetime` columns (one for the task start date/time and the other for the task completion date/time) and subtracting one from the other, but it is simpler to use a single `time` column.

[Chapter 7](#) explores how to work with each of these temporal data types.

Table Creation

Now that you have a firm grasp on what data types may be stored in a MySQL database, it's time to see how to use these types in table definitions. Let's start by defining a table to hold information about a person.

Step 1: Design

A good way to start designing a table is to do a bit of brainstorming to see what kind of information would be helpful to include. Here's what I came up with after thinking for a short time about the types of information that describe a person:

- Name
- Eye color
- Birth date

- Address
- Favorite foods

This is certainly not an exhaustive list, but it's good enough for now. The next step is to assign column names and data types. [Table 2-6](#) shows my initial attempt.

Table 2-6. Person table, first pass

Column	Type	Allowable values
name	varchar(40)	
eye_color	char(2)	BL, BR, GR
birth_date	date	
address	varchar(100)	
favorite_foods	varchar(200)	

The `name`, `address`, and `favorite_foods` columns are of type `varchar` and allow for free-form data entry. The `eye_color` column allows two characters that should equal only `BR`, `BL`, or `GR`. The `birth_date` column is of type `date`, since a time component is not needed.

Step 2: Refinement

In [Chapter 1](#), you were introduced to the concept of *normalization*, which is the process of ensuring that there are no duplicate (other than foreign keys) or compound columns in your database design. In looking at the columns in the `person` table a second time, the following issues arise:

- The `name` column is actually a compound object consisting of a first name and a last name.
- Since multiple people can have the same name, eye color, birth date, and so forth, there are no columns in the `person` table that guarantee uniqueness.
- The `address` column is also a compound object consisting of street, city, state/province, country, and postal code.
- The `favorite_foods` column is a list containing zero, one, or more independent items. It would be best to create a separate table for this data that includes a foreign key to the `person` table so that you know to which person a particular food may be attributed.

After taking these issues into consideration, [Table 2-7](#) gives a normalized version of the `person` table.

Table 2-7. Person table, second pass

Column	Type	Allowable values
person_id	smallint (unsigned)	
first_name	varchar(20)	
last_name	varchar(20)	
eye_color	char(2)	BR, BL, GR
birth_date	date	
street	varchar(30)	
city	varchar(20)	
state	varchar(20)	
country	varchar(20)	
postal_code	varchar(20)	

Now that the `person` table has a primary key (`person_id`) to guarantee uniqueness, the next step is to build a `favorite_food` table that includes a foreign key to the `person` table. [Table 2-8](#) shows the result.

Table 2-8. favorite_food table

Column	Type
person_id	smallint (unsigned)
food	varchar(20)

The `person_id` and `food` columns comprise the primary key of the `favorite_food` table, and the `person_id` column is also a foreign key to the `person` table.

How Much Is Enough?

Moving the `favorite_foods` column out of the `person` table was definitely a good idea, but are we done yet? What happens, for example, if one person lists “pasta” as a favorite food while another person lists “spaghetti”? Are they the same thing? In order to prevent this problem, you might decide that you want people to choose their favorite foods from a list of options, in which case you should create a `food` table with `food_id` and `food_name` columns and then change the `favorite_food` table to contain a foreign key to the `food` table. While this design would be fully normalized, you might decide that you simply want to store the values that the user has entered, in which case you may leave the table as is.

Step 3: Building SQL Schema Statements

Now that the design is complete for the two tables holding information about people and their favorite foods, the next step is to generate SQL statements to create the tables in the database. Here is the statement to create the `person` table:

```
CREATE TABLE person
  (person_id SMALLINT UNSIGNED,
   fname VARCHAR(20),
   lname VARCHAR(20),
   eye_color CHAR(2),
   birth_date DATE,
   street VARCHAR(30),
   city VARCHAR(20),
   state VARCHAR(20),
   country VARCHAR(20),
   postal_code VARCHAR(20),
   CONSTRAINT pk_person PRIMARY KEY (person_id)
);
```

Everything in this statement should be fairly self-explanatory except for the last item; when you define your table, you need to tell the database server what column or columns will serve as the primary key for the table. You do this by creating a *constraint* on the table. You can add several types of constraints to a table definition. This constraint is a *primary key constraint*. It is created on the `person_id` column and given the name `pk_person`.

While on the topic of constraints, there is another type of constraint that would be useful for the `person` table. In [Table 2-6](#), I added a third column to show the allowable values for certain columns (such as 'BR' and 'BL' for the `eye_color` column). Another type of constraint called a *check constraint* constrains the allowable values for a particular column. MySQL allows a check constraint to be attached to a column definition, as in the following:

```
eye_color CHAR(2) CHECK (eye_color IN ('BR','BL','GR')),
```

While check constraints operate as expected on most database servers, the MySQL server allows check constraints to be defined but does not enforce them. However, MySQL does provide another character data type called `enum` that merges the check constraint into the data type definition. Here's what it would look like for the `eye_color` column definition:

```
eye_color ENUM('BR','BL','GR'),
```

Here's how the `person` table definition looks with an `enum` data type for the `eye_color` column:

```

CREATE TABLE person
  (person_id SMALLINT UNSIGNED,
   fname VARCHAR(20),
   lname VARCHAR(20),
   eye_color ENUM('BR','BL','GR'),
   birth_date DATE,
   street VARCHAR(30),
   city VARCHAR(20),
   state VARCHAR(20),
   country VARCHAR(20),
   postal_code VARCHAR(20),
   CONSTRAINT pk_person PRIMARY KEY (person_id)
 );

```

Later in this chapter, you will see what happens if you try to add data to a column that violates its check constraint (or, in the case of MySQL, its enumeration values).

You are now ready to run the `create table` statement using the `mysql` command-line tool. Here's what it looks like:

```

mysql> CREATE TABLE person
    -> (person_id SMALLINT UNSIGNED,
    ->   fname VARCHAR(20),
    ->   lname VARCHAR(20),
    ->   eye_color ENUM('BR','BL','GR'),
    ->   birth_date DATE,
    ->   street VARCHAR(30),
    ->   city VARCHAR(20),
    ->   state VARCHAR(20),
    ->   country VARCHAR(20),
    ->   postal_code VARCHAR(20),
    ->   CONSTRAINT pk_person PRIMARY KEY (person_id)
    -> );
Query OK, 0 rows affected (0.37 sec)

```

After processing the `create table` statement, the MySQL server returns the message “Query OK, 0 rows affected,” which tells me that the statement had no syntax errors.

If you want to make sure that the `person` table does, in fact, exist, you can use the `describe` command (or `desc` for short) to look at the table definition:

```

mysql> desc person;
+-----+-----+-----+-----+-----+
| Field      | Type           | Null | Key | Default | Extra |
+-----+-----+-----+-----+-----+
| person_id  | smallint(5) unsigned | NO   | PRI | NULL    |       |
| fname       | varchar(20)        | YES  |     | NULL    |       |
| lname       | varchar(20)        | YES  |     | NULL    |       |
| eye_color   | enum('BR','BL','GR') | YES  |     | NULL    |       |
| birth_date  | date             | YES  |     | NULL    |       |
| street      | varchar(30)        | YES  |     | NULL    |       |
| city        | varchar(20)        | YES  |     | NULL    |       |
| state        | varchar(20)        | YES  |     | NULL    |       |

```

```

| country      | varchar(20)          | YES   |     | NULL    |     |
| postal_code | varchar(20)          | YES   |     | NULL    |     |
+-----+-----+-----+-----+
10 rows in set (0.00 sec)

```

Columns 1 and 2 of the `describe` output are self-explanatory. Column 3 shows whether a particular column can be omitted when data is inserted into the table. I purposefully left this topic out of the discussion for now (see the following sidebar for a short discourse), but we explore it fully in [Chapter 4](#). The fourth column shows whether a column takes part in any keys (primary or foreign); in this case, the `person_id` column is marked as the primary key. Column 5 shows whether a particular column will be populated with a default value if you omit the column when inserting data into the table. The sixth column (called “Extra”) shows any other pertinent information that might apply to a column.

What Is Null?

In some cases, it is not possible or applicable to provide a value for a particular column in your table. For example, when adding data about a new customer order, the `ship_date` column cannot yet be determined. In this case, the column is said to be *null* (note that I do not say that it *equals* null), which indicates the absence of a value. Null is used for various cases where a value cannot be supplied, such as:

- Not applicable
- Unknown
- Empty set

When designing a table, you may specify which columns are allowed to be null (the default) and which columns are not allowed to be null (designated by adding the keywords `not null` after the type definition).

Now that you have created the `person` table, your next step will be to then create the `favorite_food` table:

```

mysql> CREATE TABLE favorite_food
-> (person_id SMALLINT UNSIGNED,
-> food VARCHAR(20),
-> CONSTRAINT pk_favorite_food PRIMARY KEY (person_id, food),
-> CONSTRAINT fk_fav_food_person_id FOREIGN KEY (person_id)
-> REFERENCES person (person_id)
-> );
Query OK, 0 rows affected (0.10 sec)

```

This should look very similar to the `create table` statement for the `person` table, with the following exceptions:

- Since a person can have more than one favorite food (which is the reason this table was created in the first place), it takes more than just the `person_id` column to guarantee uniqueness in the table. This table, therefore, has a two-column primary key: `person_id` and `food`.
- The `favorite_food` table contains another type of constraint which is called a *foreign key constraint*. This constrains the values of the `person_id` column in the `favorite_food` table to include *only* values found in the `person` table. With this constraint in place, I will not be able to add a row to the `favorite_food` table indicating that `person_id` 27 likes pizza if there isn't already a row in the `person` table having a `person_id` of 27.



If you forget to create the foreign key constraint when you first create the table, you can add it later via the `alter table` statement.

`describe` shows the following after executing the `create table` statement:

```
mysql> desc favorite_food;
+-----+-----+-----+-----+
| Field | Type            | Null | Key | Default | Extra |
+-----+-----+-----+-----+
| person_id | smallint(5) unsigned | NO   | PRI | NULL    |       |
| food      | varchar(20)        | NO   | PRI | NULL    |       |
+-----+-----+-----+-----+
2 rows in set (0.00 sec)
```

Now that the tables are in place, the next logical step is to add some data.

Populating and Modifying Tables

With the `person` and `favorite_food` tables in place, you can now begin to explore the four SQL data statements: `insert`, `update`, `delete`, and `select`.

Inserting Data

Since there is not yet any data in the `person` and `favorite_food` tables, the first of the four SQL data statements to be explored will be the `insert` statement. There are three main components to an `insert` statement:

- The name of the table into which to add the data
- The names of the columns in the table to be populated

- The values with which to populate the columns

You are not required to provide data for every column in the table (unless all the columns in the table have been defined as `not null`). In some cases, those columns that are not included in the initial `insert` statement will be given a value later via an `update` statement. In other cases, a column may never receive a value for a particular row of data (such as a customer order that is canceled before being shipped, thus rendering the `ship_date` column inapplicable).

Generating numeric key data

Before inserting data into the `person` table, it would be useful to discuss how values are generated for numeric primary keys. Other than picking a number out of thin air, you have a couple of options:

- Look at the largest value currently in the table and add one.
- Let the database server provide the value for you.

Although the first option may seem valid, it proves problematic in a multiuser environment, since two users might look at the table at the same time and generate the same value for the primary key. Instead, all database servers on the market today provide a safe, robust method for generating numeric keys. In some servers, such as the Oracle Database, a separate schema object is used (called a *sequence*); in the case of MySQL, however, you simply need to turn on the *auto-increment* feature for your primary key column. Normally, you would do this at table creation, but doing it now provides the opportunity to learn another SQL schema statement, `alter table`, which is used to modify the definition of an existing table:

```
ALTER TABLE person MODIFY person_id SMALLINT UNSIGNED AUTO_INCREMENT;
```



If you are running these statements in your database, you will first need to disable the foreign key constraint on the `favorite_food` table, and then re-enable the constraints when finished. The progression of statements would be:

```
set foreign_key_checks=0;
ALTER TABLE person
    MODIFY person_id SMALLINT UNSIGNED AUTO_INCREMENT;
set foreign_key_checks=1;
```

This statement essentially redefines the `person_id` column in the `person` table. If you describe the table, you will now see the auto-increment feature listed under the “Extra” column for `person_id`:

```
mysql> DESC person;
+-----+-----+-----+-----+-----+
| Field | Type            | Null | Key | Default | Extra          |
+-----+-----+-----+-----+-----+
| person_id | smallint(5) unsigned | NO   | PRI | NULL    | auto_increment |
| .          |                         |      |     |          |                |
| .          |                         |      |     |          |                |
| .          |                         |      |     |          |                |
+-----+-----+-----+-----+-----+
```

When you insert data into the `person` table, you simply provide a `null` value for the `person_id` column, and MySQL will populate the column with the next available number (by default, MySQL starts at 1 for auto-increment columns).

The `insert` statement

Now that all the pieces are in place, it's time to add some data. The following statement creates a row in the `person` table for William Turner:

```
mysql> INSERT INTO person
    -> (person_id, fname, lname, eye_color, birth_date)
    -> VALUES (null, 'William', 'Turner', 'BR', '1972-05-27');
Query OK, 1 row affected (0.22 sec)
```

The feedback ("Query OK, 1 row affected") tells you that your statement syntax was proper and that one row was added to the database (since it was an `insert` statement). You can look at the data just added to the table by issuing a `select` statement:

```
mysql> SELECT person_id, fname, lname, birth_date
    -> FROM person;
+-----+-----+-----+-----+
| person_id | fname   | lname  | birth_date |
+-----+-----+-----+-----+
|         1 | William | Turner | 1972-05-27 |
+-----+-----+-----+-----+
1 row in set (0.06 sec)
```

As you can see, the MySQL server generated a value of 1 for the primary key. Since there is only a single row in the `person` table, I neglected to specify which row I am interested in and simply retrieved all the rows in the table. If there were more than one row in the table, however, I could add a `where` clause to specify that I want to retrieve data only for the row having a value of 1 for the `person_id` column:

```
mysql> SELECT person_id, fname, lname, birth_date
    -> FROM person
    -> WHERE person_id = 1;
+-----+-----+-----+-----+
| person_id | fname   | lname  | birth_date |
+-----+-----+-----+-----+
|         1 | William | Turner | 1972-05-27 |
+-----+-----+-----+-----+
1 row in set (0.00 sec)
```

While this query specifies a particular primary key value, you can use any column in the table to search for rows, as shown by the following query, which finds all rows with a value of 'Turner' for the `lname` column:

```
mysql> SELECT person_id, fname, lname, birth_date
-> FROM person
-> WHERE lname = 'Turner';
+-----+-----+-----+
| person_id | fname   | lname   | birth_date |
+-----+-----+-----+
|       1 | William | Turner | 1972-05-27 |
+-----+-----+-----+
1 row in set (0.00 sec)
```

Before moving on, a couple of things about the earlier `insert` statement are worth mentioning:

- Values were not provided for any of the address columns. This is fine, since `nulls` are allowed for those columns.
- The value provided for the `birth_date` column was a string. As long as you match the required format shown in [Table 2-4](#), MySQL will convert the string to a date for you.
- The column names and the values provided must correspond in number and type. If you name seven columns and provide only six values or if you provide values that cannot be converted to the appropriate data type for the corresponding column, you will receive an error.

William Turner has also provided information about his favorite three foods, so here are three `insert` statements to store his food preferences:

```
mysql> INSERT INTO favorite_food (person_id, food)
-> VALUES (1, 'pizza');
Query OK, 1 row affected (0.01 sec)
mysql> INSERT INTO favorite_food (person_id, food)
-> VALUES (1, 'cookies');
Query OK, 1 row affected (0.00 sec)
mysql> INSERT INTO favorite_food (person_id, food)
-> VALUES (1, 'nachos');
Query OK, 1 row affected (0.01 sec)
```

Here's a query that retrieves William's favorite foods in alphabetical order using an `order by` clause:

```
mysql> SELECT food
-> FROM favorite_food
-> WHERE person_id = 1
-> ORDER BY food;
+-----+
| food   |
+-----+
```

```
+-----+
| cookies |
| nachos  |
| pizza   |
+-----+
3 rows in set (0.02 sec)
```

The `order by` clause tells the server how to sort the data returned by the query. Without the `order by` clause, there is no guarantee that the data in the table will be retrieved in any particular order.

So that William doesn't get lonely, you can execute another `insert` statement to add Susan Smith to the `person` table:

```
mysql> INSERT INTO person
-> (person_id, fname, lname, eye_color, birth_date,
-> street, city, state, country, postal_code)
-> VALUES (null, 'Susan', 'Smith', 'BL', '1975-11-02',
-> '23 Maple St.', 'Arlington', 'VA', 'USA', '20220');
Query OK, 1 row affected (0.01 sec)
```

Since Susan was kind enough to provide her address, we included five more columns than when William's data was inserted. If you query the table again, you will see that Susan's row has been assigned the value 2 for its primary key value:

```
mysql> SELECT person_id, fname, lname, birth_date
-> FROM person;
+-----+-----+-----+-----+
| person_id | fname  | lname  | birth_date |
+-----+-----+-----+-----+
|        1 | William | Turner | 1972-05-27 |
|        2 | Susan   | Smith  | 1975-11-02 |
+-----+-----+-----+-----+
2 rows in set (0.00 sec)
```

Can I Get That in XML?

If you will be working with XML data, you will be happy to know that most database servers provide a simple way to generate XML output from a query. With MySQL, for example, you can use the `--xml` option when invoking the `mysql` tool, and all your output will automatically be formatted using XML. Here's what the favorite-food data looks like as an XML document:

```
C:\database> mysql -u lrngsql -p --xml bank
Enter password: xxxxxx
Welcome to the MySQL Monitor...

Mysql> SELECT * FROM favorite_food;
<?xml version="1.0"?>

<resultset statement="select * from favorite_food"
```

```
xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance">
<row>
    <field name="person_id">1</field>
    <field name="food">cookies</field>
</row>
<row>
    <field name="person_id">1</field>
    <field name="food">nachos</field>
</row>
<row>
    <field name="person_id">1</field>
    <field name="food">pizza</field>
</row>
</resultset>
3 rows in set (0.00 sec)
```

With SQL Server, you don't need to configure your command-line tool; you just need to add the `for xml` clause to the end of your query, as in:

```
SELECT * FROM favorite_food
FOR XML AUTO, ELEMENTS
```

Updating Data

When the data for William Turner was initially added to the table, data for the various address columns was not included in the `insert` statement. The next statement shows how these columns can be populated at a later time via an `update` statement:

```
mysql> UPDATE person
      -> SET street = '1225 Tremont St.',
      ->     city = 'Boston',
      ->     state = 'MA',
      ->     country = 'USA',
      ->     postal_code = '02138'
      -> WHERE person_id = 1;
Query OK, 1 row affected (0.04 sec)
Rows matched: 1  Changed: 1  Warnings: 0
```

The server responded with a two-line message: the “Rows matched: 1” item tells you that the condition in the `where` clause matched a single row in the table, and the “Changed: 1” item tells you that a single row in the table has been modified. Since the `where` clause specifies the primary key of William’s row, this is exactly what you would expect to have happen.

Depending on the conditions in your `where` clause, it is also possible to modify more than one row using a single statement. Consider, for example, what would happen if your `where` clause looked as follows:

```
WHERE person_id < 10
```

Since both William and Susan have a `person_id` value less than 10, both of their rows would be modified. If you leave off the `where` clause altogether, your update statement will modify every row in the table.

Deleting Data

It seems that William and Susan aren't getting along very well together, so one of them has got to go. Since William was there first, Susan will get the boot courtesy of the `delete` statement:

```
mysql> DELETE FROM person  
-> WHERE person_id = 2;  
Query OK, 1 row affected (0.01 sec)
```

Again, the primary key is being used to isolate the row of interest, so a single row is deleted from the table. Like the `update` statement, more than one row can be deleted depending on the conditions in your `where` clause, and all rows will be deleted if the `where` clause is omitted.

When Good Statements Go Bad

So far, all of the SQL data statements shown in this chapter have been well formed and have played by the rules. Based on the table definitions for the `person` and `favorite_food` tables, however, there are lots of ways that you can run afoul when inserting or modifying data. This section shows you some of the common mistakes that you might come across and how the MySQL server will respond.

Nonunique Primary Key

Because the table definitions include the creation of primary key constraints, MySQL will make sure that duplicate key values are not inserted into the tables. The next statement attempts to bypass the auto-increment feature of the `person_id` column and create another row in the `person` table with a `person_id` of 1:

```
mysql> INSERT INTO person  
-> (person_id, fname, lname, eye_color, birth_date)  
-> VALUES (1, 'Charles','Fulton', 'GR', '1968-01-15');  
ERROR 1062 (23000): Duplicate entry '1' for key 'PRIMARY'
```

There is nothing stopping you (with the current schema objects, at least) from creating two rows with identical names, addresses, birth dates, and so on, as long as they have different values for the `person_id` column.

Nonexistent Foreign Key

The table definition for the `favorite_food` table includes the creation of a foreign key constraint on the `person_id` column. This constraint ensures that all values of `person_id` entered into the `favorite_food` table exist in the `person` table. Here's what would happen if you tried to create a row that violates this constraint:

```
mysql> INSERT INTO favorite_food (person_id, food)
-> VALUES (999, 'lasagna');
ERROR 1452 (23000): Cannot add or update a child row: a foreign key constraint
fails ('sakila'.'favorite_food', CONSTRAINT 'fk_fav_food_person_id' FOREIGN KEY
('person_id') REFERENCES 'person' ('person_id'))
```

In this case, the `favorite_food` table is considered the *child* and the `person` table is considered the *parent*, since the `favorite_food` table is dependent on the `person` table for some of its data. If you plan to enter data into both tables, you will need to create a row in *parent* before you can enter data into `favorite_food`.



Foreign key constraints are enforced only if your tables are created using the InnoDB storage engine. We discuss MySQL's storage engines in [Chapter 12](#).

Column Value Violations

The `eye_color` column in the `person` table is restricted to the values 'BR' for brown, 'BL' for blue, and 'GR' for green. If you mistakenly attempt to set the value of the column to any other value, you will receive the following response:

```
mysql> UPDATE person
-> SET eye_color = 'ZZ'
-> WHERE person_id = 1;
ERROR 1265 (01000): Data truncated for column 'eye_color' at row 1
```

The error message is a bit confusing, but it gives you the general idea that the server is unhappy about the value provided for the `eye_color` column.

Invalid Date Conversions

If you construct a string with which to populate a date column and that string does not match the expected format, you will receive another error. Here's an example that uses a date format that does not match the default date format of YYYY-MM-DD:

```
mysql> UPDATE person
-> SET birth_date = 'DEC-21-1980'
-> WHERE person_id = 1;
ERROR 1292 (22007): Incorrect date value: 'DEC-21-1980' for column 'birth_date'
at row 1
```

In general, it is always a good idea to explicitly specify the format string rather than relying on the default format. Here's another version of the statement that uses the `str_to_date` function to specify which format string to use:

```
mysql> UPDATE person  
    -> SET birth_date = str_to_date('DEC-21-1980' , '%b-%d-%Y')  
    -> WHERE person_id = 1;  
Query OK, 1 row affected (0.12 sec)  
Rows matched: 1  Changed: 1  Warnings: 0
```

Not only is the database server happy, but William is happy as well (we just made him eight years younger, without the need for expensive cosmetic surgery!).



Earlier in the chapter, when I discussed the various temporal data types, I showed date-formatting strings such as `YYYY-MM-DD`. While many database servers use this style of formatting, MySQL uses `%Y` to indicate a four-character year. Here are a few more formatters that you might need when converting strings to `datetimes` in MySQL:

- `%a` The short weekday name, such as Sun, Mon, ...
- `%b` The short month name, such as Jan, Feb, ...
- `%c` The numeric month (00..12)
- `%d` The numeric day of the month (00..31)
- `%f` The number of microseconds (000000..999999)
- `%H` The hour of the day, in 24-hour format (00..23)
- `%h` The hour of the day, in 12-hour format (01..12)
- `%i` The minutes within the hour (00..59)
- `%j` The day of year (001..366)
- `%M` The full month name (January..December)
- `%m` The numeric month
- `%p` AM or PM
- `%s` The number of seconds (00..59)
- `%W` The full weekday name (Sunday..Saturday)
- `%w` The numeric day of the week (0=Sunday..6=Saturday)
- `%Y` The four-digit year

The Sakila Database

For the remainder of the book, most examples will use a sample database called Sakila, which is made available by the nice people at MySQL. This database models a chain of DVD rental stores, which is a bit outdated, but with a bit of imagination it can be rebranded as a video-streaming company. Some of the tables include `customer`, `film`, `actor`, `payment`, `rental`, and `category`. The entire schema and example data should have been created when you followed the final steps at the beginning of the chapter for loading the MySQL server and generating the sample data. For a diagram of the tables and their columns and relationships, see [Appendix A](#).

Table 2-9 shows some of the tables used in the Sakila schema, along with short definitions of each.

Table 2-9. Sakila schema definitions

Table name	Definition
film	A movie that has been released and can be rented
actor	A person who acts in films
customer	A person who watches films
category	A genre of films
payment	A rental of a film by a customer
language	A language spoken by the actors of a film
film_actor	An actor in a film
inventory	A film available for rental

Feel free to experiment with the tables as much as you want, including adding your own tables to expand the business functions. You can always drop the database and re-create it from the downloaded file if you want to make sure your sample data is intact. If you are using the temporary session, any changes you make will be lost when the session closes, so you may want to keep a script of your changes so you can re-create any changes you have made.

If you want to see the tables available in your database, you can use the `show tables` command, as in:

```
mysql> show tables;
+-----+
| Tables_in_sakila |
+-----+
| actor           |
| actor_info      |
| address         |
| category        |
| city            |
| country         |
| customer        |
| customer_list   |
| film            |
| film_actor      |
| film_category   |
| film_list       |
| film_text       |
| inventory       |
| language         |
| nicer_but_slower_film_list |
| payment          |
| rental           |
+-----+
```

```
| sales_by_film_category      |
| sales_by_store               |
| staff                         |
| staff_list                     |
| store                          |
+-----+
23 rows in set (0.02 sec)
```

Along with the 23 tables in the Sakila schema, your table listing may also include the two tables created in this chapter: `person` and `favorite_food`. These tables will not be used in later chapters, so feel free to drop them by issuing the following set of commands:

```
mysql> DROP TABLE favorite_food;
Query OK, 0 rows affected (0.56 sec)
mysql> DROP TABLE person;
Query OK, 0 rows affected (0.05 sec)
```

If you want to look at the columns in a table, you can use the `describe` command. Here's an example of the `describe` output for the `customer` table:

```
mysql> desc customer;
+-----+-----+-----+-----+-----+-----+
| Field      | Type       | Null | Key | Default | Extra          |
+-----+-----+-----+-----+-----+-----+
| customer_id | smallint(5) | NO   | PRI | NULL    | auto_increment |
|              | unsigned     |      |      |          |                |
| store_id    | tinyint(3)  | NO   | MUL | NULL    |                |
|              | unsigned     |      |      |          |                |
| first_name  | varchar(45) | NO   |      | NULL    |                |
| last_name   | varchar(45) | NO   | MUL | NULL    |                |
| email        | varchar(50)  | YES  |      | NULL    |                |
| address_id  | smallint(5) | NO   | MUL | NULL    |                |
|              | unsigned     |      |      |          |                |
| active       | tinyint(1)  | NO   |      | 1       |                |
| create_date  | datetime    | NO   |      | NULL    |                |
| last_update  | timestamp   | YES  |      | CURRENT_ | DEFAULT_GENERATED on
|              |             |      |      | TIMESTAMP | update CURRENT_
|              |             |      |      |           | TIMESTAMP      |
+-----+-----+-----+-----+-----+-----+
```

The more comfortable you are with the example database, the better you will understand the examples and, consequently, the concepts in the following chapters.

CHAPTER 3

Query Primer

So far, you have seen a few examples of database queries (a.k.a. `select` statements) sprinkled throughout the first two chapters. Now it's time to take a closer look at the different parts of the `select` statement and how they interact. After finishing this chapter, you should have a basic understanding of how data is retrieved, joined, filtered, grouped, and sorted; these topics will be covered in detail in Chapters 4 through 10.

Query Mechanics

Before dissecting the `select` statement, it might be interesting to look at how queries are executed by the MySQL server (or, for that matter, any database server). If you are using the `mysql` command-line tool (which I assume you are), then you have already logged in to the MySQL server by providing your username and password (and possibly a hostname if the MySQL server is running on a different computer). Once the server has verified that your username and password are correct, a *database connection* is generated for you to use. This connection is held by the application that requested it (which, in this case, is the `mysql` tool) until the application releases the connection (i.e., as a result of typing `quit`) or the server closes the connection (i.e., when the server is shut down). Each connection to the MySQL server is assigned an identifier, which is shown to you when you first log in:

```
Welcome to the MySQL monitor. Commands end with ; or \g.  
Your MySQL connection id is 11  
Server version: 8.0.15 MySQL Community Server - GPL
```

Copyright (c) 2000, 2019, Oracle and/or its affiliates. All rights reserved.
Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective

```
owners.  
Type 'help;' or '\h' for help. Type '\c' to clear the buffer.
```

In this case, my connection ID is 11. This information might be useful to your database administrator if something goes awry, such as a malformed query that runs for hours, so you might want to jot it down.

Once the server has verified your username and password and issued you a connection, you are ready to execute queries (along with other SQL statements). Each time a query is sent to the server, the server checks the following things prior to statement execution:

- Do you have permission to execute the statement?
- Do you have permission to access the desired data?
- Is your statement syntax correct?

If your statement passes these three tests, then your query is handed to the *query optimizer*, whose job it is to determine the most efficient way to execute your query. The optimizer looks at such things as the order in which to join the tables named in your `from` clause and what indexes are available, and then it picks an *execution plan*, which the server uses to execute your query.



Understanding and influencing how your database server chooses execution plans is a fascinating topic that many of you will want to explore. For those readers using MySQL, you might consider reading Baron Schwartz et al's *High Performance MySQL* (O'Reilly). Among other things, you will learn how to generate indexes, analyze execution plans, influence the optimizer via query hints, and tune your server's startup parameters. If you are using Oracle Database or SQL Server, dozens of tuning books are available.

Once the server has finished executing your query, the *result set* is returned to the calling application (which is, once again, the `mysql` tool). As I mentioned in [Chapter 1](#), a result set is just another table containing rows and columns. If your query fails to yield any results, the `mysql` tool will show you the message found at the end of the following example:

```
mysql> SELECT first_name, last_name  
-> FROM customer  
-> WHERE last_name = 'ZIEGLER';  
Empty set (0.02 sec)
```

If the query returns one or more rows, the `mysql` tool will format the results by adding column headers and constructing boxes around the columns using the `-`, `|`, and `+` symbols, as shown in the next example:

```

mysql> SELECT *
      -> FROM category;
+-----+-----+-----+
| category_id | name      | last_update          |
+-----+-----+-----+
|      1 | Action    | 2006-02-15 04:46:27 |
|      2 | Animation | 2006-02-15 04:46:27 |
|      3 | Children   | 2006-02-15 04:46:27 |
|      4 | Classics  | 2006-02-15 04:46:27 |
|      5 | Comedy    | 2006-02-15 04:46:27 |
|      6 | Documentary | 2006-02-15 04:46:27 |
|      7 | Drama     | 2006-02-15 04:46:27 |
|      8 | Family    | 2006-02-15 04:46:27 |
|      9 | Foreign   | 2006-02-15 04:46:27 |
|     10 | Games     | 2006-02-15 04:46:27 |
|     11 | Horror    | 2006-02-15 04:46:27 |
|     12 | Music     | 2006-02-15 04:46:27 |
|     13 | New       | 2006-02-15 04:46:27 |
|     14 | Sci-Fi    | 2006-02-15 04:46:27 |
|     15 | Sports    | 2006-02-15 04:46:27 |
|     16 | Travel    | 2006-02-15 04:46:27 |
+-----+-----+-----+
16 rows in set (0.02 sec)

```

This query returns all three columns for of all the rows in the `category` table. After the last row of data is displayed, the `mysql` tool displays a message telling you how many rows were returned, which, in this case, is 16.

Query Clauses

Several components or *clauses* make up the `select` statement. While only one of them is mandatory when using MySQL (the `select` clause), you will usually include at least two or three of the six available clauses. [Table 3-1](#) shows the different clauses and their purposes.

Table 3-1. Query clauses

Clause name	Purpose
<code>select</code>	Determines which columns to include in the query's result set
<code>from</code>	Identifies the tables from which to retrieve data and how the tables should be joined
<code>where</code>	Filters out unwanted data
<code>group by</code>	Used to group rows together by common column values
<code>having</code>	Filters out unwanted groups
<code>order by</code>	Sorts the rows of the final result set by one or more columns

All of the clauses shown in [Table 3-1](#) are included in the ANSI specification. The following sections delve into the uses of the six major query clauses.

The select Clause

Even though the `select` clause is the first clause of a `select` statement, it is one of the last clauses that the database server evaluates. The reason for this is that before you can determine what to include in the final result set, you need to know all of the possible columns that *could* be included in the final result set. In order to fully understand the role of the `select` clause, therefore, you will need to understand a bit about the `from` clause. Here's a query to get started:

```
mysql> SELECT *  
      -> FROM language;  
+-----+-----+-----+  
| language_id | name      | last_update          |  
+-----+-----+-----+  
| 1 | English   | 2006-02-15 05:02:19 |  
| 2 | Italian    | 2006-02-15 05:02:19 |  
| 3 | Japanese   | 2006-02-15 05:02:19 |  
| 4 | Mandarin   | 2006-02-15 05:02:19 |  
| 5 | French     | 2006-02-15 05:02:19 |  
| 6 | German     | 2006-02-15 05:02:19 |  
+-----+-----+-----+  
6 rows in set (0.03 sec)
```

In this query, the `from` clause lists a single table (`language`), and the `select` clause indicates that *all* columns (designated by `*`) in the `language` table should be included in the result set. This query could be described in English as follows:

Show me all the columns and all the rows in the language table.

In addition to specifying all the columns via the asterisk character, you can explicitly name the columns you are interested in, such as:

```
mysql> SELECT language_id, name, last_update  
      -> FROM language;  
+-----+-----+-----+  
| language_id | name      | last_update          |  
+-----+-----+-----+  
| 1 | English   | 2006-02-15 05:02:19 |  
| 2 | Italian    | 2006-02-15 05:02:19 |  
| 3 | Japanese   | 2006-02-15 05:02:19 |  
| 4 | Mandarin   | 2006-02-15 05:02:19 |  
| 5 | French     | 2006-02-15 05:02:19 |  
| 6 | German     | 2006-02-15 05:02:19 |  
+-----+-----+-----+  
6 rows in set (0.00 sec)
```

The results are identical to the first query, since all the columns in the `language` table (`language_id`, `name`, and `last_update`) are named in the `select` clause. You can choose to include only a subset of the columns in the `language` table as well:

```

mysql> SELECT name
      -> FROM language;
+-----+
| name   |
+-----+
| English |
| Italian |
| Japanese |
| Mandarin |
| French  |
| German  |
+-----+
6 rows in set (0.00 sec)

```

The job of the `select` clause, therefore, is as follows:

The `select` clause determines which of all possible columns should be included in the query's result set.

If you were limited to including only columns from the table or tables named in the `from` clause, things would be rather dull. However, you can spice things up in your `select` clause by including things such as:

- Literals, such as numbers or strings
- Expressions, such as `transaction.amount * -1`
- Built-in function calls, such as `ROUND(transaction.amount, 2)`
- User-defined function calls

The next query demonstrates the use of a table column, a literal, an expression, and a built-in function call in a single query against the `employee` table:

```

mysql> SELECT language_id,
      ->   'COMMON' language_usage,
      ->   language_id * 3.1415927 lang_pi_value,
      ->   upper(name) language_name
      -> FROM language;
+-----+-----+-----+-----+
| language_id | language_usage | lang_pi_value | language_name |
+-----+-----+-----+-----+
|       1 | COMMON        | 3.1415927    | ENGLISH      |
|       2 | COMMON        | 6.2831854    | ITALIAN      |
|       3 | COMMON        | 9.4247781    | JAPANESE    |
|       4 | COMMON        | 12.5663708   | MANDARIN    |
|       5 | COMMON        | 15.7079635   | FRENCH      |
|       6 | COMMON        | 18.8495562   | GERMAN      |
+-----+-----+-----+-----+
6 rows in set (0.04 sec)

```

We cover expressions and built-in functions in detail later, but I wanted to give you a feel for what kinds of things can be included in the `select` clause. If you only need to

execute a built-in function or evaluate a simple expression, you can skip the `from` clause entirely. Here's an example:

```
mysql> SELECT version(),
->       user(),
->       database();
+-----+-----+-----+
| version() | user()      | database() |
+-----+-----+-----+
| 8.0.15   | root@localhost | sakila    |
+-----+-----+-----+
1 row in set (0.00 sec)
```

Since this query simply calls three built-in functions and doesn't retrieve data from any tables, there is no need for a `from` clause.

Column Aliases

Although the `mysql` tool will generate labels for the columns returned by your queries, you may want to assign your own labels. While you might want to assign a new label to a column from a table (if it is poorly or ambiguously named), you will almost certainly want to assign your own labels to those columns in your result set that are generated by expressions or built-in function calls. You can do so by adding a *column alias* after each element of your `select` clause. Here's the previous query against the `language` table, which included column aliases for three of the columns:

```
mysql> SELECT language_id,
->       'COMMON' language_usage,
->       language_id * 3.1415927 lang_pi_value,
->       upper(name) language_name
->     FROM language;
+-----+-----+-----+-----+
| language_id | language_usage | lang_pi_value | language_name |
+-----+-----+-----+-----+
|          1 | COMMON        |      3.1415927 | ENGLISH      |
|          2 | COMMON        |      6.2831854 | ITALIAN      |
|          3 | COMMON        |      9.4247781 | JAPANESE    |
|          4 | COMMON        |     12.5663708 | MANDARIN    |
|          5 | COMMON        |     15.7079635 | FRENCH      |
|          6 | COMMON        |     18.8495562 | GERMAN      |
+-----+-----+-----+-----+
6 rows in set (0.04 sec)
```

If you look at the `select` clause, you can see how the column aliases `language_usage`, `lang_pi_value`, and `language_name` are added after the second, third, and fourth columns. I think you will agree that the output is easier to understand with column aliases in place, and it would be easier to work with programmatically if you were issuing the query from within Java or Python rather than interactively via the `mysql`

tool. In order to make your column aliases stand out even more, you also have the option of using the `as` keyword before the alias name, as in:

```
mysql> SELECT language_id,
   ->      'COMMON' AS language_usage,
   ->      language_id * 3.1415927 AS lang_pi_value,
   ->      upper(name) AS language_name
   -> FROM language;
```

Many people feel that including the optional `as` keyword improves readability, although I have chosen not to use it for the examples in this book.

Removing Duplicates

In some cases, a query might return duplicate rows of data. For example, if you were to retrieve the IDs of all actors who appeared in a film, you would see the following:

Since some actors appeared in more than one film, you will see the same actor ID multiple times. What you probably want in this case is the *distinct* set of actors, instead of seeing the actor IDs repeated for each film in which they appeared. You can achieve this by adding the keyword `distinct` directly after the `select` keyword, as demonstrated by the following:

```
mysql> SELECT DISTINCT actor_id FROM film_actor ORDER BY actor_id;
+-----+
| actor_id |
+-----+
|      1 |
|      2 |
|      3 |
|      4 |
|      5 |
|      6 |
|      7 |
|      8 |
|      9 |
|     10 |
...
|   192 |
|   193 |
|   194 |
|   195 |
|   196 |
|   197 |
|   198 |
|   199 |
|   200 |
+-----+
200 rows in set (0.01 sec)
```

The result set now contains 200 rows, one for each distinct actor, rather than 5,462 rows, one for each film appearance by an actor.



If you simply want a list of all actors, you can query the `actor` table rather than reading through all the rows in `film_actor` and removing duplicates.

If you do not want the server to remove duplicate data or you are sure there will be no duplicates in your result set, you can specify the `all` keyword instead of specifying `distinct`. However, the `all` keyword is the default and never needs to be explicitly named, so most programmers do not include `all` in their queries.



Keep in mind that generating a distinct set of results requires the data to be sorted, which can be time consuming for large result sets. Don't fall into the trap of using `distinct` just to be sure there are no duplicates; instead, take the time to understand the data you are working with so that you will know whether duplicates are possible.

The from Clause

Thus far, you have seen queries whose `from` clauses contain a single table. Although most SQL books define the `from` clause as simply a list of one or more tables, I would like to broaden the definition as follows:

The `from` clause defines the tables used by a query, along with the means of linking the tables together.

This definition is composed of two separate but related concepts, which we explore in the following sections.

Tables

When confronted with the term *table*, most people think of a set of related rows stored in a database. While this does describe one type of table, I would like to use the word in a more general way by removing any notion of how the data might be stored and concentrating on just the set of related rows. Four different types of tables meet this relaxed definition:

- Permanent tables (i.e., created using the `create table` statement)
- Derived tables (i.e., rows returned by a subquery and held in memory)
- Temporary tables (i.e., volatile data held in memory)
- Virtual tables (i.e., created using the `create view` statement)

Each of these table types may be included in a query's `from` clause. By now, you should be comfortable with including a permanent table in a `from` clause, so I will briefly describe the other types of tables that can be referenced in a `from` clause.

Derived (subquery-generated) tables

A subquery is a query contained within another query. Subqueries are surrounded by parentheses and can be found in various parts of a `select` statement; within the `from` clause, however, a subquery serves the role of generating a derived table that is visible from all other query clauses and can interact with other tables named in the `from` clause. Here's a simple example:

```
mysql> SELECT concat(cust.last_name, ' ', cust.first_name) full_name
-> FROM
-> (SELECT first_name, last_name, email
->   FROM customer
->   WHERE first_name = 'JESSIE'
-> ) cust;
+-----+
| full_name      |
+-----+
```

```
| BANKS, JESSIE |
| MILAM, JESSIE |
+-----+
2 rows in set (0.00 sec)
```

In this example, a subquery against the `customer` table returns three columns, and the *containing query* references two of the three available columns. The subquery is referenced by the containing query via its alias, which, in this case, is `cust`. The data in `cust` is held in memory for the duration of the query and is then discarded. This is a simplistic and not particularly useful example of a subquery in a `from` clause; you will find detailed coverage of subqueries in [Chapter 9](#).

Temporary tables

Although the implementations differ, every relational database allows the ability to define volatile, or temporary, tables. These tables look just like permanent tables, but any data inserted into a temporary table will disappear at some point (generally at the end of a transaction or when your database session is closed). Here's a simple example showing how actors whose last names start with J can be stored temporarily:

```
mysql> CREATE TEMPORARY TABLE actors_j
-> (actor_id smallint(5),
->   first_name varchar(45),
->   last_name varchar(45)
-> );
Query OK, 0 rows affected (0.00 sec)

mysql> INSERT INTO actors_j
-> SELECT actor_id, first_name, last_name
-> FROM actor
-> WHERE last_name LIKE 'J%';
Query OK, 7 rows affected (0.03 sec)
Records: 7  Duplicates: 0  Warnings: 0

mysql> SELECT * FROM actors_j;
+-----+-----+-----+
| actor_id | first_name | last_name |
+-----+-----+-----+
|      119 | WARREN    | JACKMAN   |
|     131 | JANE      | JACKMAN   |
|       8 | MATTHEW   | JOHANSSON  |
|      64 | RAY        | JOHANSSON |
|     146 | ALBERT    | JOHANSSON |
|      82 | WOODY     | JOLIE      |
|      43 | KIRK      | JOVOVICH  |
+-----+-----+-----+
7 rows in set (0.00 sec)
```

These seven rows are held in memory temporarily and will disappear after your session is closed.



Most database servers also drop the temporary table when the session ends. The exception is Oracle Database, which keeps the definition of the temporary table available for future sessions.

Views

A view is a query that is stored in the data dictionary. It looks and acts like a table, but there is no data associated with a view (this is why I call it a *virtual* table). When you issue a query against a view, your query is merged with the view definition to create a final query to be executed.

To demonstrate, here's a view definition that queries the `employee` table and includes four of the available columns:

```
mysql> CREATE VIEW cust_vw AS
    -> SELECT customer_id, first_name, last_name, active
    -> FROM customer;
Query OK, 0 rows affected (0.12 sec)
```

When the view is created, no additional data is generated or stored: the server simply tucks away the `SELECT` statement for future use. Now that the view exists, you can issue queries against it, as in:

```
mysql> SELECT first_name, last_name
    -> FROM cust_vw
    -> WHERE active = 0;
+-----+-----+
| first_name | last_name |
+-----+-----+
| SANDRA     | MARTIN   |
| JUDITH     | COX       |
| SHEILA     | WELLS    |
| ERICA      | MATTHEWS |
| HEIDI      | LARSON   |
| PENNY      | NEAL      |
| KENNETH    | GOODEN   |
| HARRY      | ARCE      |
| NATHAN     | RUNYON   |
| THEODORE   | CULP      |
| MAURICE    | CRAWLEY  |
| BEN        | EASTER   |
| CHRISTIAN  | JUNG      |
| JIMMIE     | EGGLESTON|
| TERRANCE   | ROUSH    |
+-----+-----+
15 rows in set (0.00 sec)
```

Views are created for various reasons, including to hide columns from users and to simplify complex database designs.

Table Links

The second deviation from the simple `from` clause definition is the mandate that if more than one table appears in the `from` clause, the conditions used to *link* the tables must be included as well. This is not a requirement of MySQL or any other database server, but it is the ANSI-approved method of joining multiple tables, and it is the most portable across the various database servers. We explore joining multiple tables in depth in Chapters 5 and 10, but here's a simple example in case I have piqued your curiosity:

```
mysql> SELECT customer.first_name, customer.last_name,
->     time(rental.rental_date) rental_time
->   FROM customer
->     INNER JOIN rental
->       ON customer.customer_id = rental.customer_id
->      WHERE date(rental.rental_date) = '2005-06-14';
+-----+-----+-----+
| first_name | last_name | rental_time |
+-----+-----+-----+
| JEFFERY    | PINSON    | 22:53:33   |
| ELMER      | NOE        | 22:55:13   |
| MINNIE     | ROMERO    | 23:00:34   |
| MIRIAM     | MCKINNEY  | 23:07:08   |
| DANIEL     | CABRAL    | 23:09:38   |
| TERRANCE   | ROUSH     | 23:12:46   |
| JOYCE      | EDWARDS   | 23:16:26   |
| GWENDOLYN  | MAY        | 23:16:27   |
| CATHERINE  | CAMPBELL  | 23:17:03   |
| MATTHEW    | MAHAN     | 23:25:58   |
| HERMAN     | DEVORE    | 23:35:09   |
| AMBER      | DIXON     | 23:42:56   |
| TERRENCE   | GUNDERSON | 23:47:35   |
| SONIA      | GREGORY   | 23:50:11   |
| CHARLES    | KOWALSKI  | 23:54:34   |
| JEANETTE   | GREENE    | 23:54:46   |
+-----+-----+-----+
16 rows in set (0.01 sec)
```

The previous query displays data from both the `customer` table (`first_name`, `last_name`) and the `rental` table (`rental_date`), so both tables are included in the `from` clause. The mechanism for linking the two tables (referred to as a *join*) is the `customer_id` stored in both the `customer` and `rental` tables. Thus, the database server is instructed to use the value of the `customer_id` column in the `customer` table to find all of the customer's rentals in the `rental` table. Join conditions for the two tables are found in the `on` subclause of the `from` clause; in this case, the join condition

is ON `customer.customer_id = rental.customer_id`. The `where` clause is not part of the join and is only included to keep the result set fairly small, since there are more than 16,000 rows in the `rental` table. Again, please refer to [Chapter 5](#) for a thorough discussion of joining multiple tables.

Defining Table Aliases

When multiple tables are joined in a single query, you need a way to identify which table you are referring to when you reference columns in the `select`, `where`, `group by`, `having`, and `order by` clauses. You have two choices when referencing a table outside the `from` clause:

- Use the entire table name, such as `employee.emp_id`.
- Assign each table an *alias* and use the alias throughout the query.

In the previous query, I chose to use the entire table name in the `select` and `on` clauses. Here's what the same query looks like using table aliases:

```
SELECT c.first_name, c.last_name,
       time(r.rental_date) rental_time
  FROM customer c
 INNER JOIN rental r
    ON c.customer_id = r.customer_id
   WHERE date(r.rental_date) = '2005-06-14';
```

If you look closely at the `from` clause, you will see that the `customer` table is assigned the alias `c`, and the `rental` table is assigned the alias `r`. These aliases are then used in the `on` clause when defining the join condition as well as in the `select` clause when specifying the columns to include in the result set. I hope you will agree that using aliases makes for a more compact statement without causing confusion (as long as your choices for alias names are reasonable). Additionally, you may use the `as` keyword with your table aliases, similar to what was demonstrated earlier for column aliases:

```
SELECT c.first_name, c.last_name,
       time(r.rental_date) rental_time
  FROM customer AS c
 INNER JOIN rental AS r
    ON c.customer_id = r.customer_id
   WHERE date(r.rental_date) = '2005-06-14';
```

I have found that roughly half of the database developers I have worked with use the `as` keyword with their column and table aliases, and half do not.

The where Clause

In some cases, you may want to retrieve all rows from a table, especially for small tables such as `language`. Most of the time, however, you will not want to retrieve *every* row from a table but will want a way to filter out those rows that are not of interest. This is a job for the `where` clause.

The where clause is the mechanism for filtering out unwanted rows from your result set.

For example, perhaps you are interested in renting a film but you are only interested in movies rated G that can be kept for at least a week. The following query employs a `where` clause to retrieve *only* the films meeting these criteria:

```
mysql> SELECT title
    -> FROM film
    -> WHERE rating = 'G' AND rental_duration >= 7;
+-----+
| title           |
+-----+
| BLANKET BEVERLY          |
| BORROWERS BEDAZZLED      |
| BRIDE INTRIGUE          |
| CATCH AMISTAD          |
| CITIZEN SHREK          |
| COLDBLOODED DARLING      |
| CONTROL ANTHEM          |
| CRUELTY UNFORGIVEN      |
| DARN FORRESTER          |
| DESPERATE TRAINSPOTTING  |
| DIARY PANIC            |
| DRACULA CRYSTAL          |
| EMPIRE MALKOVICH        |
| FIREHOUSE VIETNAM        |
| GILBERT PELICAN          |
| GRADUATE LORD           |
| GREASE YOUTH            |
| GUN BONNIE              |
| HOOK CHARIOTS           |
| MARRIED GO              |
| MENAGERIE RUSHMORE       |
| MUSCLE BRIGHT            |
| OPERATION OPERATION      |
| PRIMARY GLASS           |
| REBEL AIRPORT            |
| SPIKING ELEMENT          |
| TRUMAN CRAZY            |
| WAKE JAWS                |
| WAR NOTTING              |
+-----+
29 rows in set (0.00 sec)
```

In this case, the `where` clause filtered out 971 of the 1000 rows in the `film` table. This `where` clause contains two *filter conditions*, but you can include as many conditions as are required; individual conditions are separated using operators such as `and`, `or`, and `not` (see [Chapter 4](#) for a complete discussion of the `where` clause and filter conditions).

Let's see what would happen if you change the operator separating the two conditions from `and` to `or`:

```
mysql> SELECT title
-> FROM film
-> WHERE rating = 'G' OR rental_duration >= 7;
+-----+
| title           |
+-----+
| ACE GOLDFINGER
| ADAPTATION HOLES
| AFFAIR PREJUDICE
| AFRICAN EGG
| ALAMO VIDEOTAPE
| AMISTAD MIDSUMMER
| ANGELS LIFE
| ANNIE IDENTITY
| ...
| WATERSHIP FRONTIER
| WEREWOLF LOLA
| WEST LION
| WESTWARD SEABISCUIT
| WOLVES DESIRE
| WON DARES
| WORKER TARZAN
| YOUNG LANGUAGE
+-----+
340 rows in set (0.00 sec)
```

When you separate conditions using the `and` operator, *all* conditions must evaluate to `true` to be included in the result set; when you use `or`, however, only *one* of the conditions needs to evaluate to `true` for a row to be included, which explains why the size of the result set has jumped from 29 to 340 rows.

So, what should you do if you need to use both `and` and `or` operators in your `where` clause? Glad you asked. You should use parentheses to group conditions together. The next query specifies that only those films that are rated G and are available for 7 or more days, or are rated PG-13 and are available 3 or fewer days, be included in the result set:

```
mysql> SELECT title, rating, rental_duration
-> FROM film
-> WHERE (rating = 'G' AND rental_duration >= 7)
->   OR (rating = 'PG-13' AND rental_duration < 4);
```

```

+-----+-----+-----+
| title          | rating | rental_duration |
+-----+-----+-----+
| ALABAMA DEVIL    | PG-13   |            3 |
| BACKLASH UNDEFEATED | PG-13   |            3 |
| BILKO ANONYMOUS    | PG-13   |            3 |
| BLANKET BEVERLY    | G       |            7 |
| BORROWERS BEDAZZLED | G       |            7 |
| BRIDE INTRIGUE     | G       |            7 |
| CASPER DRAGONFLY    | PG-13   |            3 |
| CATCH AMISTAD      | G       |            7 |
| CITIZEN SHREK       | G       |            7 |
| COLDBLOODED DARLING | G       |            7 |
| ...
| TREASURE COMMAND     | PG-13   |            3 |
| TRUMAN CRAZY        | G       |            7 |
| WAIT CIDER          | PG-13   |            3 |
| WAKE JAWS           | G       |            7 |
| WAR NOTTING         | G       |            7 |
| WORLD LEATHERNECKS  | PG-13   |            3 |
+-----+-----+-----+

```

68 rows in set (0.00 sec)

You should always use parentheses to separate groups of conditions when mixing different operators so that you, the database server, and anyone who comes along later to modify your code will be on the same page.

The group by and having Clauses

All the queries thus far have retrieved raw data without any manipulation. Sometimes, however, you will want to find trends in your data that will require the database server to cook the data a bit before you retrieve your result set. One such mechanism is the `group by` clause, which is used to group data by column values. For example, let's say you wanted to find all of the customers who have rented 40 or more films. Rather than looking through all 16,044 rows in the `rental` table, you can write a query that instructs the server to group all rentals by customer, count the number of rentals for each customer, and then return only those customers whose rental count is at least 40. When using the `group by` clause to generate groups of rows, you may also use the `having` clause, which allows you to filter grouped data in the same way the `where` clause lets you filter raw data.

Here's what the query looks like:

```

mysql> SELECT c.first_name, c.last_name, count(*)
-> FROM customer c
-> INNER JOIN rental r
-> ON c.customer_id = r.customer_id
-> GROUP BY c.first_name, c.last_name
-> HAVING count(*) >= 40;

```

```

+-----+-----+-----+
| first_name | last_name | count(*) |
+-----+-----+-----+
| TAMMY     | SANDERS    |      41 |
| CLARA     | SHAW        |      42 |
| ELEANOR   | HUNT        |      46 |
| SUE        | PETERS      |      40 |
| MARCIA    | DEAN        |      42 |
| WESLEY    | BULL         |      40 |
| KARL      | SEAL         |      45 |
+-----+-----+-----+
7 rows in set (0.03 sec)

```

I wanted to briefly mention these two clauses so that they don't catch you by surprise later in the book, but they are a bit more advanced than the other four `select` clauses. Therefore, I ask that you wait until [Chapter 8](#) for a full description of how and when to use `group by` and `having`.

The order by Clause

In general, the rows in a result set returned from a query are not in any particular order. If you want your result set to be sorted, you will need to instruct the server to sort the results using the `order by` clause:

The `order by` clause is the mechanism for sorting your result set using either raw column data or expressions based on column data.

For example, here's another look at an earlier query that returns all customers who rented a film on June 14, 2005:

```

mysql> SELECT c.first_name, c.last_name,
->    time(r.rental_date) rental_time
->   FROM customer c
->   INNER JOIN rental r
->   ON c.customer_id = r.customer_id
-> WHERE date(r.rental_date) = '2005-06-14';
+-----+-----+-----+
| first_name | last_name | rental_time |
+-----+-----+-----+
| JEFFERY    | PINSON    | 22:53:33    |
| ELMER      | NOE        | 22:55:13    |
| MINNIE    | ROMERO    | 23:00:34    |
| MIRIAM    | MCKINNEY  | 23:07:08    |
| DANIEL    | CABRAL    | 23:09:38    |
| TERRANCE   | ROUSH     | 23:12:46    |
| JOYCE      | EDWARDS   | 23:16:26    |
| GWENDOLYN | MAY        | 23:16:27    |
| CATHERINE  | CAMPBELL  | 23:17:03    |
| MATTHEW    | MAHAN     | 23:25:58    |
| HERMAN    | DEVORE    | 23:35:09    |
| AMBER      | DIXON     | 23:42:56    |
+-----+-----+-----+

```

```

| TERRENCE    | GUNDERSON | 23:47:35      |
| SONIA       | GREGORY    | 23:50:11      |
| CHARLES     | KOWALSKI   | 23:54:34      |
| JEANETTE    | GREENE     | 23:54:46      |
+-----+-----+-----+
16 rows in set (0.01 sec)

```

If you would like the results to be in alphabetical order by last name, you can add the `last_name` column to the `order by` clause:

```

mysql> SELECT c.first_name, c.last_name,
->      time(r.rental_date) rental_time
->  FROM customer c
->  INNER JOIN rental r
->  ON c.customer_id = r.customer_id
-> WHERE date(r.rental_date) = '2005-06-14'
-> ORDER BY c.last_name;
+-----+-----+-----+
| first_name | last_name | rental_time |
+-----+-----+-----+
| DANIEL     | CABRAL    | 23:09:38      |
| CATHERINE  | CAMPBELL  | 23:17:03      |
| HERMAN     | DEVORE    | 23:35:09      |
| AMBER      | DIXON     | 23:42:56      |
| JOYCE      | EDWARDS   | 23:16:26      |
| JEANETTE   | GREENE    | 23:54:46      |
| SONIA      | GREGORY   | 23:50:11      |
| TERENCE    | GUNDERSON | 23:47:35      |
| CHARLES    | KOWALSKI  | 23:54:34      |
| MATTHEW    | MAHAN     | 23:25:58      |
| GWENDOLYN  | MAY       | 23:16:27      |
| MIRIAM     | MCKINNEY  | 23:07:08      |
| ELMER      | NOE        | 22:55:13      |
| JEFFERY    | PINSON    | 22:53:33      |
| MINNIE     | ROMERO   | 23:00:34      |
| TERRANCE   | ROUSH     | 23:12:46      |
+-----+-----+-----+
16 rows in set (0.01 sec)

```

While it is not the case in this example, large customer lists will often contain multiple people having the same last name, so you may want to extend the sort criteria to include the person's first name as well.

You can accomplish this by adding the `first_name` column after the `last_name` column in the `order by` clause:

```

mysql> SELECT c.first_name, c.last_name,
->      time(r.rental_date) rental_time
->  FROM customer c
->  INNER JOIN rental r
->  ON c.customer_id = r.customer_id
-> WHERE date(r.rental_date) = '2005-06-14'
-> ORDER BY c.last_name, c.first_name;

```

```

+-----+-----+-----+
| first_name | last_name | rental_time |
+-----+-----+-----+
| DANIEL     | CABRAL    | 23:09:38   |
| CATHERINE  | CAMPBELL  | 23:17:03   |
| HERMAN     | DEVORE    | 23:35:09   |
| AMBER      | DIXON     | 23:42:56   |
| JOYCE      | EDWARDS   | 23:16:26   |
| JEANETTE   | GREENE    | 23:54:46   |
| SONIA      | GREGORY   | 23:50:11   |
| TERRENCE   | GUNDERSON| 23:47:35   |
| CHARLES    | KOWALSKI  | 23:54:34   |
| MATTHEW    | MAHAN     | 23:25:58   |
| GWENDOLYN  | MAY       | 23:16:27   |
| MIRIAM     | MCKINNEY  | 23:07:08   |
| ELMER      | NOE       | 22:55:13   |
| JEFFERY    | PINSON    | 22:53:33   |
| MINNIE     | ROMERO    | 23:00:34   |
| TERRANCE   | ROUSH     | 23:12:46   |
+-----+-----+-----+

```

16 rows in set (0.01 sec)

The order in which columns appear in your `order by` clause does make a difference when you include more than one column. If you were to switch the order of the two columns in the `order by` clause, Amber Dixon would appear first in the result set.

Ascending Versus Descending Sort Order

When sorting, you have the option of specifying *ascending* or *descending* order via the `asc` and `desc` keywords. The default is ascending, so you will need to add the `desc` keyword if you want to use a descending sort. For example, the following query shows all customers who rented films on June 14, 2005, in descending order of rental time:

```

mysql> SELECT c.first_name, c.last_name,
->    time(r.rental_date) rental_time
->   FROM customer c
->   INNER JOIN rental r
->   ON c.customer_id = r.customer_id
->   WHERE date(r.rental_date) = '2005-06-14'
->   ORDER BY time(r.rental_date) desc;
+-----+-----+-----+
| first_name | last_name | rental_time |
+-----+-----+-----+
| JEANETTE  | GREENE    | 23:54:46   |
| CHARLES   | KOWALSKI  | 23:54:34   |
| SONIA     | GREGORY   | 23:50:11   |
| TERRENCE   | GUNDERSON| 23:47:35   |
| AMBER     | DIXON     | 23:42:56   |
| HERMAN    | DEVORE    | 23:35:09   |
| MATTHEW   | MAHAN     | 23:25:58   |
+-----+-----+-----+

```

```

| CATHERINE | CAMPBELL | 23:17:03 |
| GWENDOLYN | MAY       | 23:16:27 |
| JOYCE     | EDWARDS   | 23:16:26 |
| TERRANCE  | ROUSH     | 23:12:46 |
| DANIEL    | CABRAL    | 23:09:38 |
| MIRIAM    | MCKINNEY  | 23:07:08 |
| MINNIE    | ROMERO    | 23:00:34 |
| ELMER     | NOE       | 22:55:13 |
| JEFFERY   | PINSON    | 22:53:33 |
+-----+-----+-----+
16 rows in set (0.01 sec)

```

Descending sorts are commonly used for ranking queries, such as “show me the top five account balances.” MySQL includes a `limit` clause that allows you to sort your data and then discard all but the first X rows.

Sorting via Numeric Placeholders

If you are sorting using the columns in your `select` clause, you can opt to reference the columns by their *position* in the `select` clause rather than by name. This can be especially helpful if you are sorting on an expression, such as in the previous example. Here’s the previous example one last time, with an `order by` clause specifying a descending sort using the third element in the `select` clause:

```

mysql> SELECT c.first_name, c.last_name,
->      time(r.rental_date) rental_time
->   FROM customer c
->   INNER JOIN rental r
->   ON c.customer_id = r.customer_id
-> WHERE date(r.rental_date) = '2005-06-14'
-> ORDER BY 3 desc;
+-----+-----+-----+
| first_name | last_name | rental_time |
+-----+-----+-----+
| JEANETTE  | GREENE    | 23:54:46 |
| CHARLES   | KOWALSKI  | 23:54:34 |
| SONIA     | GREGORY   | 23:50:11 |
| TERENCE   | GUNDERSON | 23:47:35 |
| AMBER     | DIXON     | 23:42:56 |
| HERMAN    | DEVORE    | 23:35:09 |
| MATTHEW   | MAHAN     | 23:25:58 |
| CATHERINE | CAMPBELL  | 23:17:03 |
| GWENDOLYN | MAY       | 23:16:27 |
| JOYCE     | EDWARDS   | 23:16:26 |
| TERRANCE  | ROUSH     | 23:12:46 |
| DANIEL    | CABRAL    | 23:09:38 |
| MIRIAM    | MCKINNEY  | 23:07:08 |
| MINNIE    | ROMERO    | 23:00:34 |
| ELMER     | NOE       | 22:55:13 |
| JEFFERY   | PINSON    | 22:53:33 |

```

```
+-----+-----+
16 rows in set (0.01 sec)
```

You might want to use this feature sparingly, since adding a column to the `select` clause without changing the numbers in the `order by` clause can lead to unexpected results. Personally, I may reference columns positionally when writing ad hoc queries, but I always reference columns by name when writing code.

Test Your Knowledge

The following exercises are designed to strengthen your understanding of the `select` statement and its various clauses. Please see [Appendix B](#) for solutions.

Exercise 3-1

Retrieve the actor ID, first name, and last name for all actors. Sort by last name and then by first name.

Exercise 3-2

Retrieve the actor ID, first name, and last name for all actors whose last name equals '`'WILLIAMS'`' or '`'DAVIS'`'.

Exercise 3-3

Write a query against the `rental` table that returns the IDs of the customers who rented a film on July 5, 2005 (use the `rental.rental_date` column, and you can use the `date()` function to ignore the time component). Include a single row for each distinct customer ID.

Exercise 3-4

Fill in the blanks (denoted by `<#>`) for this multitable query to achieve the following results:

```
mysql> SELECT c.email, r.return_date
-> FROM customer c
->   INNER JOIN rental <1>
->     ON c.customer_id = <2>
-> WHERE date(r.rental_date) = '2005-06-14'
-> ORDER BY <3> <4>;
+-----+-----+
| email                         | return_date      |
+-----+-----+
| DANIEL.CABRAL@sakilacustomer.org | 2005-06-23 22:00:38 |
| TERRANCE.ROUSH@sakilacustomer.org | 2005-06-23 21:53:46 |
| MIRIAM.MCKINNEY@sakilacustomer.org | 2005-06-21 17:12:08 |
```

GWENDOLYN.MAY@sakilacustomer.org	2005-06-20 02:40:27
JEANETTE.GREENE@sakilacustomer.org	2005-06-19 23:26:46
HERMAN.DEVORE@sakilacustomer.org	2005-06-19 03:20:09
JEFFERY.PINSON@sakilacustomer.org	2005-06-18 21:37:33
MATTHEW.MAHAN@sakilacustomer.org	2005-06-18 05:18:58
MINNIE.ROMERO@sakilacustomer.org	2005-06-18 01:58:34
SONIA.GREGORY@sakilacustomer.org	2005-06-17 21:44:11
TERENCE.GUNDERSON@sakilacustomer.org	2005-06-17 05:28:35
ELMER.NOE@sakilacustomer.org	2005-06-17 02:11:13
JOYCE.EDWARDS@sakilacustomer.org	2005-06-16 21:00:26
AMBER.DIXON@sakilacustomer.org	2005-06-16 04:02:56
CHARLES.KOWALSKI@sakilacustomer.org	2005-06-16 02:26:34
CATHERINE.CAMPBELL@sakilacustomer.org	2005-06-15 20:43:03

+-----+-----+

16 rows in set (0.03 sec)

Sometimes you will want to work with every row in a table, such as:

- Purging all data from a table used to stage new data warehouse feeds
- Modifying all rows in a table after a new column has been added
- Retrieving all rows from a message queue table

In cases like these, your SQL statements won't need to have a `where` clause, since you don't need to exclude any rows from consideration. Most of the time, however, you will want to narrow your focus to a subset of a table's rows. Therefore, all the SQL data statements (except the `insert` statement) include an optional `where` clause containing one or more *filter conditions* used to restrict the number of rows acted on by the SQL statement. Additionally, the `select` statement includes a `having` clause in which filter conditions pertaining to grouped data may be included. This chapter explores the various types of filter conditions that you can employ in the `where` clauses of `select`, `update`, and `delete` statements; I demonstrate the use of filter conditions in the `having` clause of a `select` statement in [Chapter 8](#).

Condition Evaluation

A `where` clause may contain one or more *conditions*, separated by the operators `and` and `or`. If multiple conditions are separated only by the `and` operator, then all the conditions must evaluate to `true` for the row to be included in the result set. Consider the following `where` clause:

```
WHERE first_name = 'STEVEN' AND create_date > '2006-01-01'
```

Given these two conditions, only rows where the first name is Steven and the creation date was after January 1, 2006, will be included in the result set. Though this example

uses only two conditions, no matter how many conditions are in your `where` clause, if they are separated by the `and` operator, they must *all* evaluate to `true` for the row to be included in the result set.

If all conditions in the `where` clause are separated by the `or` operator, however, only *one* of the conditions must evaluate to `true` for the row to be included in the result set. Consider the following two conditions:

```
WHERE first_name = 'STEVEN' OR create_date > '2006-01-01'
```

There are now various ways for a given row to be included in the result set:

- The first name is Steven, and the creation date was after January 1, 2006.
- The first name is Steven, and the creation date was on or before January 1, 2006.
- The first name is anything other than Steven, but the creation date was after January 1, 2006.

Table 4-1 shows the possible outcomes for a `where` clause containing two conditions separated by the `or` operator.

Table 4-1. Two-condition evaluation using or

Intermediate result	Final result
WHERE true OR true	true
WHERE true OR false	true
WHERE false OR true	true
WHERE false OR false	false

In the case of the preceding example, the only way for a row to be excluded from the result set is if the person's first name was not Steven and the creation date was on or before January 1, 2006.

Using Parentheses

If your `where` clause includes three or more conditions using both the `and` and `or` operators, you should use parentheses to make your intent clear, both to the database server and to anyone else reading your code. Here's a `where` clause that extends the previous example by checking to make sure that the first name is Steven or the last name is Young, and the creation date is after January 1, 2006:

```
WHERE (first_name = 'STEVEN' OR last_name = 'YOUNG')
      AND create_date > '2006-01-01'
```

There are now three conditions; for a row to make it to the final result set, either the first *or* second condition (or both) must evaluate to `true`, and the third condition must evaluate to `true`. **Table 4-2** shows the possible outcomes for this `where` clause.

Table 4-2. Three-condition evaluation using and, or

Intermediate result	Final result
WHERE (true OR true) AND true	true
WHERE (true OR false) AND true	true
WHERE (false OR true) AND true	true
WHERE (false OR false) AND true	false
WHERE (true OR true) AND false	false
WHERE (true OR false) AND false	false
WHERE (false OR true) AND false	false
WHERE (false OR false) AND false	false

As you can see, the more conditions you have in your `where` clause, the more combinations there are for the server to evaluate. In this case, only three of the eight combinations yield a final result of `true`.

Using the `not` Operator

Hopefully, the previous three-condition example is fairly easy to understand. Consider the following modification, however:

```
WHERE NOT (first_name = 'STEVEN' OR last_name = 'YOUNG')  
      AND create_date > '2006-01-01'
```

Did you spot the change from the previous example? I added the `not` operator before the first set of conditions. Now, instead of looking for people with the first name of Steven or the last name of Young whose record was created after January 1, 2006, I am retrieving only rows where the first name is not Steven or the last name is not Young whose record was created after January 1, 2006. **Table 4-3** shows the possible outcomes for this example.

Table 4-3. Three-condition evaluation using and, or, and not

Intermediate result	Final result
WHERE NOT (true OR true) AND true	false
WHERE NOT (true OR false) AND true	false
WHERE NOT (false OR true) AND true	false
WHERE NOT (false OR false) AND true	true
WHERE NOT (true OR true) AND false	false

Intermediate result	Final result
WHERE NOT (true OR false) AND false	false
WHERE NOT (false OR true) AND false	false
WHERE NOT (false OR false) AND false	false

While it is easy for the database server to handle, it is typically difficult for a person to evaluate a `where` clause that includes the `not` operator, which is why you won't encounter it very often. In this case, you can rewrite the `where` clause to avoid using the `not` operator:

```
WHERE first_name <> 'STEVEN' AND last_name <> 'YOUNG'  
      AND create_date > '2006-01-01'
```

While I'm sure that the server doesn't have a preference, you probably have an easier time understanding this version of the `where` clause.

Building a Condition

Now that you have seen how the server evaluates multiple conditions, let's take a step back and look at what comprises a single condition. A condition is made up of one or more *expressions* combined with one or more *operators*. An expression can be any of the following:

- A number
- A column in a table or view
- A string literal, such as 'Maple Street'
- A built-in function, such as `concat('Learning', ' ', 'SQL')`
- A subquery
- A list of expressions, such as ('Boston', 'New York', 'Chicago')

The operators used within conditions include:

- Comparison operators, such as `=`, `!=`, `<`, `>`, `<>`, `like`, `in`, and `between`
- Arithmetic operators, such as `+`, `-`, `*`, and `/`

The following section demonstrates how you can combine these expressions and operators to manufacture the various types of conditions.

Condition Types

There are many different ways to filter out unwanted data. You can look for specific values, sets of values, or ranges of values to include or exclude, or you can use various pattern-searching techniques to look for partial matches when dealing with string data. The next four subsections explore each of these condition types in detail.

Equality Conditions

A large percentage of the filter conditions that you write or come across will be of the form '`column = expression`' as in:

```
title = 'RIVER OUTLAW'  
fed_id = '111-11-1111'  
amount = 375.25  
film_id = (SELECT film_id FROM film WHERE title = 'RIVER OUTLAW')
```

Conditions such as these are called *equality conditions* because they equate one expression to another. The first three examples equate a column to a literal (two strings and a number), and the fourth example equates a column to the value returned from a subquery. The following query uses two equality conditions, one in the `on` clause (a join condition) and the other in the `where` clause (a filter condition):

```
mysql> SELECT c.email  
      -> FROM customer c  
      -> INNER JOIN rental r  
      -> ON c.customer_id = r.customer_id  
      -> WHERE date(r.rental_date) = '2005-06-14';  
+-----+  
| email |  
+-----+  
| CATHERINE.CAMPBELL@sakilacustomer.org |  
| JOYCE.EDWARDS@sakilacustomer.org |  
| AMBER.DIXON@sakilacustomer.org |  
| JEANETTE.GREENE@sakilacustomer.org |  
| MINNIE.ROMERO@sakilacustomer.org |  
| GWENDOLYN.MAY@sakilacustomer.org |  
| SONIA.GREGORY@sakilacustomer.org |  
| MIRIAM.MCKINNEY@sakilacustomer.org |  
| CHARLES.KOWALSKI@sakilacustomer.org |  
| DANIEL.CABRAL@sakilacustomer.org |  
| MATTHEW.MAHAN@sakilacustomer.org |  
| JEFFERY.PINSON@sakilacustomer.org |  
| HERMAN.DEVORE@sakilacustomer.org |  
| ELMER.NOE@sakilacustomer.org |  
| TERRANCE.ROUSH@sakilacustomer.org |  
| TERRENCE.GUNDERSON@sakilacustomer.org |  
+-----+  
16 rows in set (0.03 sec)
```

This query shows all email addresses of every customer who rented a film on June 14, 2005.

Inequality conditions

Another fairly common type of condition is the *inequality condition*, which asserts that two expressions are *not* equal. Here's the previous query with the filter condition in the where clause changed to an inequality condition:

```
mysql> SELECT c.email
->   FROM customer c
->     INNER JOIN rental r
->       ON c.customer_id = r.customer_id
-> WHERE date(r.rental_date) <> '2005-06-14';

+-----+
| email
+-----+
| MARY.SMITH@sakilacustomer.org |
...
| AUSTIN.CINTRON@sakilacustomer.org |
+-----+
16028 rows in set (0.03 sec)
```

This query returns all email addresses for films rented on any other date than June 14, 2005. When building inequality conditions, you may choose to use either the != or <> operator.

Data modification using equality conditions

Equality/inequality conditions are commonly used when modifying data. For example, let's say that the movie rental company has a policy of removing old account rows

once per year. Your task is to remove rows from the `rental` table where the rental date was in 2004. Here's one way to tackle it:

```
DELETE FROM rental
WHERE year(rental_date) = 2004;
```

This statement includes a single equality condition; here's an example that uses two inequality conditions to remove any rows where the rental date was not in 2005 or 2006:

```
DELETE FROM rental
WHERE year(rental_date) <> 2005 AND year(rental_date) <> 2006;
```



When crafting examples of delete and update statements, I try to write each statement such that no rows are modified. That way, when you execute the statements, your data will remain unchanged, and your output from select statements will always match that shown in this book.

Since MySQL sessions are in auto-commit mode by default (see [Chapter 12](#)), you would not be able to roll back (undo) any changes made to the example data if one of my statements modified the data. You may, of course, do whatever you want with the example data, including wiping it clean and rerunning the scripts to populate the tables, but I try to leave it intact.

Range Conditions

Along with checking that an expression is equal to (or not equal to) another expression, you can build conditions that check whether an expression falls within a certain range. This type of condition is common when working with numeric or temporal data. Consider the following query:

```
mysql> SELECT customer_id, rental_date
->   FROM rental
-> WHERE rental_date < '2005-05-25';
+-----+-----+
| customer_id | rental_date      |
+-----+-----+
|      130 | 2005-05-24 22:53:30 |
|      459 | 2005-05-24 22:54:33 |
|      408 | 2005-05-24 23:03:39 |
|      333 | 2005-05-24 23:04:41 |
|      222 | 2005-05-24 23:05:21 |
|      549 | 2005-05-24 23:08:07 |
|      269 | 2005-05-24 23:11:53 |
|      239 | 2005-05-24 23:31:46 |
+-----+-----+
8 rows in set (0.00 sec)
```

This query finds all film rentals prior to May 25, 2005. As well as specifying an upper limit for the rental date, you may also want to specify a lower range:

```
mysql> SELECT customer_id, rental_date
-> FROM rental
-> WHERE rental_date <= '2005-06-16'
-> AND rental_date >= '2005-06-14';
+-----+-----+
| customer_id | rental_date      |
+-----+-----+
|        416 | 2005-06-14 22:53:33 |
|        516 | 2005-06-14 22:55:13 |
|        239 | 2005-06-14 23:00:34 |
|        285 | 2005-06-14 23:07:08 |
|        310 | 2005-06-14 23:09:38 |
|        592 | 2005-06-14 23:12:46 |
...
|       148 | 2005-06-15 23:20:26 |
|       237 | 2005-06-15 23:36:37 |
|       155 | 2005-06-15 23:55:27 |
|       341 | 2005-06-15 23:57:20 |
|       149 | 2005-06-15 23:58:53 |
+-----+-----+
364 rows in set (0.00 sec)
```

This version of the query retrieves all films rented on June 14 or 15 of 2005.

The between operator

When you have *both* an upper and lower limit for your range, you may choose to use a single condition that utilizes the `between` operator rather than using two separate conditions, as in:

```
mysql> SELECT customer_id, rental_date
-> FROM rental
-> WHERE rental_date BETWEEN '2005-06-14' AND '2005-06-16';
+-----+-----+
| customer_id | rental_date      |
+-----+-----+
|        416 | 2005-06-14 22:53:33 |
|        516 | 2005-06-14 22:55:13 |
|        239 | 2005-06-14 23:00:34 |
|        285 | 2005-06-14 23:07:08 |
|        310 | 2005-06-14 23:09:38 |
|        592 | 2005-06-14 23:12:46 |
...
|       148 | 2005-06-15 23:20:26 |
|       237 | 2005-06-15 23:36:37 |
|       155 | 2005-06-15 23:55:27 |
|       341 | 2005-06-15 23:57:20 |
|       149 | 2005-06-15 23:58:53 |
+-----+-----+
```

```
+-----+-----+
364 rows in set (0.00 sec)
```

When using the `between` operator, there are a couple of things to keep in mind. You should always specify the lower limit of the range first (after `between`) and the upper limit of the range second (after `and`). Here's what happens if you mistakenly specify the upper limit first:

```
mysql> SELECT customer_id, rental_date
-> FROM rental
-> WHERE rental_date BETWEEN '2005-06-16' AND '2005-06-14';
Empty set (0.00 sec)
```

As you can see, no data is returned. This is because the server is, in effect, generating two conditions from your single condition using the `<=` and `>=` operators, as in:

```
SELECT customer_id, rental_date
-> FROM rental
-> WHERE rental_date >= '2005-06-16'
-> AND rental_date <= '2005-06-14'
Empty set (0.00 sec)
```

Since it is impossible to have a date that is *both* greater than June 16, 2005, and less than June 14, 2005, the query returns an empty set. This brings me to the second pitfall when using `between`, which is to remember that your upper and lower limits are *inclusive*, meaning that the values you provide are included in the range limits. In this case, I want to return any films rented on June 14 or 15, so I specify `2005-06-14` as the lower end of the range and `2005-06-16` as the upper end. Since I am not specifying the time component of the date, the time defaults to midnight, so the effective range is `2005-06-14 00:00:00` to `2005-06-16 00:00:00`, which will include any rentals made on June 14 or 15.

Along with dates, you can also build conditions to specify ranges of numbers. Numeric ranges are fairly easy to grasp, as demonstrated by the following:

```
mysql> SELECT customer_id, payment_date, amount
-> FROM payment
-> WHERE amount BETWEEN 10.0 AND 11.99;
+-----+-----+-----+
| customer_id | payment_date       | amount |
+-----+-----+-----+
|      2 | 2005-07-30 13:47:43 | 10.99 |
|      3 | 2005-07-27 20:23:12 | 10.99 |
|     12 | 2005-08-01 06:50:26 | 10.99 |
|     13 | 2005-07-29 22:37:41 | 11.99 |
|     21 | 2005-06-21 01:04:35 | 10.99 |
|     29 | 2005-07-09 21:55:19 | 10.99 |
...
|    571 | 2005-06-20 08:15:27 | 10.99 |
|    572 | 2005-06-17 04:05:12 | 10.99 |
|    573 | 2005-07-31 12:14:19 | 10.99 |
```

```

|      591 | 2005-07-07 20:45:51 | 11.99 |
|      592 | 2005-07-06 22:58:31 | 11.99 |
|      595 | 2005-07-31 11:51:46 | 10.99 |
+-----+-----+-----+
114 rows in set (0.01 sec)

```

All payments between \$10 and \$11.99 are returned. Again, make sure that you specify the lower amount first.

String ranges

While ranges of dates and numbers are easy to understand, you can also build conditions that search for ranges of strings, which are a bit harder to visualize. Say, for example, you are searching for customers whose last name falls within a range. Here's a query that returns customers whose last name falls between FA and FR:

```

mysql> SELECT last_name, first_name
    -> FROM customer
    -> WHERE last_name BETWEEN 'FA' AND 'FR';
+-----+-----+
| last_name | first_name |
+-----+-----+
| FARNSWORTH | JOHN |
| FENNELL | ALEXANDER |
| FERGUSON | BERTHA |
| FERNANDEZ | MELINDA |
| FIELDS | VICKI |
| FISHER | CINDY |
| FLEMING | MYRTLE |
| FLETCHER | MAE |
| FLORES | JULIA |
| FORD | CRYSTAL |
| FORMAN | MICHEAL |
| FORSYTHE | ENRIQUE |
| FORTIER | RAUL |
| FORTNER | HOWARD |
| FOSTER | PHYLLIS |
| FOUST | JACK |
| FOWLER | JO |
| FOX | HOLLY |
+-----+-----+
18 rows in set (0.00 sec)

```

While there are five customers whose last name starts with FR, they are not included in the results, since a name like FRANKLIN is outside of the range. However, we can pick up four of the five customers by extending the righthand range to FRB:

```

mysql> SELECT last_name, first_name
    -> FROM customer
    -> WHERE last_name BETWEEN 'FA' AND 'FRB';
+-----+-----+
| last_name | first_name |
+-----+-----+

```

```

+-----+-----+
| FARNSWORTH | JOHN      |
| FENNELL    | ALEXANDER |
| FERGUSON    | BERTHA    |
| FERNANDEZ   | MELINDA   |
| FIELDS      | VICKI     |
| FISHER      | CINDY     |
| FLEMING     | MYRTLE   |
| FLETCHER    | MAE       |
| FLORES      | JULIA     |
| FORD        | CRYSTAL   |
| FORMAN      | MICHEAL   |
| FORSYTHE   | ENRIQUE   |
| FORTIER     | RAUL      |
| FORTNER     | HOWARD    |
| FOSTER      | PHYLLIS   |
| FOUST       | JACK      |
| FOWLER      | JO        |
| FOX         | HOLLY     |
| FRALEY      | JUAN      |
| FRANCISCO  | JOEL      |
| FRANKLIN   | BETH      |
| FRAZIER     | GLENDA   |
+-----+-----+
22 rows in set (0.00 sec)

```

To work with string ranges, you need to know the order of the characters within your character set (the order in which the characters within a character set are sorted is called a *collation*).

Membership Conditions

In some cases, you will not be restricting an expression to a single value or range of values but rather to a finite set of values. For example, you might want to locate all films that have a rating of either 'G' or 'PG':

```

mysql> SELECT title, rating
    -> FROM film
    -> WHERE rating = 'G' OR rating = 'PG';
+-----+-----+
| title           | rating |
+-----+-----+
| ACADEMY DINOSAUR | PG     |
| ACE GOLDFINGER  | G      |
| AFFAIR PREJUDICE | G      |
| AFRICAN EGG     | G      |
| AGENT TRUMAN    | PG     |
| ALAMO VIDEOTAPE | G      |
| ALASKA PHANTOM   | PG     |
| ALI FOREVER      | PG     |
| AMADEUS HOLY     | PG     |
+-----+-----+

```

```

...
| WEDDING APOLLO      | PG    |
| WEREWOLF LOLA       | G     |
| WEST LION            | G     |
| WIZARD COLDBLOODED  | PG    |
| WON DARES           | PG    |
| WONDERLAND CHRISTMAS| PG    |
| WORDS HUNTER        | PG    |
| WORST BANGER         | PG    |
| YOUNG LANGUAGE       | G     |
+-----+-----+
372 rows in set (0.00 sec)

```

While this `where` clause (two conditions `or'd` together) wasn't too tedious to generate, imagine if the set of expressions contained 10 or 20 members. For these situations, you can use the `in` operator instead:

```

SELECT title, rating
FROM film
WHERE rating IN ('G', 'PG');

```

With the `in` operator, you can write a single condition no matter how many expressions are in the set.

Using subqueries

Along with writing your own set of expressions, such as `('G', 'PG')`, you can use a subquery to generate a set for you on the fly. For example, if you can assume that any film whose title includes the string `'PET'` would be safe for family viewing, you could execute a subquery against the `film` table to retrieve all ratings associated with these films and then retrieve all films having any of these ratings:

```

mysql> SELECT title, rating
-> FROM film
-> WHERE rating IN (SELECT rating FROM film WHERE title LIKE '%PET%');
+-----+-----+
| title          | rating |
+-----+-----+
| ACADEMY DINOSAUR | PG    |
| ACE GOLDFINGER   | G     |
| AFFAIR PREJUDICE | G     |
| AFRICAN EGG      | G     |
| AGENT TRUMAN     | PG    |
| ALAMO VIDEOTAPE  | G     |
| ALASKA PHANTOM    | PG    |
| ALI FOREVER       | PG    |
| AMADEUS HOLY      | PG    |
...
| WEDDING APOLLO      | PG    |
| WEREWOLF LOLA       | G     |
| WEST LION            | G     |

```

```

| WIZARD COLDBLOODED      | PG    |
| WON DARES               | PG    |
| WONDERLAND CHRISTMAS    | PG    |
| WORDS HUNTER            | PG    |
| WORST BANGER             | PG    |
| YOUNG LANGUAGE           | G     |
+-----+-----+
372 rows in set (0.00 sec)

```

The subquery returns the set 'G' and 'PG', and the main query checks to see whether the value of the `rating` column can be found in the set returned by the subquery.

Using `not in`

Sometimes you want to see whether a particular expression exists within a set of expressions, and sometimes you want to see whether the expression does *not* exist within the set. For these situations, you can use the `not in` operator:

```

SELECT title, rating
  FROM film
 WHERE rating NOT IN ('PG-13','R', 'NC-17');

```

This query finds all accounts that are *not* rated 'PG-13', 'R', or 'NC-17', which will return the same set of 372 rows as the previous queries.

Matching Conditions

So far, you have been introduced to conditions that identify an exact string, a range of strings, or a set of strings; the final condition type deals with partial string matches. You may, for example, want to find all customers whose last name begins with Q. You could use a built-in function to strip off the first letter of the `last_name` column, as in the following:

```

mysql> SELECT last_name, first_name
->   FROM customer
->  WHERE left(last_name, 1) = 'Q';
+-----+-----+
| last_name | first_name |
+-----+-----+
| QUALLS    | STEPHEN    |
| QUINTANILLA | ROGER    |
| QUIGLEY   | TROY      |
+-----+-----+
3 rows in set (0.00 sec)

```

While the built-in function `left()` does the job, it doesn't give you much flexibility. Instead, you can use wildcard characters to build search expressions, as demonstrated in the next section.

Using wildcards

When searching for partial string matches, you might be interested in:

- Strings beginning/ending with a certain character
- Strings beginning/ending with a substring
- Strings containing a certain character anywhere within the string
- Strings containing a substring anywhere within the string
- Strings with a specific format, regardless of individual characters

You can build search expressions to identify these and many other partial string matches by using the wildcard characters shown in [Table 4-4](#).

Table 4-4. Wildcard characters

Wildcard character	Matches
_	Exactly one character
%	Any number of characters (including 0)

The underscore character takes the place of a single character, while the percent sign can take the place of a variable number of characters. When building conditions that utilize search expressions, you use the `like` operator, as in:

```
mysql> SELECT last_name, first_name
    -> FROM customer
    -> WHERE last_name LIKE '_A_T%';
+-----+-----+
| last_name | first_name |
+-----+-----+
| MATTHEWS | ERICA      |
| WALTERS  | CASSANDRA  |
| WATTS    | SHELLY     |
+-----+-----+
3 rows in set (0.00 sec)
```

The search expression in the previous example specifies strings containing an *A* in the second position and a *T* in the fourth position, followed by any number of characters and ending in *S*. [Table 4-5](#) shows some more search expressions and their interpretations.

Table 4-5. Sample search expressions

Search expression	Interpretation
F%	Strings beginning with <i>F</i>
%t	Strings ending with <i>t</i>
%bas%	Strings containing the substring 'bas'

Search expression	Interpretation
_ _t_	Four-character strings with a <i>t</i> in the third position
_ _ - _ - _ - _	11-character strings with dashes in the fourth and seventh positions

The wildcard characters work fine for building simple search expressions; if your needs are a bit more sophisticated, however, you can use multiple search expressions, as demonstrated by the following:

```
mysql> SELECT last_name, first_name
-> FROM customer
-> WHERE last_name LIKE 'Q%' OR last_name LIKE 'Y%';
+-----+-----+
| last_name | first_name |
+-----+-----+
| QUALLS   | STEPHEN    |
| QUIGLEY  | TROY        |
| QUINTANILLA | ROGER      |
| YANEZ    | LUIS         |
| YEE      | MARVIN      |
| YOUNG    | CYNTHIA    |
+-----+-----+
6 rows in set (0.00 sec)
```

This query finds all customers whose last name begins with Q or Y.

Using regular expressions

If you find that the wildcard characters don't provide enough flexibility, you can use regular expressions to build search expressions. A regular expression is, in essence, a search expression on steroids. If you are new to SQL but have coded using programming languages such as Perl, then you might already be intimately familiar with regular expressions. If you have never used regular expressions, then you may want to consult Jeffrey E. F. Friedl's *Mastering Regular Expressions* (O'Reilly), since it is far too large a topic to try to cover in this book.

Here's what the previous query (find all customers whose last name starts with Q or Y) would look like using the MySQL implementation of regular expressions:

```
mysql> SELECT last_name, first_name
-> FROM customer
-> WHERE last_name REGEXP '^[QY]';
+-----+-----+
| last_name | first_name |
+-----+-----+
| YOUNG    | CYNTHIA    |
| QUALLS   | STEPHEN    |
| QUINTANILLA | ROGER      |
| YANEZ    | LUIS         |
| YEE      | MARVIN      |
| QUIGLEY  | TROY        |
+-----+-----+
```

```
+-----+-----+
6 rows in set (0.16 sec)
```

The `regexp` operator takes a regular expression ('`^QY`' in this example) and applies it to the expression on the lefthand side of the condition (the column `last_name`). The query now contains a single condition using a regular expression rather than two conditions using wildcard characters.

Both Oracle Database and Microsoft SQL Server also support regular expressions. With Oracle Database, you would use the `regexp_like` function instead of the `regexp` operator shown in the previous example, whereas SQL Server allows regular expressions to be used with the `like` operator.

Null: That Four-Letter Word

I put it off as long as I could, but it's time to broach a topic that tends to be met with fear, uncertainty, and dread: the `null` value. `null` is the absence of a value; before an employee is terminated, for example, her `end_date` column in the `employee` table should be `null`. There is no value that can be assigned to the `end_date` column that would make sense in this situation. `null` is a bit slippery, however, as there are various flavors of `null`:

Not applicable

Such as the employee ID column for a transaction that took place at an ATM machine

Value not yet known

Such as when the federal ID is not known at the time a customer row is created

Value undefined

Such as when an account is created for a product that has not yet been added to the database



Some theorists argue that there should be a different expression to cover each of these (and more) situations, but most practitioners would agree that having multiple `null` values would be far too confusing.

When working with `null`, you should remember:

- An expression can *be* `null`, but it can never *equal* `null`.
- Two `null`s are never equal to each other.

To test whether an expression is `null`, you need to use the `is null` operator, as demonstrated by the following:

```
mysql> SELECT rental_id, customer_id
    -> FROM rental
    -> WHERE return_date IS NULL;
+-----+-----+
| rental_id | customer_id |
+-----+-----+
| 11496 | 155 |
| 11541 | 335 |
| 11563 | 83 |
| 11577 | 219 |
| 11593 | 99 |
...
| 15867 | 505 |
| 15875 | 41 |
| 15894 | 168 |
| 15966 | 374 |
+-----+
183 rows in set (0.01 sec)
```

This query finds all film rentals that were never returned. Here's the same query using `= null` instead of `is null`:

```
mysql> SELECT rental_id, customer_id
    -> FROM rental
    -> WHERE return_date = NULL;
Empty set (0.01 sec)
```

As you can see, the query parses and executes but does not return any rows. This is a common mistake made by inexperienced SQL programmers, and the database server will not alert you to your error, so be careful when constructing conditions that test for `null`.

If you want to see whether a value has been assigned to a column, you can use the `is not null` operator, as in:

```
mysql> SELECT rental_id, customer_id, return_date
    -> FROM rental
    -> WHERE return_date IS NOT NULL;
+-----+-----+-----+
| rental_id | customer_id | return_date |
+-----+-----+-----+
| 1 | 130 | 2005-05-26 22:04:30 |
| 2 | 459 | 2005-05-28 19:40:33 |
| 3 | 408 | 2005-06-01 22:12:39 |
| 4 | 333 | 2005-06-03 01:43:41 |
| 5 | 222 | 2005-06-02 04:33:21 |
| 6 | 549 | 2005-05-27 01:32:07 |
| 7 | 269 | 2005-05-29 20:34:53 |
...
...
```

```

|    16043 |      526 | 2005-08-31 03:09:03 |
|    16044 |      468 | 2005-08-25 04:08:39 |
|    16045 |       14 | 2005-08-25 23:54:26 |
|    16046 |       74 | 2005-08-27 18:02:47 |
|    16047 |      114 | 2005-08-25 02:48:48 |
|    16048 |      103 | 2005-08-31 21:33:07 |
|    16049 |      393 | 2005-08-30 01:01:12 |
+-----+
15861 rows in set (0.02 sec)

```

This version of the query returns all rentals that were returned, which is the majority of the rows in the table (15,861 out of 16,044).

Before putting `null` aside for a while, it would be helpful to investigate one more potential pitfall. Suppose that you have been asked to find all rentals that were not returned during May through August of 2005. Your first instinct might be to do the following:

```

mysql> SELECT rental_id, customer_id, return_date
-> FROM rental
-> WHERE return_date NOT BETWEEN '2005-05-01' AND '2005-09-01';
+-----+-----+-----+
| rental_id | customer_id | return_date        |
+-----+-----+-----+
|    15365 |      327 | 2005-09-01 03:14:17 |
|    15388 |       50 | 2005-09-01 03:50:23 |
|    15392 |      410 | 2005-09-01 01:14:15 |
|    15401 |      103 | 2005-09-01 03:44:10 |
|    15415 |      204 | 2005-09-01 02:05:56 |
...
|    15977 |      550 | 2005-09-01 22:12:10 |
|    15982 |      370 | 2005-09-01 21:51:31 |
|    16005 |      466 | 2005-09-02 02:35:22 |
|    16020 |      311 | 2005-09-01 18:17:33 |
|    16033 |      226 | 2005-09-01 02:36:15 |
|    16037 |       45 | 2005-09-01 02:48:04 |
|    16040 |      195 | 2005-09-02 02:19:33 |
+-----+
62 rows in set (0.01 sec)

```

While it is true that these 62 rentals were returned outside of the May to August window, if you look carefully at the data, you will see that all of the rows returned have a non-null return date. But what about the 183 rentals that were never returned? One might argue that these 183 rows were also not returned between May and August, so they should also be included in the result set. To answer the question correctly, therefore, you need to account for the possibility that some rows might contain a `null` in the `return_date` column:

```

mysql> SELECT rental_id, customer_id, return_date
-> FROM rental
-> WHERE return_date IS NULL

```

```

-> OR return_date NOT BETWEEN '2005-05-01' AND '2005-09-01';
+-----+-----+
| rental_id | customer_id | return_date      |
+-----+-----+
| 11496    |      155 | NULL           |
| 11541    |      335 | NULL           |
| 11563    |       83 | NULL           |
| 11577    |      219 | NULL           |
| 11593    |       99 | NULL           |
...
| 15939    |      382 | 2005-09-01 17:25:21 |
| 15942    |      210 | 2005-09-01 18:39:40 |
| 15966    |      374 | NULL           |
| 15971    |      187 | 2005-09-02 01:28:33 |
| 15973    |      343 | 2005-09-01 20:08:41 |
| 15977    |      550 | 2005-09-01 22:12:10 |
| 15982    |      370 | 2005-09-01 21:51:31 |
| 16005    |      466 | 2005-09-02 02:35:22 |
| 16020    |      311 | 2005-09-01 18:17:33 |
| 16033    |      226 | 2005-09-01 02:36:15 |
| 16037    |       45 | 2005-09-01 02:48:04 |
| 16040    |      195 | 2005-09-02 02:19:33 |
+-----+-----+
245 rows in set (0.01 sec)

```

The result set now includes the 62 rentals that were returned outside of the May to August window, along with the 183 rentals that were never returned, for a total of 245 rows. When working with a database that you are not familiar with, it is a good idea to find out which columns in a table allow `nulls` so that you can take appropriate measures with your filter conditions to keep data from slipping through the cracks.

Test Your Knowledge

The following exercises test your understanding of filter conditions. Please see [Appendix B](#) for solutions.

You'll need to refer to the following subset of rows from the `payment` table for the first two exercises:

```

+-----+-----+-----+
| payment_id | customer_id | amount | date(payment_date) |
+-----+-----+-----+
| 101 |      4 | 8.99 | 2005-08-18 |
| 102 |      4 | 1.99 | 2005-08-19 |
| 103 |      4 | 2.99 | 2005-08-20 |
| 104 |      4 | 6.99 | 2005-08-20 |
| 105 |      4 | 4.99 | 2005-08-21 |
| 106 |      4 | 2.99 | 2005-08-22 |
| 107 |      4 | 1.99 | 2005-08-23 |
| 108 |      5 | 0.99 | 2005-05-29 |
| 109 |      5 | 6.99 | 2005-05-31 |

```

110	5	1.99	2005-05-31
111	5	3.99	2005-06-15
112	5	2.99	2005-06-16
113	5	4.99	2005-06-17
114	5	2.99	2005-06-19
115	5	4.99	2005-06-20
116	5	4.99	2005-07-06
117	5	2.99	2005-07-08
118	5	4.99	2005-07-09
119	5	5.99	2005-07-09
120	5	1.99	2005-07-09

Exercise 4-1

Which of the payment IDs would be returned by the following filter conditions?

```
customer_id <> 5 AND (amount > 8 OR date(payment_date) = '2005-08-23')
```

Exercise 4-2

Which of the payment IDs would be returned by the following filter conditions?

```
customer_id = 5 AND NOT (amount > 6 OR date(payment_date) = '2005-06-19')
```

Exercise 4-3

Construct a query that retrieves all rows from the payments table where the amount is either 1.98, 7.98, or 9.98.

Exercise 4-4

Construct a query that finds all customers whose last name contains an A in the second position and a W anywhere after the A.

Querying Multiple Tables

Back in [Chapter 2](#), I demonstrated how related concepts are broken into separate pieces through a process known as normalization. The end result of this exercise was two tables: `person` and `favorite_food`. If, however, you want to generate a single report showing a person's name, address, *and* favorite foods, you will need a mechanism to bring the data from these two tables back together again; this mechanism is known as a *join*, and this chapter concentrates on the simplest and most common join, the *inner join*. [Chapter 10](#) demonstrates all of the different join types.

What Is a Join?

Queries against a single table are certainly not rare, but you will find that most of your queries will require two, three, or even more tables. To illustrate, let's look at the definitions for the `customer` and `address` tables and then define a query that retrieves data from both tables:

```
mysql> desc customer;
+-----+-----+-----+-----+
| Field | Type | Null | Key | Default |
+-----+-----+-----+-----+
| customer_id | smallint(5) unsigned | NO | PRI | NULL |
| store_id | tinyint(3) unsigned | NO | MUL | NULL |
| first_name | varchar(45) | NO |  | NULL |
| last_name | varchar(45) | NO | MUL | NULL |
| email | varchar(50) | YES |  | NULL |
| address_id | smallint(5) unsigned | NO | MUL | NULL |
| active | tinyint(1) | NO |  | 1 |
| create_date | datetime | NO |  | NULL |
| last_update | timestamp | YES |  | CURRENT_TIMESTAMP |
+-----+-----+-----+-----+
mysql> desc address;
```

Field	Type	Null	Key	Default
address_id	smallint(5) unsigned	NO	PRI	NULL
address	varchar(50)	NO		NULL
address2	varchar(50)	YES		NULL
district	varchar(20)	NO		NULL
city_id	smallint(5) unsigned	NO	MUL	NULL
postal_code	varchar(10)	YES		NULL
phone	varchar(20)	NO		NULL
location	geometry	NO	MUL	NULL
last_update	timestamp	NO		CURRENT_TIMESTAMP

Let's say you want to retrieve the first and last names of each customer, along with their street address. Your query will therefore need to retrieve the `customer.first_name`, `customer.last_name`, and `address.address` columns. But how can you retrieve data from both tables in the same query? The answer lies in the `customer.address_id` column, which holds the ID of the customer's record in the `address` table (in more formal terms, the `customer.address_id` column is the *foreign key* to the `address` table). The query, which you will see shortly, instructs the server to use the `customer.address_id` column as the *transportation* between the `customer` and `address` tables, thereby allowing columns from both tables to be included in the query's result set. This type of operation is known as a *join*.



A foreign key constraint can optionally be created to verify that the values in one table exist in another table. For the previous example, a foreign key constraint could be created on the `customer` table to ensure that any values inserted into the `customer.address_id` column can be found in the `address.address_id` column. Please note that it is not necessary to have a foreign key constraint in place in order to join two tables.

Cartesian Product

The easiest way to start is to put the `customer` and `address` tables into the `from` clause of a query and see what happens. Here's a query that retrieves the customer's first and last names along with the street address, with a `from` clause naming both tables separated by the `join` keyword:

```
mysql> SELECT c.first_name, c.last_name, a.address
-> FROM customer c JOIN address a;
+-----+-----+
| first_name | last_name | address          |
+-----+-----+
| MARY      | SMITH     | 47 MySakila Drive |
| PATRICIA  | JOHNSON   | 47 MySakila Drive |
```

```

| LINDA      | WILLIAMS   | 47 MySakila Drive |
| BARBARA    | JONES       | 47 MySakila Drive |
| ELIZABETH  | BROWN       | 47 MySakila Drive |
| JENNIFER   | DAVIS        | 47 MySakila Drive |
| MARIA      | MILLER      | 47 MySakila Drive |
| SUSAN      | WILSON      | 47 MySakila Drive |
|
| ...         |             |                   |
| SETH        | HANNON      | 1325 Fukuyama Street |
| KENT        | ARSENAULT   | 1325 Fukuyama Street |
| TERRANCE   | ROUSH        | 1325 Fukuyama Street |
| RENE        | MCALISTER   | 1325 Fukuyama Street |
| EDUARDO    | HIATT        | 1325 Fukuyama Street |
| TERRENCE   | GUNDERSON  | 1325 Fukuyama Street |
| ENRIQUE    | FORSYTHE    | 1325 Fukuyama Street |
| FREDDIE    | DUGGAN       | 1325 Fukuyama Street |
| WADE        | DELVALLE    | 1325 Fukuyama Street |
| AUSTIN     | CINTRON     | 1325 Fukuyama Street |
+-----+-----+-----+

```

361197 rows in set (0.03 sec)

Hmmm...there are only 599 customers and 603 rows in the `address` table, so how did the result set end up with 361,197 rows? Looking more closely, you can see that many of the customers seem to have the same street address. Because the query didn't specify *how* the two tables should be joined, the database server generated the *Cartesian product*, which is *every* permutation of the two tables (599 customers x 603 addresses = 361,197 permutations). This type of join is known as a *cross join*, and it is rarely used (on purpose, at least). Cross joins are one of the join types that we study in [Chapter 10](#).

Inner Joins

To modify the previous query so that only a single row is returned for each customer, you need to describe how the two tables are related. Earlier, I showed that the `customer.address_id` column serves as the link between the two tables, so this information needs to be added to the `on` subclause of the `from` clause:

```

mysql> SELECT c.first_name, c.last_name, a.address
-> FROM customer c JOIN address a
-> ON c.address_id = a.address_id;
+-----+-----+-----+
| first_name | last_name | address          |
+-----+-----+-----+
| MARY       | SMITH     | 1913 Hanoi Way   |
| PATRICIA   | JOHNSON   | 1121 Loja Avenue |
| LINDA      | WILLIAMS  | 692 Joliet Street|
| BARBARA    | JONES     | 1566 Inegl Manor |
| ELIZABETH  | BROWN     | 53 Idfu Parkway  |
| JENNIFER   | DAVIS     | 1795 Santiago de Compostela Way |
| MARIA      | MILLER    | 900 Santiago de Compostela Parkway |
| SUSAN      | WILSON    | 478 Joliet Way   |

```

MARGARET	MOORE	613 Korolev Drive
...		
TERRANCE	ROUSH	42 Fontana Avenue
RENE	MCALISTER	1895 Zhezqazghan Drive
EDUARDO	HIATT	1837 Kaduna Parkway
TERRENCE	GUNDERSON	844 Bucuresti Place
ENRIQUE	FORSYTHE	1101 Bucuresti Boulevard
FREDDIE	DUGGAN	1103 Quilmes Boulevard
WADE	DELVALLE	1331 Usak Boulevard
AUSTIN	CINTRON	1325 Fukuyama Street

599 rows in set (0.00 sec)

Instead of 361,197 rows, you now have the expected 599 rows due to the addition of the `on` subclause, which instructs the server to join the `customer` and `address` tables by using the `address_id` column to traverse from one table to the other. For example, Mary Smith's row in the `customer` table contains a value of 5 in the `address_id` column (not shown in the example). The server uses this value to look up the row in the `address` table having a value of 5 in its `address_id` column and then retrieves the value '1913 Hanoi Way' from the `address` column in that row.

If a value exists for the `address_id` column in one table but *not* the other, then the join fails for the rows containing that value, and those rows are excluded from the result set. This type of join is known as an *inner join*, and it is the most commonly used type of join. To clarify, if a row in the `customer` table has the value 999 in the `address_id` column and there's no row in the `address` table with a value of 999 in the `address_id` column, then that customer row would not be included in the result set. If you want to include all rows from one table or the other regardless of whether a match exists, you need to specify an *outer join*, but this will be explored [Chapter 10](#).

In the previous example, I did not specify in the `from` clause which type of join to use. However, when you wish to join two tables using an inner join, you should explicitly specify this in your `from` clause; here's the same example, with the addition of the join type (note the keyword `inner`):

```
SELECT c.first_name, c.last_name, a.address
FROM customer c INNER JOIN address a
ON c.address_id = a.address_id;
```

If you do not specify the type of join, then the server will do an inner join by default. As you will see later in the book, however, there are several types of joins, so you should get in the habit of specifying the exact type of join that you require, especially for the benefit of any other people who might use/maintain your queries in the future.

If the names of the columns used to join the two tables are identical, which is true in the previous query, you can use the `using` subclause instead of the `on` subclause, as in:

```
SELECT c.first_name, c.last_name, a.address
FROM customer c INNER JOIN address a
    USING (address_id);
```

Since `using` is a shorthand notation that you can use in only a specific situation, I prefer always to use the `on` subclause to avoid confusion.

The ANSI Join Syntax

The notation used throughout this book for joining tables was introduced in the SQL92 version of the ANSI SQL standard. All the major databases (Oracle Database, Microsoft SQL Server, MySQL, IBM DB2 Universal Database, and Sybase Adaptive Server) have adopted the SQL92 join syntax. Because most of these servers have been around since before the release of the SQL92 specification, they all include an older join syntax as well. For example, all these servers would understand the following variation of the previous query:

```
mysql> SELECT c.first_name, c.last_name, a.address
-> FROM customer c, address a
-> WHERE c.address_id = a.address_id;
+-----+-----+-----+
| first_name | last_name | address
+-----+-----+-----+
| MARY      | SMITH     | 1913 Hanoi Way
| PATRICIA  | JOHNSON   | 1121 Loja Avenue
| LINDA     | WILLIAMS  | 692 Joliet Street
| BARBARA   | JONES     | 1566 Inegl Manor
| ELIZABETH | BROWN     | 53 Idfu Parkway
| JENNIFER  | DAVIS     | 1795 Santiago de Compostela Way
| MARIA     | MILLER    | 900 Santiago de Compostela Parkway
| SUSAN     | WILSON    | 478 Joliet Way
| MARGARET  | MOORE     | 613 Korolev Drive
...
| TERRANCE  | ROUSH     | 42 Fontana Avenue
| RENE      | MCALISTER | 1895 Zhezqazghan Drive
| EDUARDO   | HIATT     | 1837 Kaduna Parkway
| TERRENCE  | GUNDERSON | 844 Bucuresti Place
| ENRIQUE   | FORSYTHE  | 1101 Bucuresti Boulevard
| FREDDIE   | DUGGAN    | 1103 Quilmes Boulevard
| WADE      | DELVALLE  | 1331 Usak Boulevard
| AUSTIN    | CINTRON   | 1325 Fukuyama Street
+-----+-----+-----+
599 rows in set (0.00 sec)
```

This older method of specifying joins does not include the `on` subclause; instead, tables are named in the `from` clause separated by commas, and join conditions are included in the `where` clause. While you may decide to ignore the SQL92 syntax in favor of the older join syntax, the ANSI join syntax has the following advantages:

- Join conditions and filter conditions are separated into two different clauses (the `on` subclause and the `where` clause, respectively), making a query easier to understand.
- The join conditions for each pair of tables are contained in their own `on` clause, making it less likely that part of a join will be mistakenly omitted.
- Queries that use the SQL92 join syntax are portable across database servers, whereas the older syntax is slightly different across the different servers.

The benefits of the SQL92 join syntax are easier to identify for complex queries that include both join and filter conditions. Consider the following query, which returns only those customers whose postal code is 52137:

```
mysql> SELECT c.first_name, c.last_name, a.address
-> FROM customer c, address a
-> WHERE c.address_id = a.address_id
-> AND a.postal_code = 52137;
+-----+-----+-----+
| first_name | last_name | address          |
+-----+-----+-----+
| JAMES      | GANNON    | 1635 Kuwana Boulevard |
| FREDDIE    | DUGGAN    | 1103 Quilmes Boulevard |
+-----+-----+-----+
2 rows in set (0.01 sec)
```

At first glance, it is not so easy to determine which conditions in the `where` clause are join conditions and which are filter conditions. It is also not readily apparent which type of join is being employed (to identify the type of join, you would need to look closely at the join conditions in the `where` clause to see whether any special characters are employed), nor is it easy to determine whether any join conditions have been mistakenly left out. Here's the same query using the SQL92 join syntax:

```
mysql> SELECT c.first_name, c.last_name, a.address
-> FROM customer c INNER JOIN address a
-> ON c.address_id = a.address_id
-> WHERE a.postal_code = 52137;
+-----+-----+-----+
| first_name | last_name | address          |
+-----+-----+-----+
| JAMES      | GANNON    | 1635 Kuwana Boulevard |
| FREDDIE    | DUGGAN    | 1103 Quilmes Boulevard |
+-----+-----+-----+
2 rows in set (0.00 sec)
```

With this version, it is clear which condition is used for the join and which condition is used for filtering. Hopefully, you will agree that the version using SQL92 join syntax is easier to understand.

Joining Three or More Tables

Joining three tables is similar to joining two tables, but with one slight wrinkle. With a two-table join, there are two tables and one join type in the `from` clause, and a single `on` subclause to define how the tables are joined. With a three-table join, there are three tables and two join types in the `from` clause, and two `on` subclauses.

To illustrate, let's change the previous query to return the customer's city rather than their street address. The city name, however, is not stored in the `address` table but is accessed via a foreign key to the `city` table. Here are the table definitions:

```
mysql> desc address;
+-----+-----+-----+-----+
| Field | Type      | Null | Key | Default |
+-----+-----+-----+-----+
| address_id | smallint(5) unsigned | NO | PRI | NULL |
| address | varchar(50) | NO |   | NULL |
| address2 | varchar(50) | YES |   | NULL |
| district | varchar(20) | NO |   | NULL |
| city_id | smallint(5) unsigned | NO | MUL | NULL |
| postal_code | varchar(10) | YES |   | NULL |
| phone | varchar(20) | NO |   | NULL |
| location | geometry | NO | MUL | NULL |
| last_update | timestamp | NO |   | CURRENT_TIMESTAMP |
+-----+-----+-----+-----+

mysql> desc city;
+-----+-----+-----+-----+
| Field | Type      | Null | Key | Default |
+-----+-----+-----+-----+
| city_id | smallint(5) unsigned | NO | PRI | NULL |
| city | varchar(50) | NO |   | NULL |
| country_id | smallint(5) unsigned | NO | MUL | NULL |
| last_update | timestamp | NO |   | CURRENT_TIMESTAMP |
+-----+-----+-----+-----+
```

To show each customer's city, you will need to traverse from the `customer` table to the `address` table using the `address_id` column and then from the `address` table to the `city` table using the `city_id` column. The query would look like the following:

```
mysql> SELECT c.first_name, c.last_name, ct.city
    -> FROM customer c
    ->     INNER JOIN address a
    ->         ON c.address_id = a.address_id
    ->     INNER JOIN city ct
    ->         ON a.city_id = ct.city_id;
+-----+-----+-----+
| first_name | last_name | city |
+-----+-----+-----+
| JULIE     | SANCHEZ  | A Corua (La Corua) |
| PEGGY     | MYERS    | Abha   |
+-----+-----+-----+
```

TOM	MILNER	Abu Dhabi
GLEN	TALBERT	Acua
LARRY	THRASHER	Adana
SEAN	DOUGLASS	Addis Abeba
...		
MICHELE	GRANT	Yuncheng
GARY	COY	Yuzhou
PHYLLIS	FOSTER	Zalantun
CHARLENE	ALVAREZ	Zanzibar
FRANKLIN	TROUTMAN	Zaoyang
FLOYD	GANDY	Zapopan
CONSTANCE	REID	Zaria
JACK	FOUST	Zelevnogorsk
BYRON	BOX	Zhezqazghan
GUY	BROWNLEE	Zhoushan
RONNIE	RICKETTS	Ziguinchor

599 rows in set (0.03 sec)

For this query, there are three tables, two join types, and two `on` subclauses in the `from` clause, so things have gotten quite a bit busier. At first glance, it might seem like the order in which the tables appear in the `from` clause is important, but if you switch the table order, you will get the exact same results. All three of these variations return the same results:

```

SELECT c.first_name, c.last_name, ct.city
FROM customer c
    INNER JOIN address a
        ON c.address_id = a.address_id
    INNER JOIN city ct
        ON a.city_id = ct.city_id;

SELECT c.first_name, c.last_name, ct.city
FROM city ct
    INNER JOIN address a
        ON a.city_id = ct.city_id
    INNER JOIN customer c
        ON c.address_id = a.address_id;

SELECT c.first_name, c.last_name, ct.city
FROM address a
    INNER JOIN city ct
        ON a.city_id = ct.city_id
    INNER JOIN customer c
        ON c.address_id = a.address_id;

```

The only difference you may see would be the order in which the rows are returned, since there is no `order by` clause to specify how the results should be ordered.

Does Join Order Matter?

If you are confused about why all three versions of the `customer/address/city` query yield the same results, keep in mind that SQL is a nonprocedural language, meaning that you describe what you want to retrieve and which database objects need to be involved, but it is up to the database server to determine how best to execute your query. Using statistics gathered from your database objects, the server must pick one of three tables as a starting point (the chosen table is thereafter known as the *driving table*) and then decide in which order to join the remaining tables. Therefore, the order in which tables appear in your `from` clause is not significant.

If, however, you believe that the tables in your query should always be joined in a particular order, you can place the tables in the desired order and then specify the keyword `straight_join` in MySQL, request the `force_order` option in SQL Server, or use either the `ordered` or the `leading` optimizer hint in Oracle Database. For example, to tell the MySQL server to use the `city` table as the driving table and to then join the `address` and `customer` tables, you could do the following:

```
SELECT STRAIGHT_JOIN c.first_name, c.last_name, ct.city
  FROM city ct
    INNER JOIN address a
      ON a.city_id = ct.city_id
    INNER JOIN customer c
      ON c.address_id = a.address_id
```

Using Subqueries as Tables

You have already seen several examples of queries that include multiple tables, but there is one variation worth mentioning: what to do if some of the data sets are generated by subqueries. Subqueries are the focus of [Chapter 9](#), but I already introduced the concept of a subquery in the `from` clause in the previous chapter. The following query joins the `customer` table to a subquery against the `address` and `city` tables:

```
mysql> SELECT c.first_name, c.last_name, addr.address, addr.city
->   FROM customer c
->     INNER JOIN
->       (SELECT a.address_id, a.address, ct.city
->         FROM address a
->           INNER JOIN city ct
->             ON a.city_id = ct.city_id
->             WHERE a.district = 'California'
->       ) addr
->     ON c.address_id = addr.address_id;
+-----+-----+-----+-----+
| first_name | last_name | address          | city      |
+-----+-----+-----+-----+
| PATRICIA  | JOHNSON   | 1121 Loja Avenue | San Bernardino |
```

```

| BETTY      | WHITE      | 770 Bydgoszcz Avenue | Citrus Heights |
| ALICE      | STEWART    | 1135 Izumisano Parkway | Fontana        |
| ROSA       | REYNOLDS   | 793 Cam Ranh Avenue  | Lancaster     |
| RENEE      | LANE        | 533 al-Ayn Boulevard | Compton       |
| KRISTIN    | JOHNSTON   | 226 Brest Manor     | Sunnyvale     |
| CASSANDRA  | WALTERS    | 920 Kumbakonam Loop  | Salinas       |
| JACOB      | LANCE       | 1866 al-Qatif Avenue | El Monte      |
| RENE       | MCALISTER   | 1895 Zhezqazghan Drive | Garden Grove  |
+-----+-----+-----+-----+
9 rows in set (0.00 sec)

```

The subquery, which starts on line 4 and is given the alias `addr`, finds all addresses that are in California. The outer query joins the subquery results to the `customer` table to return the first name, last name, street address, and city of all customers who live in California. While this query could have been written without the use of a subquery by simply joining the three tables, it can sometimes be advantageous from a performance and/or readability aspect to use one or more subqueries.

One way to visualize what is going on is to run the subquery by itself and look at the results. Here are the results of the subquery from the prior example:

```

mysql> SELECT a.address_id, a.address, ct.city
-> FROM address a
-> INNER JOIN city ct
-> ON a.city_id = ct.city_id
-> WHERE a.district = 'California';
+-----+-----+-----+
| address_id | address           | city      |
+-----+-----+-----+
|       6    | 1121 Loja Avenue  | San Bernardino |
|      18    | 770 Bydgoszcz Avenue | Citrus Heights |
|      55    | 1135 Izumisano Parkway | Fontana        |
|     116    | 793 Cam Ranh Avenue | Lancaster     |
|     186    | 533 al-Ayn Boulevard | Compton       |
|     218    | 226 Brest Manor    | Sunnyvale     |
|     274    | 920 Kumbakonam Loop | Salinas       |
|     425    | 1866 al-Qatif Avenue | El Monte      |
|     599    | 1895 Zhezqazghan Drive | Garden Grove  |
+-----+-----+-----+
9 rows in set (0.00 sec)

```

This result set consists of all nine California addresses. When joined to the `customer` table via the `address_id` column, your result set will contain information about the customers assigned to these addresses.

Using the Same Table Twice

If you are joining multiple tables, you might find that you need to join the same table more than once. In the sample database, for example, actors are related to the films in which they appeared via the `film_actor` table. If you want to find all of the films in

which two specific actors appear, you could write a query such as this one, which joins the `film` table to the `film_actor` table to the `actor` table:

```
mysql> SELECT f.title
->   FROM film f
->   INNER JOIN film_actor fa
->   ON f.film_id = fa.film_id
->   INNER JOIN actor a
->   ON fa.actor_id = a.actor_id
->   WHERE ((a.first_name = 'CATE' AND a.last_name = 'MCQUEEN')
->          OR (a.first_name = 'CUBA' AND a.last_name = 'BIRCH'));
+-----+
| title      |
+-----+
| ATLANTIS CAUSE    |
| BLOOD ARGONAUTS  |
| COMMANDMENTS EXPRESS |
| DYNAMITE TARZAN    |
| EDGE KISSING      |
...
| TOWERS HURRICANE  |
| TROJAN TOMORROW   |
| VIRGIN DAISY       |
| VOLCANO TEXAS     |
| WATERSHIP FRONTIER |
+-----+
54 rows in set (0.00 sec)
```

This query returns all movies in which either Cate McQueen or Cuba Birch appeared. However, let's say that you want to retrieve only those films in which *both* of these actors appeared. To accomplish this, you will need to find all rows in the `film` table that have two rows in the `film_actor` table, one of which is associated with Cate McQueen, and the other associated with Cuba Birch. Therefore, you will need to include the `film_actor` and `actor` tables twice, each with a different alias so that the server knows which one you are referring to in the various clauses:

```
mysql> SELECT f.title
->   FROM film f
->   INNER JOIN film_actor fa1
->   ON f.film_id = fa1.film_id
->   INNER JOIN actor a1
->   ON fa1.actor_id = a1.actor_id
->   INNER JOIN film_actor fa2
->   ON f.film_id = fa2.film_id
->   INNER JOIN actor a2
->   ON fa2.actor_id = a2.actor_id
->   WHERE (a1.first_name = 'CATE' AND a1.last_name = 'MCQUEEN')
->          AND (a2.first_name = 'CUBA' AND a2.last_name = 'BIRCH');
+-----+
| title      |
+-----+
```

```
| BLOOD ARGONAUTS  |
| TOWERS HURRICANE |
+-----+
2 rows in set (0.00 sec)
```

Between them, the two actors appeared in 52 different films, but there are only two films in which both actors appeared. This is one example of a query that *requires* the use of table aliases, since the same tables are used multiple times.

Self-Joins

Not only can you include the same table more than once in the same query, but you can actually join a table to itself. This might seem like a strange thing to do at first, but there are valid reasons for doing so. Some tables include a *self-referencing foreign key*, which means that it includes a column that points to the primary key within the same table. While the sample database doesn't include such a relationship, let's imagine that the `film` table includes the column `prequel_film_id`, which points to the film's parent (e.g., the film *Fiddler Lost II* would use this column to point to the parent film *Fiddler Lost*). Here's what the table would look like if we were to add this additional column:

```
mysql> desc film;
+-----+-----+-----+-----+-----+
| Field          | Type           | Null | Key | Default |
+-----+-----+-----+-----+-----+
| film_id        | smallint(5) unsigned | NO   | PRI | NULL    |
| title          | varchar(255)      | NO   | MUL | NULL    |
| description     | text             | YES  |      | NULL    |
| release_year    | year(4)          | YES  |      | NULL    |
| language_id     | tinyint(3) unsigned | NO   | MUL | NULL    |
| original_language_id | tinyint(3) unsigned | YES  | MUL | NULL    |
| rental_duration | tinyint(3) unsigned | NO   |      | 3       |
| rental_rate     | decimal(4,2)       | NO   |      | 4.99    |
| length          | smallint(5) unsigned | YES  |      | NULL    |
| replacement_cost | decimal(5,2)       | NO   |      | 19.99   |
| rating          | enum('G','PG','PG-13',
          |                  'R','NC-17') | YES  |      | G       |
| special_features | set('Trailers',...,
          |                  'Behind the Scenes') | YES  |      | NULL    | |
| last_update      | timestamp         | NO   |      | CURRENT_TIMESTAMP |
| prequel_film_id  | smallint(5) unsigned | YES  | MUL | NULL    |
+-----+-----+-----+-----+-----+
```

Using a *self-join*, you can write a query that lists every film that has a prequel, along with the prequel's title:

```
mysql> SELECT f.title, f_prnt.title prequel
-> FROM film f
-> INNER JOIN film f_prnt
```

```

-> ON f_prnt.film_id = f.prequel_film_id
-> WHERE f.prequel_film_id IS NOT NULL;
+-----+-----+
| title          | prequel      |
+-----+-----+
| FIDDLER LOST II | FIDDLER LOST |
+-----+-----+
1 row in set (0.00 sec)

```

This query joins the `film` table to itself using the `prequel_film_id` foreign key, and the table aliases `f` and `f_prnt` are assigned in order to make it clear which table is used for which purpose.

Test Your Knowledge

The following exercises are designed to test your understanding of inner joins. Please see [Appendix B](#) for the solutions to these exercises.

Exercise 5-1

Fill in the blanks (denoted by `<#>`) for the following query to obtain the results that follow:

```

mysql> SELECT c.first_name, c.last_name, a.address, ct.city
-> FROM customer c
-> INNER JOIN address <1>
-> ON c.address_id = a.address_id
-> INNER JOIN city ct
-> ON a.city_id = <2>
-> WHERE a.district = 'California';
+-----+-----+-----+-----+
| first_name | last_name | address           | city      |
+-----+-----+-----+-----+
| PATRICIA   | JOHNSON   | 1121 Loja Avenue    | San Bernardino |
| BETTY      | WHITE      | 770 Bydgoszcz Avenue | Citrus Heights |
| ALICE      | STEWART    | 1135 Izumisano Parkway | Fontana     |
| ROSA       | REYNOLDS   | 793 Cam Ranh Avenue | Lancaster   |
| RENEE      | LANE       | 533 al-Ayn Boulevard | Compton     |
| KRISTIN    | JOHNSTON   | 226 Brest Manor     | Sunnyvale   |
| CASSANDRA  | WALTERS    | 920 Kumbakonam Loop | Salinas     |
| JACOB      | LANCE      | 1866 al-Qatif Avenue | El Monte    |
| RENE       | MCALISTER  | 1895 Zhezqazghan Drive | Garden Grove |
+-----+-----+-----+-----+
9 rows in set (0.00 sec)

```

Exercise 5-2

Write a query that returns the title of every film in which an actor with the first name JOHN appeared.

Exercise 5-3

Construct a query that returns all addresses that are in the same city. You will need to join the address table to itself, and each row should include two different addresses.

CHAPTER 6

Working with Sets

Although you can interact with the data in a database one row at a time, relational databases are really all about sets. This chapter explores how you can combine multiple result sets using various set operators. After a quick overview of set theory, I'll demonstrate how to use the set operators `union`, `intersect`, and `except` to blend multiple data sets together.

Set Theory Primer

In many parts of the world, basic set theory is included in elementary-level math curriculums. Perhaps you recall looking at something like what is shown in [Figure 6-1](#).

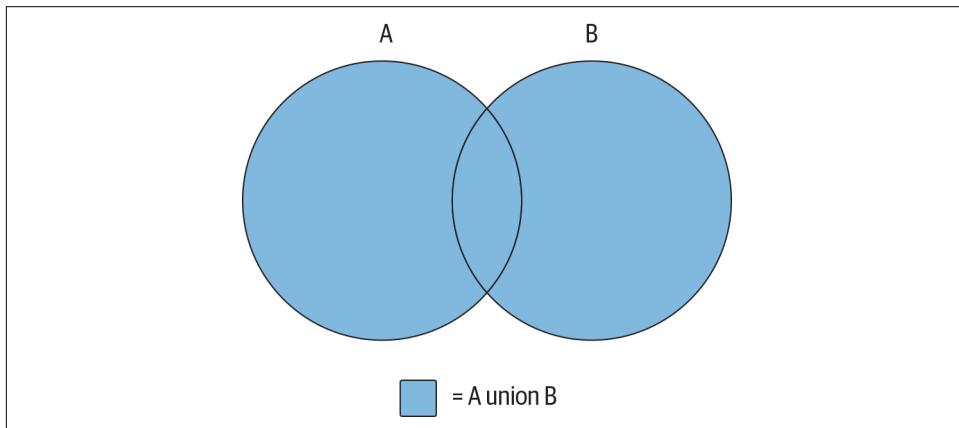


Figure 6-1. The union operation

The shaded area in [Figure 6-1](#) represents the *union* of sets A and B, which is the combination of the two sets (with any overlapping regions included only once). Is this

starting to look familiar? If so, then you'll finally get a chance to put that knowledge to use; if not, don't worry, because it's easy to visualize using a couple of diagrams.

Using circles to represent two data sets (A and B), imagine a subset of data that is common to both sets; this common data is represented by the overlapping area shown in [Figure 6-1](#). Since set theory is rather uninteresting without an overlap between data sets, I use the same diagram to illustrate each set operation. There is another set operation that is concerned *only* with the overlap between two data sets; this operation is known as the *intersection* and is demonstrated in [Figure 6-2](#).

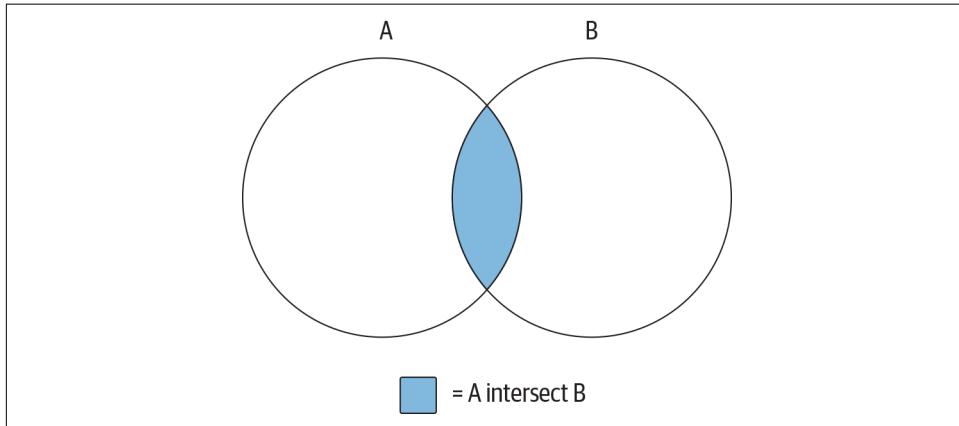


Figure 6-2. The intersection operation

The data set generated by the intersection of sets A and B is just the area of overlap between the two sets. If the two sets have no overlap, then the intersection operation yields the empty set.

The third and final set operation, which is demonstrated in [Figure 6-3](#), is known as the *except* operation.

[Figure 6-3](#) shows the results of A except B, which is the whole of set A minus any overlap with set B. If the two sets have no overlap, then the operation A except B yields the whole of set A.

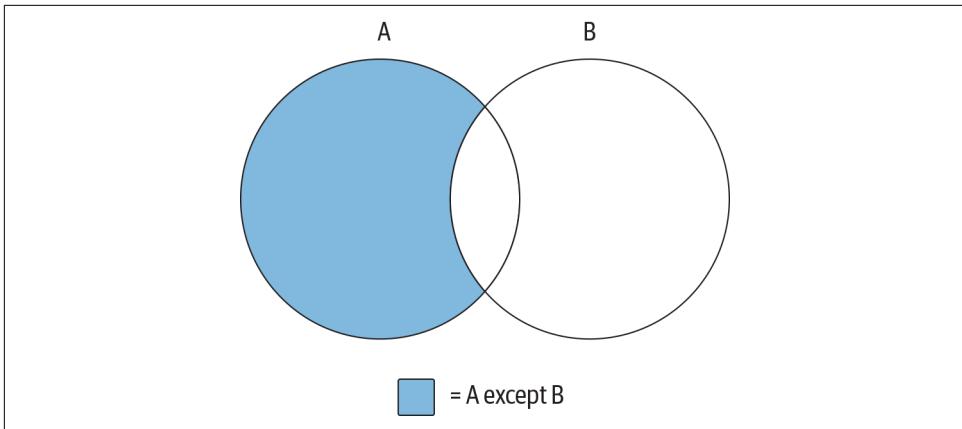


Figure 6-3. The except operation

Using these three operations, or by combining different operations together, you can generate whatever results you need. For example, imagine that you want to build a set demonstrated by [Figure 6-4](#).

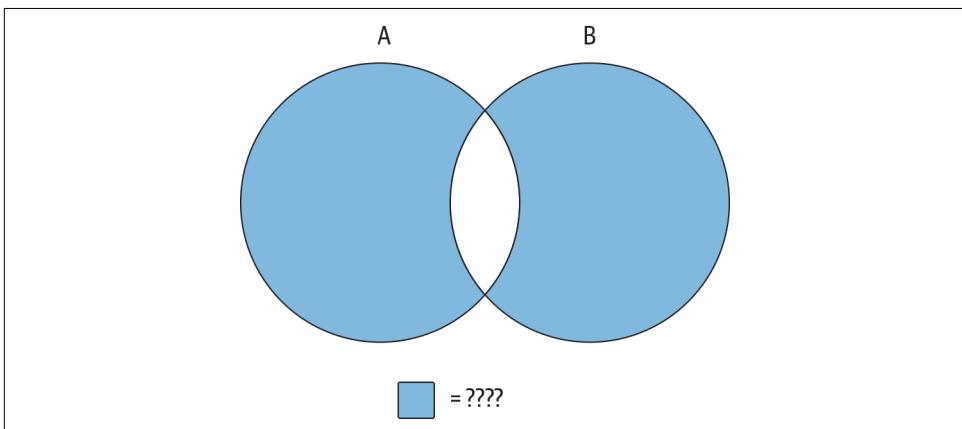


Figure 6-4. Mystery data set

The data set you are looking for includes all of sets A and B *without* the overlapping region. You can't achieve this outcome with just one of the three operations shown earlier; instead, you will need to first build a data set that encompasses all of sets A and B, and then utilize a second operation to remove the overlapping region. If the combined set is described as `A union B`, and the overlapping region is described as `A intersect B`, then the operation needed to generate the data set represented by [Figure 6-4](#) would look as follows:

`(A union B) except (A intersect B)`

Of course, there are often multiple ways to achieve the same results; you could reach a similar outcome using the following operation:

```
(A except B) union (B except A)
```

While these concepts are fairly easy to understand using diagrams, the next sections show you how these concepts are applied to a relational database using the SQL set operators.

Set Theory in Practice

The circles used in the previous section's diagrams to represent data sets don't convey anything about what the data sets comprise. When dealing with actual data, however, there is a need to describe the composition of the data sets involved if they are to be combined. Imagine, for example, what would happen if you tried to generate the union of the `customer` table and the `city` table, whose definitions are as follows:

```
mysql> desc customer;
+-----+-----+-----+-----+
| Field | Type      | Null | Key | Default |
+-----+-----+-----+-----+
| customer_id | smallint(5) unsigned | NO   | PRI  | NULL    |
| store_id    | tinyint(3) unsigned  | NO   | MUL  | NULL    |
| first_name  | varchar(45)        | NO   |       | NULL    |
| last_name   | varchar(45)        | NO   | MUL  | NULL    |
| email        | varchar(50)         | YES  |       | NULL    |
| address_id  | smallint(5) unsigned | NO   | MUL  | NULL    |
| active       | tinyint(1)          | NO   |       | 1       |
| create_date  | datetime             | NO   |       | NULL    |
| last_update  | timestamp            | YES  |       | CURRENT_TIMESTAMP |
+-----+-----+-----+-----+
mysql> desc city;
+-----+-----+-----+-----+
| Field | Type      | Null | Key | Default |
+-----+-----+-----+-----+
| city_id | smallint(5) unsigned | NO   | PRI  | NULL    |
| city    | varchar(50)        | NO   |       | NULL    |
| country_id | smallint(5) unsigned | NO   | MUL  | NULL    |
| last_update | timestamp           | NO   |       | CURRENT_TIMESTAMP |
+-----+-----+-----+-----+
```

When combined, the first column in the result set would include both the `customer.customer_id` and `city.city_id` columns, the second column would be the combination of the `customer.store_id` and `city.city` columns, and so forth. While some column pairs are easy to combine (e.g., two numeric columns), it is unclear how other column pairs should be combined, such as a numeric column with a string column or a string column with a date column. Additionally, the fifth through ninth columns of the combined tables would include data from only the

`customer` table's fifth through ninth columns, since the `city` table has only four columns. Clearly, there needs to be some commonality between two data sets that you wish to combine.

Therefore, when performing set operations on two data sets, the following guidelines must apply:

- Both data sets must have the same number of columns.
- The data types of each column across the two data sets must be the same (or the server must be able to convert one to the other).

With these rules in place, it is easier to envision what “overlapping data” means in practice; each column pair from the two sets being combined must contain the same string, number, or date for rows in the two tables to be considered the same.

You perform a set operation by placing a *set operator* between two `select` statements, as demonstrated by the following:

```
mysql> SELECT 1 num, 'abc' str
      -> UNION
      -> SELECT 9 num, 'xyz' str;
+-----+
| num | str |
+-----+
|   1 | abc |
|   9 | xyz |
+-----+
2 rows in set (0.02 sec)
```

Each of the individual queries yields a data set consisting of a single row having a numeric column and a string column. The set operator, which in this case is `union`, tells the database server to combine all rows from the two sets. Thus, the final set includes two rows of two columns. This query is known as a *compound query* because it comprises multiple, otherwise-independent queries. As you will see later, compound queries may include *more* than two queries if multiple set operations are needed to attain the final results.

Set Operators

The SQL language includes three set operators that allow you to perform each of the various set operations described earlier in the chapter. Additionally, each set operator has two flavors, one that includes duplicates and another that removes duplicates (but not necessarily *all* of the duplicates). The following subsections define each operator and demonstrate how they are used.

The union Operator

The `union` and `union all` operators allow you to combine multiple data sets. The difference between the two is that `union` sorts the combined set and removes duplicates, whereas `union all` does not. With `union all`, the number of rows in the final data set will always equal the sum of the number of rows in the sets being combined. This operation is the simplest set operation to perform (from the server's point of view), since there is no need for the server to check for overlapping data. The following example demonstrates how you can use the `union all` operator to generate a set of first and last names from multiple tables:

```
mysql> SELECT 'CUST' typ, c.first_name, c.last_name
-> FROM customer c
-> UNION ALL
-> SELECT 'ACTR' typ, a.first_name, a.last_name
-> FROM actor a;
+-----+-----+
| typ | first_name | last_name |
+-----+-----+
| CUST | MARY      | SMITH     |
| CUST | PATRICIA  | JOHNSON   |
| CUST | LINDA     | WILLIAMS  |
| CUST | BARBARA   | JONES     |
| CUST | ELIZABETH | BROWN     |
| CUST | JENNIFER  | DAVIS     |
| CUST | MARIA     | MILLER    |
| CUST | SUSAN     | WILSON    |
| CUST | MARGARET | MOORE     |
| CUST | DOROTHY   | TAYLOR    |
| CUST | LISA      | ANDERSON  |
| CUST | NANCY     | THOMAS    |
| CUST | KAREN     | JACKSON   |
...
| ACTR | BURT      | TEMPLE    |
| ACTR | MERYL     | ALLEN     |
| ACTR | JAYNE     | SILVERSTONE |
| ACTR | BELA      | WALKEN    |
| ACTR | REESE     | WEST      |
| ACTR | MARY      | KEITEL    |
| ACTR | JULIA     | FAWCETT   |
| ACTR | THORA     | TEMPLE    |
+-----+-----+
799 rows in set (0.00 sec)
```

The query returns 799 names, with 599 rows coming from the `customer` table and the other 200 coming from the `actor` table. The first column, which has the alias `typ`, is not necessary, but was added to show the source of each name returned by the query.

Just to drive home the point that the `union all` operator doesn't remove duplicates, here's another version of the previous example, but with two identical queries against the `actor` table:

```
mysql> SELECT 'ACTR' typ, a.first_name, a.last_name
      -> FROM actor a
      -> UNION ALL
      -> SELECT 'ACTR' typ, a.first_name, a.last_name
      -> FROM actor a;
+-----+-----+-----+
| typ | first_name | last_name |
+-----+-----+-----+
| ACTR | PENELOPE | GUINNESS |
| ACTR | NICK     | WAHLBERG |
| ACTR | ED       | CHASE    |
| ACTR | JENNIFER | DAVIS    |
| ACTR | JOHNNY   | LOLLOBRIGIDA |
| ACTR | BETTE   | NICHOLSON |
| ACTR | GRACE   | MOSTEL   |
...
| ACTR | BURT     | TEMPLE   |
| ACTR | MERYL   | ALLEN    |
| ACTR | JAYNE   | SILVERSTONE |
| ACTR | BELA     | WALKEN   |
| ACTR | REESE   | WEST     |
| ACTR | MARY     | KEITEL   |
| ACTR | JULIA   | FAWCETT  |
| ACTR | THORA   | TEMPLE   |
+-----+-----+-----+
400 rows in set (0.00 sec)
```

As you can see by the results, the 200 rows from the `actor` table are included twice, for a total of 400 rows.

While you are unlikely to repeat the same query twice in a compound query, here is another compound query that returns duplicate data:

```
mysql> SELECT c.first_name, c.last_name
      -> FROM customer c
      -> WHERE c.first_name LIKE 'J%' AND c.last_name LIKE 'D%'
      -> UNION ALL
      -> SELECT a.first_name, a.last_name
      -> FROM actor a
      -> WHERE a.first_name LIKE 'J%' AND a.last_name LIKE 'D%';
+-----+-----+
| first_name | last_name |
+-----+-----+
| JENNIFER  | DAVIS    |
| JENNIFER  | DAVIS    |
| JUDY       | DEAN     |
| JODIE      | DEGENERES |
| JULIANNE   | DENCH    |
+-----+-----+
```

```
+-----+-----+
5 rows in set (0.00 sec)
```

Both queries return the names of people having the initials JD. Of the five rows in the result set, one of them is a duplicate (Jennifer Davis). If you would like your combined table to *exclude* duplicate rows, you need to use the `union` operator instead of `union all`:

```
mysql> SELECT c.first_name, c.last_name
-> FROM customer c
-> WHERE c.first_name LIKE 'J%' AND c.last_name LIKE 'D%'
-> UNION
-> SELECT a.first_name, a.last_name
-> FROM actor a
-> WHERE a.first_name LIKE 'J%' AND a.last_name LIKE 'D%';
+-----+-----+
| first_name | last_name |
+-----+-----+
| JENNIFER   | DAVIS      |
| JUDY        | DEAN       |
| JODIE       | DEGENERES |
| JULIANNE    | DENCH      |
+-----+-----+
4 rows in set (0.00 sec)
```

For this version of the query, only the four distinct names are included in the result set, rather than the five rows returned when using `union all`.

The `intersect` Operator

The ANSI SQL specification includes the `intersect` operator for performing intersections. Unfortunately, version 8.0 of MySQL does not implement the `intersect` operator. If you are using Oracle or SQL Server 2008, you will be able to use `intersect`; since I am using MySQL for all examples in this book, however, the result sets for the example queries in this section are fabricated and cannot be executed with any versions up to and including version 8.0. I also refrain from showing the MySQL prompt (`mysql>`), since the statements are not being executed by the MySQL server.

If the two queries in a compound query return nonoverlapping data sets, then the intersection will be an empty set. Consider the following query:

```
SELECT c.first_name, c.last_name
FROM customer c
WHERE c.first_name LIKE 'D%' AND c.last_name LIKE 'T%'
INTERSECT
SELECT a.first_name, a.last_name
FROM actor a
WHERE a.first_name LIKE 'D%' AND a.last_name LIKE 'T%';
Empty set (0.04 sec)
```

While there are both actors and customers having the initials DT, these sets are completely nonoverlapping, so the intersection of the two sets yields the empty set. If we switch back to the initials JD, however, the intersection will yield a single row:

```
SELECT c.first_name, c.last_name
FROM customer c
WHERE c.first_name LIKE 'J%' AND c.last_name LIKE 'D%'
INTERSECT
SELECT a.first_name, a.last_name
FROM actor a
WHERE a.first_name LIKE 'J%' AND a.last_name LIKE 'D%';
+-----+-----+
| first_name | last_name |
+-----+-----+
| JENNIFER   | DAVIS      |
+-----+-----+
1 row in set (0.00 sec)
```

The intersection of these two queries yields Jennifer Davis, which is the only name found in both queries' result sets.

Along with the `intersect` operator, which removes any duplicate rows found in the overlapping region, the ANSI SQL specification calls for an `intersect all` operator, which does not remove duplicates. The only database server that currently implements the `intersect all` operator is IBM's DB2 Universal Server.

The `except` Operator

The ANSI SQL specification includes the `except` operator for performing the `except` operation. Once again, unfortunately, version 8.0 of MySQL does not implement the `except` operator, so the same rules apply for this section as for the previous section.



If you are using Oracle Database, you will need to use the non-ANSI-compliant `minus` operator instead.

The `except` operator returns the first result set minus any overlap with the second result set. Here's the example from the previous section, but using `except` instead of `intersect`, and with the order of the queries reversed:

```
SELECT a.first_name, a.last_name
FROM actor a
WHERE a.first_name LIKE 'J%' AND a.last_name LIKE 'D%'
EXCEPT
SELECT c.first_name, c.last_name
FROM customer c
WHERE c.first_name LIKE 'J%' AND c.last_name LIKE 'D%';
```

```

+-----+
| first_name | last_name |
+-----+
| JUDY       | DEAN      |
| JODIE      | DEGENERES |
| JULIANNE   | DENCH     |
+-----+
3 rows in set (0.00 sec)

```

In this version of the query, the result set consists of the three rows from the first query minus Jennifer Davis, who is found in the result sets from both queries. There is also an `except all` operator specified in the ANSI SQL specification, but once again, only IBM's DB2 Universal Server has implemented the `except all` operator.

The `except all` operator is a bit tricky, so here is an example that demonstrates how duplicate data is handled. Let's say you have two data sets that look like the following:

Set A

```

+-----+
| actor_id |
+-----+
|    10 |
|    11 |
|    12 |
|    10 |
|    10 |
+-----+

```

Set B

```

+-----+
| actor_id |
+-----+
|    10 |
|    10 |
+-----+

```

The operation `A except B` yields the following:

```

+-----+
| actor_id |
+-----+
|    11 |
|    12 |
+-----+

```

If you change the operation to `A except all B`, you will see the following:

```
+-----+
| actor_id |
+-----+
|      10 |
|      11 |
|      12 |
+-----+
```

Therefore, the difference between the two operations is that `except` removes all occurrences of duplicate data from set A, whereas `except all` removes only one occurrence of duplicate data from set A for every occurrence in set B.

Set Operation Rules

The following sections outline some rules that you must follow when working with compound queries.

Sorting Compound Query Results

If you want the results of your compound query to be sorted, you can add an `order by` clause after the last query. When specifying column names in the `order by` clause, you will need to choose from the column names in the first query of the compound query. Frequently, the column names are the same for both queries in a compound query, but this does not need to be the case, as demonstrated by the following:

```
mysql> SELECT a.first_name fname, a.last_name lname
-> FROM actor a
-> WHERE a.first_name LIKE 'J%' AND a.last_name LIKE 'D%'
-> UNION ALL
-> SELECT c.first_name, c.last_name
-> FROM customer c
-> WHERE c.first_name LIKE 'J%' AND c.last_name LIKE 'D%'
-> ORDER BY lname, fname;
+-----+-----+
| fname | lname  |
+-----+-----+
| JENNIFER | DAVIS   |
| JENNIFER | DAVIS   |
| JUDY    | DEAN    |
| JODIE   | DEGENERES |
| JULIANNE | DENCH   |
+-----+-----+
5 rows in set (0.00 sec)
```

The column names specified in the two queries are different in this example. If you specify a column name from the second query in your `order by` clause, you will see the following error:

```
mysql> SELECT a.first_name fname, a.last_name lname
-> FROM actor a
-> WHERE a.first_name LIKE 'J%' AND a.last_name LIKE 'D%'
-> UNION ALL
-> SELECT c.first_name, c.last_name
-> FROM customer c
-> WHERE c.first_name LIKE 'J%' AND c.last_name LIKE 'D%'
-> ORDER BY last_name, first_name;
ERROR 1054 (42S22): Unknown column 'last_name' in 'order clause'
```

I recommend giving the columns in both queries identical column aliases in order to avoid this issue.

Set Operation Precedence

If your compound query contains more than two queries using different set operators, you need to think about the order in which to place the queries in your compound statement to achieve the desired results. Consider the following three-query compound statement:

```
mysql> SELECT a.first_name, a.last_name
-> FROM actor a
-> WHERE a.first_name LIKE 'J%' AND a.last_name LIKE 'D%'
-> UNION ALL
-> SELECT a.first_name, a.last_name
-> FROM actor a
-> WHERE a.first_name LIKE 'M%' AND a.last_name LIKE 'T%'
-> UNION
-> SELECT c.first_name, c.last_name
-> FROM customer c
-> WHERE c.first_name LIKE 'J%' AND c.last_name LIKE 'D%';
+-----+-----+
| first_name | last_name |
+-----+-----+
| JENNIFER   | DAVIS      |
| JUDY        | DEAN       |
| JODIE       | DEGENERES |
| JULIANNE    | DENCH      |
| MARY        | TANDY      |
| MENA        | TEMPLE     |
+-----+-----+
6 rows in set (0.00 sec)
```

This compound query includes three queries that return sets of nonunique names; the first and second queries are separated with the `union all` operator, while the second and third queries are separated with the `union` operator. While it might not seem to make much difference where the `union` and `union all` operators are placed, it does, in fact, make a difference. Here's the same compound query with the set operators reversed:

```

mysql> SELECT a.first_name, a.last_name
      -> FROM actor a
      -> WHERE a.first_name LIKE 'J%' AND a.last_name LIKE 'D%'
      -> UNION
      -> SELECT a.first_name, a.last_name
      -> FROM actor a
      -> WHERE a.first_name LIKE 'M%' AND a.last_name LIKE 'T%'
      -> UNION ALL
      -> SELECT c.first_name, c.last_name
      -> FROM customer c
      -> WHERE c.first_name LIKE 'J%' AND c.last_name LIKE 'D%';
+-----+-----+
| first_name | last_name |
+-----+-----+
| JENNIFER   | DAVIS      |
| JUDY        | DEAN       |
| JODIE       | DEGENERES |
| JULIANNE    | DENCH      |
| MARY        | TANDY      |
| MENA        | TEMPLE     |
| JENNIFER   | DAVIS      |
+-----+-----+
7 rows in set (0.00 sec)

```

Looking at the results, it's obvious that it *does* make a difference how the compound query is arranged when using different set operators. In general, compound queries containing three or more queries are evaluated in order from top to bottom, but with the following caveats:

- The ANSI SQL specification calls for the `intersect` operator to have precedence over the other set operators.
- You may dictate the order in which queries are combined by enclosing multiple queries in parentheses.

MySQL does not yet allow parentheses in compound queries, but if you are using a different database server, you can wrap adjoining queries in parentheses to override the default top-to-bottom processing of compound queries, as in:

```

SELECT a.first_name, a.last_name
FROM actor a
WHERE a.first_name LIKE 'J%' AND a.last_name LIKE 'D%'
UNION
(SELECT a.first_name, a.last_name
 FROM actor a
 WHERE a.first_name LIKE 'M%' AND a.last_name LIKE 'T%'
 UNION ALL
 SELECT c.first_name, c.last_name
 FROM customer c
 WHERE c.first_name LIKE 'J%' AND c.last_name LIKE 'D%'
)

```

For this compound query, the second and third queries would be combined using the `union all` operator, then the results would be combined with the first query using the `union` operator.

Test Your Knowledge

The following exercises are designed to test your understanding of set operations. See [Appendix B](#) for the answers to these exercises.

Exercise 6-1

If set A = {L M N O P} and set B = {P Q R S T}, what sets are generated by the following operations?

- A `union` B
- A `union all` B
- A `intersect` B
- A `except` B

Exercise 6-2

Write a compound query that finds the first and last names of all actors and customers whose last name starts with L.

Exercise 6-3

Sort the results from Exercise 6-2 by the `last_name` column.

Data Generation, Manipulation, and Conversion

As I mentioned in the preface, this book strives to teach generic SQL techniques that can be applied across multiple database servers. This chapter, however, deals with the generation, conversion, and manipulation of string, numeric, and temporal data, and the SQL language does not include commands covering this functionality. Rather, built-in functions are used to facilitate data generation, conversion, and manipulation, and while the SQL standard does specify some functions, the database vendors often do not comply with the function specifications.

Therefore, my approach for this chapter is to show you some of the common ways in which data is generated and manipulated within SQL statements and then demonstrate some of the built-in functions implemented by Microsoft SQL Server, Oracle Database, and MySQL. Along with reading this chapter, I strongly recommend you download a reference guide covering all the functions implemented by your server. If you work with more than one database server, there are several reference guides that cover multiple servers, such as Kevin Kline et al.'s *SQL in a Nutshell* and Jonathan Gennick's *SQL Pocket Guide*, both from O'Reilly.

Working with String Data

When working with string data, you will be using one of the following character data types:

CHAR

Holds fixed-length, blank-padded strings. MySQL allows CHAR values up to 255 characters in length, Oracle Database permits up to 2,000 characters, and SQL Server allows up to 8,000 characters.

`varchar`

Holds variable-length strings. MySQL permits up to 65,535 characters in a `varchar` column, Oracle Database (via the `varchar2` type) allows up to 4,000 characters, and SQL Server allows up to 8,000 characters.

`text` (*MySQL and SQL Server*) or `clob` (*Oracle Database*)

Holds very large variable-length strings (generally referred to as documents in this context). MySQL has multiple text types (`tinytext`, `text`, `mediumtext`, and `longtext`) for documents up to 4 GB in size. SQL Server has a single `text` type for documents up to 2 GB in size, and Oracle Database includes the `clob` data type, which can hold documents up to a whopping 128 TB. SQL Server 2005 also includes the `varchar(max)` data type and recommends its use instead of the `text` type, which will be removed from the server in some future release.

To demonstrate how you can use these various types, I use the following table for some of the examples in this section:

```
CREATE TABLE string_tbl
  (char_fld CHAR(30),
   varchar_fld VARCHAR(30),
   text_fld TEXT
  );
```

The next two subsections show how you can generate and manipulate string data.

String Generation

The simplest way to populate a character column is to enclose a string in quotes, as in the following examples:

```
mysql> INSERT INTO string_tbl (char_fld, varchar_fld, text_fld)
    -> VALUES ('This is char data',
    ->           'This is varchar data',
    ->           'This is text data');
Query OK, 1 row affected (0.00 sec)
```

When inserting string data into a table, remember that if the length of the string exceeds the maximum size for the character column (either the designated maximum or the maximum allowed for the data type), the server will throw an exception. Although this is the default behavior for all three servers, you can configure MySQL and SQL Server to silently truncate the string instead of throwing an exception. To demonstrate how MySQL handles this situation, the following `update` statement attempts to modify the `varchar_fld` column, whose maximum length is defined as 30, with a string that is 46 characters in length:

```
mysql> UPDATE string_tbl
    -> SET varchar_fld = 'This is a piece of extremely long varchar data';
ERROR 1406 (22001): Data too long for column 'varchar_fld' at row 1
```

Since MySQL 6.0, the default behavior is now “strict” mode, which means that exceptions are thrown when problems arise, whereas in older versions of the server the string would have been truncated and a warning issued. If you would rather have the engine truncate the string and issue a warning instead of raising an exception, you can opt to be in ANSI mode. The following example shows how to check which mode you are in and then how to change the mode using the `set` command:

```
mysql> SELECT @@session.sql_mode;
+-----+
| @@session.sql_mode |
+-----+
| STRICT_TRANS_TABLES,NO_ENGINE_SUBSTITUTION |
+-----+
1 row in set (0.00 sec)

mysql> SET sql_mode='ansi';
Query OK, 0 rows affected (0.08 sec)

mysql> SELECT @@session.sql_mode;
+-----+
| @@session.sql_mode |
+-----+
| REAL_AS_FLOAT,PIPES_AS_CONCAT,ANSI_QUOTES,IGNORE_SPACE,ONLY_FULL_GROUP_BY,ANSI |
+-----+
1 row in set (0.00 sec)
```

If you rerun the previous update statement, you will find that the column has been modified, but the following warning is generated:

```
mysql> SHOW WARNINGS;
+-----+-----+-----+
| Level | Code | Message |
+-----+-----+-----+
| Warning | 1265 | Data truncated for column 'vchar_fld' at row 1 |
+-----+-----+-----+
1 row in set (0.00 sec)
```

If you retrieve the `vchar_fld` column, you will see that the string has indeed been truncated:

```
mysql> SELECT vchar_fld
-> FROM string_tbl;
+-----+
| vchar_fld |
+-----+
| This is a piece of extremely l |
+-----+
1 row in set (0.05 sec)
```

As you can see, only the first 30 characters of the 46-character string made it into the `vchar_fld` column. The best way to avoid string truncation (or exceptions, in the

case of Oracle Database or MySQL in strict mode) when working with `varchar` columns is to set the upper limit of a column to a high enough value to handle the longest strings that might be stored in the column (keeping in mind that the server allocates only enough space to store the string, so it is not wasteful to set a high upper limit for `varchar` columns).

Including single quotes

Since strings are demarcated by single quotes, you will need to be alert for strings that include single quotes or apostrophes. For example, you won't be able to insert the following string because the server will think that the apostrophe in the word *doesn't* marks the end of the string:

```
UPDATE string_tbl  
SET text_fld = 'This string doesn't work';
```

To make the server ignore the apostrophe in the word *doesn't*, you will need to add an *escape* to the string so that the server treats the apostrophe like any other character in the string. All three servers allow you to escape a single quote by adding another single quote directly before, as in:

```
mysql> UPDATE string_tbl  
-> SET text_fld = 'This string didn''t work, but it does now';  
Query OK, 1 row affected (0.01 sec)  
Rows matched: 1  Changed: 1  Warnings: 0
```



Oracle Database and MySQL users may also choose to escape a single quote by adding a backslash character immediately before, as in:

```
UPDATE string_tbl SET text_fld =  
'This string didn\'t work, but it does now'
```

If you retrieve a string for use in a screen or report field, you don't need to do anything special to handle embedded quotes:

```
mysql> SELECT text_fld  
-> FROM string_tbl;  
+-----+  
| text_fld |  
+-----+  
| This string didn't work, but it does now |  
+-----+  
1 row in set (0.00 sec)
```

However, if you are retrieving the string to add to a file that another program will read, you may want to include the escape as part of the retrieved string. If you are using MySQL, you can use the built-in function `quote()`, which places quotes around

the entire string *and* adds escapes to any single quotes/apostrophes within the string. Here's what our string looks like when retrieved via the `quote()` function:

```
mysql> SELECT quote(text_fld)
    -> FROM string_tbl;
+-----+
| QUOTE(text_fld) |
+-----+
| 'This string didn\'t work, but it does now' |
+-----+
1 row in set (0.04 sec)
```

When retrieving data for data export, you may want to use the `quote()` function for all non-system-generated character columns, such as a `customer_notes` column.

Including special characters

If your application is multinational in scope, you might find yourself working with strings that include characters that do not appear on your keyboard. When working with the French and German languages, for example, you might need to include accented characters such as é and ö. The SQL Server and MySQL servers include the built-in function `char()` so that you can build strings from any of the 255 characters in the ASCII character set (Oracle Database users can use the `chr()` function). To demonstrate, the next example retrieves a typed string and its equivalent built via individual characters:

```
mysql> SELECT 'abcdefg', CHAR(97,98,99,100,101,102,103);
+-----+-----+
| abcdefg | CHAR(97,98,99,100,101,102,103) |
+-----+-----+
| abcdefg | abcdefg           |
+-----+-----+
1 row in set (0.01 sec)
```

Thus, the 97th character in the ASCII character set is the letter *a*. While the characters shown in the preceding example are not special, the following examples show the location of the accented characters along with other special characters, such as currency symbols:

```
mysql> SELECT CHAR(128,129,130,131,132,133,134,135,136,137);
+-----+
| CHAR(128,129,130,131,132,133,134,135,136,137) |
+-----+
| Çüéâäàâçéë |
+-----+
1 row in set (0.01 sec)

mysql> SELECT CHAR(138,139,140,141,142,143,144,145,146,147);
+-----+
| CHAR(138,139,140,141,142,143,144,145,146,147) |
+-----+
```

```
+-----+
| èïîìÃÃÉæÆô |
+-----+
1 row in set (0.01 sec)

mysql> SELECT CHAR(148,149,150,151,152,153,154,155,156,157);
+-----+
| CHAR(148,149,150,151,152,153,154,155,156,157) |
+-----+
| öðûùýðÙøEØ |
+-----+
1 row in set (0.00 sec)

mysql> SELECT CHAR(158,159,160,161,162,163,164,165);
+-----+
| CHAR(158,159,160,161,162,163,164,165) |
+-----+
| xfáíóÚñÑ |
+-----+
1 row in set (0.01 sec)
```



I am using the `utf8mb4` character set for the examples in this section. If your session is configured for a different character set, you will see a different set of characters than what is shown here. The same concepts apply, but you will need to familiarize yourself with the layout of your character set to locate specific characters.

Building strings character by character can be quite tedious, especially if only a few of the characters in the string are accented. Fortunately, you can use the `concat()` function to concatenate individual strings, some of which you can type while others you can generate via the `char()` function. For example, the following shows how to build the phrase *danke schön* using the `concat()` and `char()` functions:

```
mysql> SELECT CONCAT('danke sch', CHAR(148), 'n');
+-----+
| CONCAT('danke sch', CHAR(148), 'n') |
+-----+
| danke schön |
+-----+
1 row in set (0.00 sec)
```



Oracle Database users can use the concatenation operator (||) instead of the concat() function, as in:

```
SELECT 'danke sch' || CHR(148) || 'n'  
FROM dual;
```

SQL Server does not include a concat() function, so you will need to use the concatenation operator (+), as in:

```
SELECT 'danke sch' + CHAR(148) + 'n'
```

If you have a character and need to find its ASCII equivalent, you can use the ascii() function, which takes the leftmost character in the string and returns a number:

```
mysql> SELECT ASCII('ö');  
+-----+  
| ASCII('ö') |  
+-----+  
|      148 |  
+-----+  
1 row in set (0.00 sec)
```

Using the char(), ascii(), and concat() functions (or concatenation operators), you should be able to work with any Roman language even if you are using a keyboard that does not include accented or special characters.

String Manipulation

Each database server includes many built-in functions for manipulating strings. This section explores two types of string functions: those that return numbers and those that return strings. Before I begin, however, I reset the data in the `string_tbl` table to the following:

```
mysql> DELETE FROM string_tbl;  
Query OK, 1 row affected (0.02 sec)  
  
mysql> INSERT INTO string_tbl (char_fld, varchar_fld, text_fld)  
    -> VALUES ('This string is 28 characters',  
    ->   'This string is 28 characters',  
    ->   'This string is 28 characters');  
Query OK, 1 row affected (0.00 sec)
```

String functions that return numbers

Of the string functions that return numbers, one of the most commonly used is the `length()` function, which returns the number of characters in the string (SQL Server users will need to use the `len()` function). The following query applies the `length()` function to each column in the `string_tbl` table:

```
mysql> SELECT LENGTH(char_fld) char_length,
->      LENGTH(vchar_fld) varchar_length,
->      LENGTH(text_fld) text_length
->   FROM string_tbl;
+-----+-----+-----+
| char_length | varchar_length | text_length |
+-----+-----+-----+
|        28 |           28 |          28 |
+-----+-----+-----+
1 row in set (0.00 sec)
```

While the lengths of the `varchar` and `text` columns are as expected, you might have expected the length of the `char` column to be 30, since I told you that strings stored in `char` columns are right-padded with spaces. The MySQL server removes trailing spaces from `char` data when it is retrieved, however, so you will see the same results from all string functions regardless of the type of column in which the strings are stored.

Along with finding the length of a string, you might want to find the location of a substring within a string. For example, if you want to find the position at which the string '`characters`' appears in the `vchar_fld` column, you could use the `position()` function, as demonstrated by the following:

```
mysql> SELECT POSITION('characters' IN vchar_fld)
->   FROM string_tbl;
+-----+
| POSITION('characters' IN vchar_fld) |
+-----+
|                   19 |
+-----+
1 row in set (0.12 sec)
```

If the substring cannot be found, the `position()` function returns 0.



For those of you who program in a language such as C or C++, where the first element of an array is at position 0, remember when working with databases that the first character in a string is at position 1. A return value of 0 from `instr()` indicates that the substring could not be found, not that the substring was found at the first position in the string.

If you want to start your search at something other than the first character of your target string, you will need to use the `locate()` function, which is similar to the `position()` function except that it allows an optional third parameter, which is used to define the search's start position. The `locate()` function is also proprietary, whereas the `position()` function is part of the SQL:2003 standard. Here's an example asking

for the position of the string 'is' starting at the fifth character in the varchar_fld column:

```
mysql> SELECT LOCATE('is', varchar_fld, 5)
   -> FROM string_tbl;
+-----+
| LOCATE('is', varchar_fld, 5) |
+-----+
|                  13 |
+-----+
1 row in set (0.02 sec)
```



Oracle Database does not include the `position()` or `locate()` function, but it does include the `instr()` function, which mimics the `position()` function when provided with two arguments and mimics the `locate()` function when provided with three arguments. SQL Server also doesn't include a `position()` or `locate()` function, but it does include the `charindex()` function, which also accepts either two or three arguments similar to Oracle's `instr()` function.

Another function that takes strings as arguments and returns numbers is the string comparison function `strcmp()`. `strcmp()`, which is implemented only by MySQL and has no analog in Oracle Database or SQL Server, takes two strings as arguments and returns one of the following:

- `-1` if the first string comes before the second string in sort order
- `0` if the strings are identical
- `1` if the first string comes after the second string in sort order

To illustrate how the function works, I first show the sort order of five strings using a query and then show how the strings compare to one another using `strcmp()`. Here are the five strings that I insert into the `string_tbl` table:

```
mysql> DELETE FROM string_tbl;
Query OK, 1 row affected (0.00 sec)

mysql> INSERT INTO string_tbl(varchar_fld)
   -> VALUES ('abcd'),
   ->           ('xyz'),
   ->           ('QRSTUV'),
   ->           ('qrstuv'),
   ->           ('12345');
Query OK, 5 rows affected (0.05 sec)
Records: 5  Duplicates: 0  Warnings: 0
```

Here are the five strings in their sort order:

```
mysql> SELECT varchar_fld
-> FROM string_tbl
-> ORDER BY varchar_fld;
+-----+
| varchar_fld |
+-----+
| 12345      |
| abcd       |
| QRSTUV    |
| qrstuv    |
| xyz        |
+-----+
5 rows in set (0.00 sec)
```

The next query makes six comparisons among the five different strings:

```
mysql> SELECT STRCMP('12345','12345') 12345_12345,
->   STRCMP('abcd','xyz') abcd_xyz,
->   STRCMP('abcd','QRSTUV') abcd_QRSTUV,
->   STRCMP('qrstuv','QRSTUV') qrstuv_QRSTUV,
->   STRCMP('12345','xyz') 12345_xyz,
->   STRCMP('xyz','qrstuv') xyz_qrstuv;
+-----+-----+-----+-----+-----+
| 12345_12345 | abcd_xyz | abcd_QRSTUV | qrstuv_QRSTUV | 12345_xyz | xyz_qrstuv |
+-----+-----+-----+-----+-----+
|          0 |      -1 |         -1 |           0 |      -1 |         1 |
+-----+-----+-----+-----+-----+
1 row in set (0.00 sec)
```

The first comparison yields 0, which is to be expected since I compared a string to itself. The fourth comparison also yields 0, which is a bit surprising, since the strings are composed of the same letters, with one string all uppercase and the other all lowercase. The reason for this result is that MySQL's `strcmp()` function is case-insensitive, which is something to remember when using the function. The other four comparisons yield either -1 or 1 depending on whether the first string comes before or after the second string in sort order. For example, `strcmp('abcd','xyz')` yields -1, since the string 'abcd' comes before the string 'xyz'.

Along with the `strcmp()` function, MySQL also allows you to use the `like` and `regexp` operators to compare strings in the `select` clause. Such comparisons will yield 1 (for `true`) or 0 (for `false`). Therefore, these operators allow you to build expressions that return a number, much like the functions described in this section. Here's an example using `like`:

```
mysql> SELECT name, name LIKE '%y' ends_in_y
-> FROM category;
+-----+-----+
| name      | ends_in_y |
+-----+-----+
| Action     |      0 |
| Animation  |      0 |
```

```

| Children      |      0 |
| Classics     |      0 |
| Comedy        |      1 |
| Documentary  |      1 |
| Drama         |      0 |
| Family        |      1 |
| Foreign       |      0 |
| Games         |      0 |
| Horror        |      0 |
| Music         |      0 |
| New           |      0 |
| Sci-Fi        |      0 |
| Sports         |      0 |
| Travel         |      0 |
+-----+-----+
16 rows in set (0.00 sec)

```

This example retrieves all the category names, along with an expression that returns 1 if the name ends in “y” or 0 otherwise. If you want to perform more complex pattern matches, you can use the `regexp` operator, as demonstrated by the following:

```

mysql> SELECT name, name REGEXP 'y$' ends_in_y
    -> FROM category;
+-----+-----+
| name      | ends_in_y |
+-----+-----+
| Action    |      0 |
| Animation |      0 |
| Children  |      0 |
| Classics  |      0 |
| Comedy    |      1 |
| Documentary |      1 |
| Drama    |      0 |
| Family   |      1 |
| Foreign  |      0 |
| Games    |      0 |
| Horror   |      0 |
| Music    |      0 |
| New      |      0 |
| Sci-Fi   |      0 |
| Sports   |      0 |
| Travel   |      0 |
+-----+-----+
16 rows in set (0.00 sec)

```

The second column of this query returns 1 if the value stored in the `name` column matches the given regular expression.



Microsoft SQL Server and Oracle Database users can achieve similar results by building case expressions, which I describe in detail in [Chapter 11](#).

String functions that return strings

In some cases, you will need to modify existing strings, either by extracting part of the string or by adding additional text to the string. Every database server includes multiple functions to help with these tasks. Before I begin, I once again reset the data in the `string_tbl` table:

```
mysql> DELETE FROM string_tbl;
Query OK, 5 rows affected (0.00 sec)

mysql> INSERT INTO string_tbl (text_fld)
-> VALUES ('This string was 29 characters');
Query OK, 1 row affected (0.01 sec)
```

Earlier in the chapter, I demonstrated the use of the `concat()` function to help build words that include accented characters. The `concat()` function is useful in many other situations, including when you need to append additional characters to a stored string. For instance, the following example modifies the string stored in the `text_fld` column by tacking an additional phrase on the end:

```
mysql> UPDATE string_tbl
-> SET text_fld = CONCAT(text_fld, ', but now it is longer');
Query OK, 1 row affected (0.03 sec)
Rows matched: 1  Changed: 1  Warnings: 0
```

The contents of the `text_fld` column are now as follows:

```
mysql> SELECT text_fld
-> FROM string_tbl;
+-----+
| text_fld           |
+-----+
| This string was 29 characters, but now it is longer |
+-----+
1 row in set (0.00 sec)
```

Thus, like all functions that return a string, you can use `concat()` to replace the data stored in a character column.

Another common use for the `concat()` function is to build a string from individual pieces of data. For example, the following query generates a narrative string for each customer:

```
mysql> SELECT concat(first_name, ' ', last_name,
->   ' has been a customer since ', date(create_date)) cust_narrative
-> FROM customer;
```

```

+-----+
| cust_narrative |
+-----+
| MARY SMITH has been a customer since 2006-02-14 |
| PATRICIA JOHNSON has been a customer since 2006-02-14 |
| LINDA WILLIAMS has been a customer since 2006-02-14 |
| BARBARA JONES has been a customer since 2006-02-14 |
| ELIZABETH BROWN has been a customer since 2006-02-14 |
| JENNIFER DAVIS has been a customer since 2006-02-14 |
| MARIA MILLER has been a customer since 2006-02-14 |
| SUSAN WILSON has been a customer since 2006-02-14 |
| MARGARET MOORE has been a customer since 2006-02-14 |
| DOROTHY TAYLOR has been a customer since 2006-02-14 |
|
| ...
| RENE MCALISTER has been a customer since 2006-02-14 |
| EDUARDO HIATT has been a customer since 2006-02-14 |
| TERRENCE GUNDERSON has been a customer since 2006-02-14 |
| ENRIQUE FORSYTHE has been a customer since 2006-02-14 |
| FREDDIE DUGGAN has been a customer since 2006-02-14 |
| WADE DELVALLE has been a customer since 2006-02-14 |
| AUSTIN CINTRON has been a customer since 2006-02-14 |
+
599 rows in set (0.00 sec)

```

The `concat()` function can handle any expression that returns a string and will even convert numbers and dates to string format, as evidenced by the date column (`create_date`) used as an argument. Although Oracle Database includes the `concat()` function, it will accept only two string arguments, so the previous query will not work on Oracle. Instead, you would need to use the concatenation operator (`||`) rather than a function call, as in:

```

SELECT first_name || ' ' || last_name ||
    ' has been a customer since ' || date(create_date) cust_narrative
FROM customer;

```

SQL Server does not include a `concat()` function, so you would need to use the same approach as the previous query, except that you would use SQL Server's concatenation operator (+) instead of `||`.

While `concat()` is useful for adding characters to the beginning or end of a string, you may also have a need to add or replace characters in the *middle* of a string. All three database servers provide functions for this purpose, but all of them are different, so I demonstrate the MySQL function and then show the functions from the other two servers.

MySQL includes the `insert()` function, which takes four arguments: the original string, the position at which to start, the number of characters to replace, and the replacement string. Depending on the value of the third argument, the function may be used to either insert or replace characters in a string. With a value of 0 for the third

argument, the replacement string is inserted, and any trailing characters are pushed to the right, as in:

```
mysql> SELECT INSERT('goodbye world', 9, 0, 'cruel ') string;
+-----+
| string      |
+-----+
| goodbye cruel world |
+-----+
1 row in set (0.00 sec)
```

In this example, all characters starting from position 9 are pushed to the right, and the string 'cruel' is inserted. If the third argument is greater than zero, then that number of characters is replaced with the replacement string, as in:

```
mysql> SELECT INSERT('goodbye world', 1, 7, 'hello') string;
+-----+
| string      |
+-----+
| hello world |
+-----+
1 row in set (0.00 sec)
```

For this example, the first seven characters are replaced with the string 'hello'. Oracle Database does not provide a single function with the flexibility of MySQL's `insert()` function, but Oracle does provide the `replace()` function, which is useful for replacing one substring with another. Here's the previous example reworked to use `replace()`:

```
SELECT REPLACE('goodbye world', 'goodbye', 'hello')
FROM dual;
```

All instances of the string 'goodbye' will be replaced with the string 'hello', resulting in the string 'hello world'. The `replace()` function will replace *every* instance of the search string with the replacement string, so you need to be careful that you don't end up with more replacements than you anticipated.

SQL Server also includes a `replace()` function with the same functionality as Oracle's, but SQL Server also includes a function called `stuff()` with similar functionality to MySQL's `insert()` function. Here's an example:

```
SELECT STUFF('hello world', 1, 5, 'goodbye cruel')
```

When executed, five characters are removed starting at position 1, and then the string 'goodbye cruel' is inserted at the starting position, resulting in the string 'goodbye cruel world'.

Along with inserting characters into a string, you may have a need to *extract* a substring from a string. For this purpose, all three servers include the `substring()` function (although Oracle Database's version is called `substr()`), which extracts a

specified number of characters starting at a specified position. The following example extracts five characters from a string starting at the ninth position:

```
mysql> SELECT SUBSTRING('goodbye cruel world', 9, 5);
+-----+
| SUBSTRING('goodbye cruel world', 9, 5) |
+-----+
| cruel |
+-----+
1 row in set (0.00 sec)
```

Along with the functions demonstrated here, all three servers include many more built-in functions for manipulating string data. While many of them are designed for very specific purposes, such as generating the string equivalent of octal or hexadecimal numbers, there are many other general-purpose functions as well, such as functions that remove or add trailing spaces. For more information, consult your server's SQL reference guide, or a general-purpose SQL reference guide such as *SQL in a Nutshell* (O'Reilly).

Working with Numeric Data

Unlike string data (and temporal data, as you will see shortly), numeric data generation is quite straightforward. You can type a number, retrieve it from another column, or generate it via a calculation. All the usual arithmetic operators (+, -, *, /) are available for performing calculations, and parentheses may be used to dictate precedence, as in:

```
mysql> SELECT (37 * 59) / (78 - (8 * 6));
+-----+
| (37 * 59) / (78 - (8 * 6)) |
+-----+
| 72.77 |
+-----+
1 row in set (0.00 sec)
```

As I mentioned in [Chapter 2](#), the main concern when storing numeric data is that numbers might be rounded if they are larger than the specified size for a numeric column. For example, the number 9.96 will be rounded to 10.0 if stored in a column defined as `float(3,1)`.

Performing Arithmetic Functions

Most of the built-in numeric functions are used for specific arithmetic purposes, such as determining the square root of a number. [Table 7-1](#) lists some of the common numeric functions that take a single numeric argument and return a number.

Table 7-1. Single-argument numeric functions

Function name	Description
<code>acos(x)</code>	Calculates the arc cosine of x
<code>asin(x)</code>	Calculates the arc sine of x
<code>atan(x)</code>	Calculates the arc tangent of x
<code>cos(x)</code>	Calculates the cosine of x
<code>cot(x)</code>	Calculates the cotangent of x
<code>exp(x)</code>	Calculates e^x
<code>ln(x)</code>	Calculates the natural log of x
<code>sin(x)</code>	Calculates the sine of x
<code>sqrt(x)</code>	Calculates the square root of x
<code>tan(x)</code>	Calculates the tangent of x

These functions perform very specific tasks, and I refrain from showing examples for these functions (if you don't recognize a function by name or description, then you probably don't need it). Other numeric functions used for calculations, however, are a bit more flexible and deserve some explanation.

For example, the `modulo` operator, which calculates the remainder when one number is divided into another number, is implemented in MySQL and Oracle Database via the `mod()` function. The following example calculates the remainder when 10 is divided by 4:

```
mysql> SELECT MOD(10,4);
+-----+
| MOD(10,4) |
+-----+
|      2     |
+-----+
1 row in set (0.02 sec)
```

While the `mod()` function is typically used with integer arguments, with MySQL you can also use real numbers, as in:

```
mysql> SELECT MOD(22.75, 5);
+-----+
| MOD(22.75, 5) |
+-----+
|      2.75     |
+-----+
1 row in set (0.02 sec)
```



SQL Server does not have a `mod()` function. Instead, the operator `%` is used for finding remainders. The expression `10 % 4` will therefore yield the value 2.

Another numeric function that takes two numeric arguments is the `pow()` function (or `power()` if you are using Oracle Database or SQL Server), which returns one number raised to the power of a second number, as in:

```
mysql> SELECT POW(2,8);
+-----+
| POW(2,8) |
+-----+
|      256 |
+-----+
1 row in set (0.03 sec)
```

Thus, `pow(2,8)` is the MySQL equivalent of specifying 2^8 . Since computer memory is allocated in chunks of 2^x bytes, the `pow()` function can be a handy way to determine the exact number of bytes in a certain amount of memory:

```
mysql> SELECT POW(2,10) kilobyte, POW(2,20) megabyte,
->     POW(2,30) gigabyte, POW(2,40) terabyte;
+-----+-----+-----+-----+
| kilobyte | megabyte | gigabyte | terabyte |
+-----+-----+-----+-----+
|      1024 |    1048576 | 1073741824 | 1099511627776 |
+-----+-----+-----+-----+
1 row in set (0.00 sec)
```

I don't know about you, but I find it easier to remember that a gigabyte is 2^{30} bytes than to remember the number 1,073,741,824.

Controlling Number Precision

When working with floating-point numbers, you may not always want to interact with or display a number with its full precision. For example, you may store monetary transaction data with a precision to six decimal places, but you might want to round to the nearest hundredth for display purposes. Four functions are useful when limiting the precision of floating-point numbers: `ceil()`, `floor()`, `round()`, and `truncate()`. All three servers include these functions, although Oracle Database includes `trunc()` instead of `truncate()`, and SQL Server includes `ceiling()` instead of `ceil()`.

The `ceil()` and `floor()` functions are used to round either up or down to the closest integer, as demonstrated by the following:

```
mysql> SELECT CEIL(72.445), FLOOR(72.445);
+-----+-----+
| CEIL(72.445) | FLOOR(72.445) |
+-----+-----+
|          73 |            72 |
+-----+-----+
1 row in set (0.06 sec)
```

Thus, any number between 72 and 73 will be evaluated as 73 by the `ceil()` function and 72 by the `floor()` function. Remember that `ceil()` will round up even if the decimal portion of a number is very small, and `floor()` will round down even if the decimal portion is quite significant, as in:

```
mysql> SELECT CEIL(72.000000001), FLOOR(72.999999999);
+-----+-----+
| CEIL(72.000000001) | FLOOR(72.999999999) |
+-----+-----+
|          73 |            72 |
+-----+-----+
1 row in set (0.00 sec)
```

If this is a bit too severe for your application, you can use the `round()` function to round up or down from the *midpoint* between two integers, as in:

```
mysql> SELECT ROUND(72.49999), ROUND(72.5), ROUND(72.50001);
+-----+-----+-----+
| ROUND(72.49999) | ROUND(72.5) | ROUND(72.50001) |
+-----+-----+-----+
|          72 |          73 |          73 |
+-----+-----+-----+
1 row in set (0.00 sec)
```

Using `round()`, any number whose decimal portion is halfway or more between two integers will be rounded up, whereas the number will be rounded down if the decimal portion is anything less than halfway between the two integers.

Most of the time, you will want to keep at least some part of the decimal portion of a number rather than rounding to the nearest integer; the `round()` function allows an optional second argument to specify how many digits to the right of the decimal place to round to. The next example shows how you can use the second argument to round the number 72.0909 to one, two, and three decimal places:

```
mysql> SELECT ROUND(72.0909, 1), ROUND(72.0909, 2), ROUND(72.0909, 3);
+-----+-----+-----+
| ROUND(72.0909, 1) | ROUND(72.0909, 2) | ROUND(72.0909, 3) |
+-----+-----+-----+
|          72.1 |          72.09 |        72.091 |
+-----+-----+-----+
1 row in set (0.00 sec)
```

Like the `round()` function, the `truncate()` function allows an optional second argument to specify the number of digits to the right of the decimal, but `truncate()` simply discards the unwanted digits without rounding. The next example shows how the number 72.0909 would be truncated to one, two, and three decimal places:

```
mysql> SELECT TRUNCATE(72.0909, 1), TRUNCATE(72.0909, 2),
->      TRUNCATE(72.0909, 3);
+-----+-----+-----+
| TRUNCATE(72.0909, 1) | TRUNCATE(72.0909, 2) | TRUNCATE(72.0909, 3) |
+-----+-----+-----+
|          72.0 |           72.09 |        72.090 |
+-----+-----+-----+
1 row in set (0.00 sec)
```



SQL Server does not include a `truncate()` function. Instead, the `round()` function allows for an optional third argument that, if present and nonzero, calls for the number to be truncated rather than rounded.

Both `truncate()` and `round()` also allow a *negative* value for the second argument, meaning that numbers to the *left* of the decimal place are truncated or rounded. This might seem like a strange thing to do at first, but there are valid applications. For example, you might sell a product that can be purchased only in units of 10. If a customer were to order 17 units, you could choose from one of the following methods to modify the customer's order quantity:

```
mysql> SELECT ROUND(17, -1), TRUNCATE(17, -1);
+-----+-----+
| ROUND(17, -1) | TRUNCATE(17, -1) |
+-----+-----+
|          20 |           10 |
+-----+-----+
1 row in set (0.00 sec)
```

If the product in question is thumbtacks, then it might not make much difference to your bottom line whether you sold the customer 10 or 20 thumbtacks when only 17 were requested; if you are selling Rolex watches, however, your business may fare better by rounding.

Handling Signed Data

If you are working with numeric columns that allow negative values (in [Chapter 2](#), I showed how a numeric column may be labeled *unsigned*, meaning that only positive numbers are allowed), several numeric functions might be of use. Let's say, for example, that you are asked to generate a report showing the current status of a set of bank accounts using the following data from the `account` table:

```
+-----+-----+-----+
| account_id | acct_type     | balance |
+-----+-----+-----+
|      123 | MONEY MARKET |   785.22 |
|      456 | SAVINGS       |    0.00 |
|      789 | CHECKING     | -324.22 |
+-----+-----+-----+
```

The following query returns three columns useful for generating the report:

```
mysql> SELECT account_id, SIGN(balance), ABS(balance)
-> FROM account;
+-----+-----+-----+
| account_id | SIGN(balance) | ABS(balance) |
+-----+-----+-----+
|      123 |           1 |      785.22 |
|      456 |           0 |       0.00 |
|      789 |          -1 |      324.22 |
+-----+-----+-----+
3 rows in set (0.00 sec)
```

The second column uses the `sign()` function to return `-1` if the account balance is negative, `0` if the account balance is zero, and `1` if the account balance is positive. The third column returns the absolute value of the account balance via the `abs()` function.

Working with Temporal Data

Of the three types of data discussed in this chapter (character, numeric, and temporal), temporal data is the most involved when it comes to data generation and manipulation. Some of the complexity of temporal data is caused by the myriad ways in which a single date and time can be described. For example, the date on which I wrote this paragraph can be described in all the following ways:

- Wednesday, June 5, 2019
- 6/05/2019 2:14:56 P.M. EST
- 6/05/2019 19:14:56 GMT
- 1562019 (Julian format)
- Star date [-4] 97026.79 14:14:56 (*Star Trek* format)

While some of these differences are purely a matter of formatting, most of the complexity has to do with your frame of reference, which we explore in the next section.

Dealing with Time Zones

Because people around the world prefer that noon coincides roughly with the sun's peak at their location, there has never been a serious attempt to coerce everyone to

use a universal clock. Instead, the world has been sliced into 24 imaginary sections, called *time zones*; within a particular time zone, everyone agrees on the current time, whereas people in different time zones do not. While this seems simple enough, some geographic regions shift their time by one hour twice a year (implementing what is known as *daylight saving time*) and some do not, so the time difference between two points on Earth might be four hours for one-half of the year and five hours for the other half of the year. Even within a single time zone, different regions may or may not adhere to daylight saving time, causing different clocks in the same time zone to agree for one-half of the year but be one hour different for the rest of the year.

While the computer age has exacerbated the issue, people have been dealing with time zone differences since the early days of naval exploration. To ensure a common point of reference for timekeeping, fifteenth-century navigators set their clocks to the time of day in Greenwich, England. This became known as *Greenwich Mean Time*, or GMT. All other time zones can be described by the number of hours' difference from GMT; for example, the time zone for the Eastern United States, known as *Eastern Standard Time*, can be described as GMT -5:00, or five hours earlier than GMT.

Today, we use a variation of GMT called *Coordinated Universal Time*, or UTC, which is based on an atomic clock (or, to be more precise, the average time of 200 atomic clocks in 50 locations worldwide, which is referred to as *Universal Time*). Both SQL Server and MySQL provide functions that will return the current UTC timestamp (`getutcdate()` for SQL Server and `utc_timestamp()` for MySQL).

Most database servers default to the time zone setting of the server on which it resides and provide tools for modifying the time zone if needed. For example, a database used to store stock exchange transactions from around the world would generally be configured to use UTC time, whereas a database used to store transactions at a particular retail establishment might use the server's time zone.

MySQL keeps two different time zone settings: a global time zone and a session time zone, which may be different for each user logged in to a database. You can see both settings via the following query:

```
mysql> SELECT @@global.time_zone, @@session.time_zone;
+-----+-----+
| @@global.time_zone | @@session.time_zone |
+-----+-----+
| SYSTEM           | SYSTEM           |
+-----+-----+
1 row in set (0.00 sec)
```

A value of `system` tells you that the server is using the time zone setting from the server on which the database resides.

If you are sitting at a computer in Zurich, Switzerland, and you open a session across the network to a MySQL server situated in New York, you may want to change the time zone setting for your session, which you can do via the following command:

```
mysql> SET time_zone = 'Europe/Zurich';
Query OK, 0 rows affected (0.18 sec)
```

If you check the time zone settings again, you will see the following:

```
mysql> SELECT @@global.time_zone, @@session.time_zone;
+-----+-----+
| @@global.time_zone | @@session.time_zone |
+-----+-----+
| SYSTEM           | Europe/Zurich      |
+-----+-----+
1 row in set (0.00 sec)
```

All dates displayed in your session will now conform to Zurich time.



Oracle Database users can change the time zone setting for a session via the following command:

```
ALTER SESSION TIMEZONE = 'Europe/Zurich'
```

Generating Temporal Data

You can generate temporal data via any of the following means:

- Copying data from an existing `date`, `datetime`, or `time` column
- Executing a built-in function that returns a `date`, `datetime`, or `time`
- Building a string representation of the temporal data to be evaluated by the server

To use the last method, you will need to understand the various components used in formatting dates.

String representations of temporal data

Table 2-4 in Chapter 2 presented the more popular date components; to refresh your memory, Table 7-2 shows these same components.

Table 7-2. Date format components

Component	Definition	Range
YYYY	Year, including century	1000 to 9999
MM	Month	01 (January) to 12 (December)
DD	Day	01 to 31

Component	Definition	Range
HH	Hour	00 to 23
HHH	Hours (elapsed)	-838 to 838
MI	Minute	00 to 59
SS	Second	00 to 59

To build a string that the server can interpret as a `date`, `datetime`, or `time`, you need to put the various components together in the order shown in [Table 7-3](#).

Table 7-3. Required date components

Type	Default format
<code>date</code>	YYYY-MM-DD
<code>datetime</code>	YYYY-MM-DD HH:MI:SS
<code>timestamp</code>	YYYY-MM-DD HH:MI:SS
<code>time</code>	HHH:MI:SS

Thus, to populate a `datetime` column with 3:30 P.M. on September 17, 2019, you will need to build the following string:

```
'2019-09-17 15:30:00'
```

If the server is expecting a `datetime` value, such as when updating a `datetime` column or when calling a built-in function that takes a `datetime` argument, you can provide a properly formatted string with the required date components, and the server will do the conversion for you. For example, here's a statement used to modify the return date of a film rental:

```
UPDATE rental
SET return_date = '2019-09-17 15:30:00'
WHERE rental_id = 99999;
```

The server determines that the string provided in the `set` clause must be a `datetime` value, since the string is being used to populate a `datetime` column. Therefore, the server will attempt to convert the string for you by parsing the string into the six components (year, month, day, hour, minute, second) included in the default `date time` format.

String-to-date conversions

If the server is *not* expecting a `datetime` value or if you would like to represent the `datetime` using a nondefault format, you will need to tell the server to convert the string to a `datetime`. For example, here is a simple query that returns a `datetime` value using the `cast()` function:

```
mysql> SELECT CAST('2019-09-17 15:30:00' AS DATETIME);
+-----+
| CAST('2019-09-17 15:30:00' AS DATETIME) |
+-----+
| 2019-09-17 15:30:00 |
+-----+
1 row in set (0.00 sec)
```

We cover the `cast()` function at the end of this chapter. While this example demonstrates how to build `datetime` values, the same logic applies to the `date` and `time` types as well. The following query uses the `cast()` function to generate a `date` value and a `time` value:

```
mysql> SELECT CAST('2019-09-17' AS DATE) date_field,
->     CAST('108:17:57' AS TIME) time_field;
+-----+-----+
| date_field | time_field |
+-----+-----+
| 2019-09-17 | 108:17:57 |
+-----+-----+
1 row in set (0.00 sec)
```

You might, of course, explicitly convert your strings even when the server is expecting a `date`, `datetime`, or `time` value, rather than allowing the server to do an implicit conversion.

When strings are converted to temporal values—whether explicitly or implicitly—you must provide all the date components in the required order. While some servers are quite strict regarding the date format, the MySQL server is quite lenient about the separators used between the components. For example, MySQL will accept all of the following strings as valid representations of 3:30 P.M. on September 17, 2019:

```
'2019-09-17 15:30:00'
'2019/09/17 15:30:00'
'2019,09,17,15,30,00'
'20190917153000'
```

Although this gives you a bit more flexibility, you may find yourself trying to generate a temporal value *without* the default date components; the next section demonstrates a built-in function that is far more flexible than the `cast()` function.

Functions for generating dates

If you need to generate temporal data from a string and the string is not in the proper form to use the `cast()` function, you can use a built-in function that allows you to provide a format string along with the date string. MySQL includes the `str_to_date()` function for this purpose. Say, for example, that you pull the string '`September 17, 2019`' from a file and need to use it to update a `date` column. Since

the string is not in the required YYYY-MM-DD format, you can use `str_to_date()` instead of reformatting the string so that you can use the `cast()` function, as in:

```
UPDATE rental
SET return_date = STR_TO_DATE('September 17, 2019', '%M %d, %Y')
WHERE rental_id = 99999;
```

The second argument in the call to `str_to_date()` defines the format of the date string, with, in this case, a month name (%M), a numeric day (%d), and a four-digit numeric year (%Y). While there are more than 30 recognized format components, **Table 7-4** defines the dozen or so of the most commonly used components.

Table 7-4. Date format components

Format component	Description
%M	Month name (January to December)
%m	Month numeric (01 to 12)
%d	Day numeric (01 to 31)
%j	Day of year (001 to 366)
%W	Weekday name (Sunday to Saturday)
%Y	Year, four-digit numeric
%y	Year, two-digit numeric
%H	Hour (00 to 23)
%h	Hour (01 to 12)
%i	Minutes (00 to 59)
%s	Seconds (00 to 59)
%f	Microseconds (000000 to 999999)
%p	A.M. or P.M.

The `str_to_date()` function returns a `datetime`, `date`, or `time` value depending on the contents of the format string. For example, if the format string includes only %H, %i, and %s, then a `time` value will be returned.



Oracle Database users can use the `to_date()` function in the same manner as MySQL's `str_to_date()` function. SQL Server includes a `convert()` function that is not quite as flexible as MySQL and Oracle Database; rather than supplying a custom format string, your date string must conform to one of 21 predefined formats.

If you are trying to generate the *current* date/time, then you won't need to build a string, because the following built-in functions will access the system clock and return the current date and/or time as a string for you:

```
mysql> SELECT CURRENT_DATE(), CURRENT_TIME(), CURRENT_TIMESTAMP();
+-----+-----+-----+
| CURRENT_DATE() | CURRENT_TIME() | CURRENT_TIMESTAMP() |
+-----+-----+-----+
| 2019-06-05     | 16:54:36      | 2019-06-05 16:54:36 |
+-----+-----+-----+
1 row in set (0.12 sec)
```

The values returned by these functions are in the default format for the temporal type being returned. Oracle Database will include `current_date()` and `current_timestamp()` but not `current_time()`, and Microsoft SQL Server includes only the `current_timestamp()` function.

Manipulating Temporal Data

This section explores the built-in functions that take date arguments and return dates, strings, or numbers.

Temporal functions that return dates

Many of the built-in temporal functions take one date as an argument and return another date. MySQL's `date_add()` function, for example, allows you to add any kind of interval (e.g., days, months, years) to a specified date to generate another date. Here's an example that demonstrates how to add five days to the current date:

```
mysql> SELECT DATE_ADD(CURRENT_DATE(), INTERVAL 5 DAY);
+-----+
| DATE_ADD(CURRENT_DATE(), INTERVAL 5 DAY) |
+-----+
| 2019-06-10                                |
+-----+
1 row in set (0.06 sec)
```

The second argument is composed of three elements: the `interval` keyword, the desired quantity, and the type of interval. [Table 7-5](#) shows some of the commonly used interval types.

Table 7-5. Common interval types

Interval name	Description
second	Number of seconds
minute	Number of minutes
hour	Number of hours
day	Number of days
month	Number of months
year	Number of years
minute_second	Number of minutes and seconds, separated by ":"

Interval name	Description
hour_second	Number of hours, minutes, and seconds, separated by ":"
year_month	Number of years and months, separated by "-"

While the first six types listed in [Table 7-5](#) are pretty straightforward, the last three types require a bit more explanation since they have multiple elements. For example, if you are told that a film was actually returned 3 hours, 27 minutes, and 11 seconds later than what was originally specified, you can fix it via the following:

```
UPDATE rental
SET return_date = DATE_ADD(return_date, INTERVAL '3:27:11' HOUR_SECOND)
WHERE rental_id = 99999;
```

In this example, the `date_add()` function takes the value in the `return_date` column and adds 3 hours, 27 minutes, and 11 seconds to it. Then it uses the value that results to modify the `return_date` column.

Or, if you work in HR and found out that employee ID 4789 claimed to be older than he actually is, you could add 9 years and 11 months to his birth date, as in:

```
UPDATE employee
SET birth_date = DATE_ADD(birth_date, INTERVAL '9-11' YEAR_MONTH)
WHERE emp_id = 4789;
```



SQL Server users can accomplish the previous example using the `dateadd()` function:

```
UPDATE employee
SET birth_date =
    DATEADD(MONTH, 119, birth_date)
WHERE emp_id = 4789
```

SQL Server doesn't have combined intervals (i.e., `year_month`), so I converted 9 years and 11 months to 119 months.

Oracle Database users can use the `add_months()` function for this example, as in:

```
UPDATE employee
SET birth_date = ADD_MONTHS(birth_date, 119)
WHERE emp_id = 4789;
```

There are some cases where you want to add an interval to a date, and you know where you want to arrive but not how many days it takes to get there. For example, let's say that a bank customer logs on to the online banking system and schedules a transfer for the end of the month. Rather than writing some code that figures out the current month and then looks up the number of days in that month, you can call the `last_day()` function, which does the work for you (both MySQL and Oracle Database include the `last_day()` function; SQL Server has no comparable function). If

the customer asks for the transfer on September 17, 2019, you could find the last day of September via the following:

```
mysql> SELECT LAST_DAY('2019-09-17');
+-----+
| LAST_DAY('2019-09-17') |
+-----+
| 2019-09-30           |
+-----+
1 row in set (0.10 sec)
```

Whether you provide a `date` or `datetime` value, the `last_day()` function always returns a `date`. Although this function may not seem like an enormous time-saver, the underlying logic can be tricky if you're trying to find the last day of February and need to figure out whether the current year is a leap year.

Temporal functions that return strings

Most of the temporal functions that return string values are used to extract a portion of a date or time. For example, MySQL includes the `dayname()` function to determine which day of the week a certain date falls on, as in:

```
mysql> SELECT DAYNAME('2019-09-18');
+-----+
| DAYNAME('2019-09-18') |
+-----+
| Wednesday            |
+-----+
1 row in set (0.00 sec)
```

Many such functions are included with MySQL for extracting information from date values, but I recommend that you use the `extract()` function instead, since it's easier to remember a few variations of one function than to remember a dozen different functions. Additionally, the `extract()` function is part of the SQL:2003 standard and has been implemented by Oracle Database as well as MySQL.

The `extract()` function uses the same interval types as the `date_add()` function (see [Table 7-5](#)) to define which element of the date interests you. For example, if you want to extract just the year portion of a `datetime` value, you can do the following:

```
mysql> SELECT EXTRACT(YEAR FROM '2019-09-18 22:19:05');
+-----+
| EXTRACT(YEAR FROM '2019-09-18 22:19:05') |
+-----+
| 2019 |
+-----+
1 row in set (0.00 sec)
```



SQL Server doesn't include an implementation of `extract()`, but it does include the `datepart()` function. Here's how you would extract the year from a `datetime` value using `datepart()`:

```
SELECT DATEPART(YEAR, GETDATE())
```

Temporal functions that return numbers

Earlier in this chapter, I showed you a function used to add a given interval to a date value, thus generating another date value. Another common activity when working with dates is to take *two* date values and determine the number of intervals (days, weeks, years) *between* the two dates. For this purpose, MySQL includes the function `datediff()`, which returns the number of full days between two dates. For example, if I want to know the number of days that my kids will be out of school this summer, I can do the following:

```
mysql> SELECT DATEDIFF('2019-09-03', '2019-06-21');
+-----+
| DATEDIFF('2019-09-03', '2019-06-21') |
+-----+
|                               74 |
+-----+
1 row in set (0.00 sec)
```

Thus, I will have to endure 74 days of poison ivy, mosquito bites, and scraped knees before the kids are safely back at school. The `datediff()` function ignores the time of day in its arguments. Even if I include a time of day, setting it to one second until midnight for the first date and to one second after midnight for the second date, those times will have no effect on the calculation:

```
mysql> SELECT DATEDIFF('2019-09-03 23:59:59', '2019-06-21 00:00:01');
+-----+
| DATEDIFF('2019-09-03 23:59:59', '2019-06-21 00:00:01') |
+-----+
|                               74 |
+-----+
1 row in set (0.00 sec)
```

If I switch the arguments and have the earlier date first, `datediff()` will return a negative number, as in:

```
mysql> SELECT DATEDIFF('2019-06-21', '2019-09-03');
+-----+
| DATEDIFF('2019-06-21', '2019-09-03') |
+-----+
|                               -74 |
+-----+
1 row in set (0.00 sec)
```



SQL Server also includes the `datediff()` function, but it is more flexible than the MySQL implementation in that you can specify the interval type (i.e., year, month, day, hour) instead of counting only the number of days between two dates. Here's how SQL Server would accomplish the previous example:

```
SELECT DATEDIFF(DAY, '2019-06-21', '2019-09-03')
```

Oracle Database allows you to determine the number of days between two dates simply by subtracting one date from another.

Conversion Functions

Earlier in this chapter, I showed you how to use the `cast()` function to convert a string to a `datetime` value. While every database server includes a number of proprietary functions used to convert data from one type to another, I recommend using the `cast()` function, which is included in the SQL:2003 standard and has been implemented by MySQL, Oracle Database, and Microsoft SQL Server.

To use `cast()`, you provide a value or expression, the `as` keyword, and the type to which you want the value converted. Here's an example that converts a string to an integer:

```
mysql> SELECT CAST('1456328' AS SIGNED INTEGER);
+-----+
| CAST('1456328' AS SIGNED INTEGER) |
+-----+
|                      1456328 |
+-----+
1 row in set (0.01 sec)
```

When converting a string to a number, the `cast()` function will attempt to convert the entire string from left to right; if any nonnumeric characters are found in the string, the conversion halts without an error. Consider the following example:

```
mysql> SELECT CAST('999ABC111' AS UNSIGNED INTEGER);
+-----+
| CAST('999ABC111' AS UNSIGNED INTEGER) |
+-----+
|                           999 |
+-----+
1 row in set, 1 warning (0.08 sec)

mysql> show warnings;
+-----+-----+-----+
| Level | Code | Message           |
+-----+-----+-----+
| Warning | 1292 | Truncated incorrect INTEGER value: '999ABC111' |
+-----+-----+-----+
1 row in set (0.07 sec)
```

In this case, the first three digits of the string are converted, whereas the rest of the string is discarded, resulting in a value of 999. The server did, however, issue a warning to let you know that not all the string was converted.

If you are converting a string to a date, `time`, or `datetime` value, then you will need to stick with the default formats for each type, since you can't provide the `cast()` function with a format string. If your date string is not in the default format (i.e., `YYYY-MM-DD HH:MI:SS` for `datetime` types), then you will need to resort to using another function, such as MySQL's `str_to_date()` function described earlier in the chapter.

Test Your Knowledge

These exercises are designed to test your understanding of some of the built-in functions shown in this chapter. See [Appendix B](#) for the answers.

Exercise 7-1

Write a query that returns the 17th through 25th characters of the string 'Please find the substring in this string'.

Exercise 7-2

Write a query that returns the absolute value and sign (-1, 0, or 1) of the number -25.76823. Also return the number rounded to the nearest hundredth.

Exercise 7-3

Write a query to return just the month portion of the current date.

Grouping and Aggregates

Data is generally stored at the lowest level of granularity needed by any of a database's users; if Chuck in accounting needs to look at individual customer transactions, then there needs to be a table in the database that stores individual transactions. That doesn't mean, however, that all users must deal with the data as it is stored in the database. The focus of this chapter is on how data can be grouped and aggregated to allow users to interact with it at some higher level of granularity than what is stored in the database.

Grouping Concepts

Sometimes you will want to find trends in your data that will require the database server to cook the data a bit before you can generate the results you are looking for. For example, let's say that you are in charge of sending coupons for free rentals to your best customers. You could issue a simple query to look at the raw data:

```
mysql> SELECT customer_id FROM rental;
+-----+
| customer_id |
+-----+
|      1 |
|      1 |
|      1 |
|      1 |
|      1 |
|      1 |
|      1 |
|
|      ...    |
|      599 |
|      599 |
|      599 |
|      599 |
```

```
|      599 |
|      599 |
+-----+
16044 rows in set (0.01 sec)
```

With 599 customers spanning more than 16,000 rental records, it isn't feasible to determine which customers have rented the most films by looking at the raw data. Instead, you can ask the database server to group the data for you by using the `group by` clause. Here's the same query but employing a `group by` clause to group the rental data by customer ID:

```
mysql> SELECT customer_id
-> FROM rental
-> GROUP BY customer_id;
+-----+
| customer_id |
+-----+
|      1 |
|      2 |
|      3 |
|      4 |
|      5 |
|      6 |
...
|      594 |
|      595 |
|      596 |
|      597 |
|      598 |
|      599 |
+-----+
599 rows in set (0.00 sec)
```

The result set contains one row for each distinct value in the `customer_id` column, resulting in 599 rows instead of the full 16,044 rows. The reason for the smaller result set is that some of the customers rented more than one film. To see how many films each customer rented, you can use an *aggregate function* in the `select` clause to count the number of rows in each group:

```
mysql> SELECT customer_id, count(*)
-> FROM rental
-> GROUP BY customer_id;
+-----+-----+
| customer_id | count(*) |
+-----+-----+
|      1 |      32 |
|      2 |      27 |
|      3 |      26 |
|      4 |      22 |
|      5 |      38 |
|      6 |      28 |
```

```

...
|      594 |      27 |
|      595 |      30 |
|      596 |      28 |
|      597 |      25 |
|      598 |      22 |
|      599 |      19 |
+-----+-----+
599 rows in set (0.01 sec)

```

The aggregate function `count()` counts the number of rows in each group, and the asterisk tells the server to count everything in the group. Using the combination of a `group by` clause and the `count()` aggregate function, you are able to generate exactly the data needed to answer the business question without having to look at the raw data.

Looking at the results, you can see that 32 films were rented by customer ID 1, and 25 films were rented by the customer ID 597. In order to determine which customers have rented the most films, simply add an `order by` clause:

```

mysql> SELECT customer_id, count(*)
    -> FROM rental
    -> GROUP BY customer_id
    -> ORDER BY 2 DESC;
+-----+-----+
| customer_id | count(*) |
+-----+-----+
|      148 |      46 |
|      526 |      45 |
|      236 |      42 |
|      144 |      42 |
|      75 |      41 |
...
|      248 |      15 |
|      110 |      14 |
|      281 |      14 |
|       61 |      14 |
|      318 |      12 |
+-----+-----+
599 rows in set (0.01 sec)

```

Now that the results are sorted, you can easily see that customer ID 148 has rented the most films (46), while customer ID 318 has rented the fewest films (12).

When grouping data, you may need to filter out undesired data from your result set based on groups of data rather than based on the raw data. Since the `group by` clause runs *after* the `where` clause has been evaluated, you cannot add filter conditions to your `where` clause for this purpose. For example, here's an attempt to filter out any customers who have rented fewer than 40 films:

```
mysql> SELECT customer_id, count(*)
-> FROM rental
-> WHERE count(*) >= 40
-> GROUP BY customer_id;
ERROR 1111 (HY000): Invalid use of group function
```

You cannot refer to the aggregate function `count(*)` in your `where` clause, because the groups have not yet been generated at the time the `where` clause is evaluated. Instead, you must put your group filter conditions in the `having` clause. Here's what the query would look like using `having`:

```
mysql> SELECT customer_id, count(*)
-> FROM rental
-> GROUP BY customer_id
-> HAVING count(*) >= 40;
+-----+-----+
| customer_id | count(*) |
+-----+-----+
|      75 |      41 |
|     144 |      42 |
|     148 |      46 |
|     197 |      40 |
|    236 |      42 |
|    469 |      40 |
|    526 |      45 |
+-----+-----+
7 rows in set (0.01 sec)
```

Because those groups containing fewer than 40 members have been filtered out via the `having` clause, the result set now contains only those customers who have rented 40 or more films.

Aggregate Functions

Aggregate functions perform a specific operation over all rows in a group. Although every database server has its own set of specialty aggregate functions, the common aggregate functions implemented by all major servers include:

`max()`
Returns the maximum value within a set

`min()`
Returns the minimum value within a set

`avg()`
Returns the average value across a set

`sum()`
Returns the sum of the values across a set

`count()`

Returns the number of values in a set

Here's a query that uses all of the common aggregate functions to analyze the data on film rental payments:

```
mysql> SELECT MAX(amount) max_amt,
->   MIN(amount) min_amt,
->   AVG(amount) avg_amt,
->   SUM(amount) tot_amt,
->   COUNT(*) num_payments
->   FROM payment;
+-----+-----+-----+-----+
| max_amt | min_amt | avg_amt | tot_amt | num_payments |
+-----+-----+-----+-----+
| 11.99 | 0.00 | 4.200667 | 67416.51 | 16049 |
+-----+-----+-----+-----+
1 row in set (0.09 sec)
```

The results from this query tell you that, across the 16,049 rows in the `payment` table, the maximum amount paid to rent a film was \$11.99, the minimum amount was \$0, the average payment was \$4.20, and the total of all rental payments was \$67,416.51. Hopefully, this gives you an appreciation for the role of these aggregate functions; the next subsections further clarify how you can utilize these functions.

Implicit Versus Explicit Groups

In the previous example, every value returned by the query is generated by an aggregate function. Since there is no `group by` clause, there is a single, *implicit* group (all rows in the `payment` table).

In most cases, however, you will want to retrieve additional columns along with columns generated by aggregate functions. What if, for example, you wanted to extend the previous query to execute the same five aggregate functions for *each* customer, instead of across all customers? For this query, you would want to retrieve the `customer_id` column along with the five aggregate functions, as in:

```
SELECT customer_id,
      MAX(amount) max_amt,
      MIN(amount) min_amt,
      AVG(amount) avg_amt,
      SUM(amount) tot_amt,
      COUNT(*) num_payments
   FROM payment;
```

However, if you try to execute the query, you will receive the following error:

```
ERROR 1140 (42000): In aggregated query without GROUP BY,
expression #1 of SELECT list contains nonaggregated column
```

While it may be obvious to you that you want the aggregate functions applied to each customer found in the `payment` table, this query fails because you have not *explicitly* specified how the data should be grouped. Therefore, you will need to add a `group by` clause to specify over which group of rows the aggregate functions should be applied:

```
mysql> SELECT customer_id,
->   MAX(amount) max_amt,
->   MIN(amount) min_amt,
->   AVG(amount) avg_amt,
->   SUM(amount) tot_amt,
->   COUNT(*) num_payments
-> FROM payment
-> GROUP BY customer_id;
```

customer_id	max_amt	min_amt	avg_amt	tot_amt	num_payments
1	9.99	0.99	3.708750	118.68	32
2	10.99	0.99	4.767778	128.73	27
3	10.99	0.99	5.220769	135.74	26
4	8.99	0.99	3.717273	81.78	22
5	9.99	0.99	3.805789	144.62	38
6	7.99	0.99	3.347143	93.72	28
...					
594	8.99	0.99	4.841852	130.73	27
595	10.99	0.99	3.923333	117.70	30
596	6.99	0.99	3.454286	96.72	28
597	8.99	0.99	3.990000	99.75	25
598	7.99	0.99	3.808182	83.78	22
599	9.99	0.99	4.411053	83.81	19

599 rows in set (0.04 sec)

With the inclusion of the `group by` clause, the server knows to group together rows having the same value in the `customer_id` column first and then to apply the five aggregate functions to each of the 599 groups.

Counting Distinct Values

When using the `count()` function to determine the number of members in each group, you have your choice of counting *all* members in the group or counting only the *distinct* values for a column across all members of the group.

For example, consider the following query, which uses the `count()` function with the `customer_id` column in two different ways:

```
mysql> SELECT COUNT(customer_id) num_rows,
->   COUNT(DISTINCT customer_id) num_customers
-> FROM payment;
```

num_rows	num_customers

```
|      16049 |          599 |
+-----+-----+
1 row in set (0.01 sec)
```

The first column in the query simply counts the number of rows in the `payment` table, whereas the second column examines the values in the `customer_id` column and counts only the number of unique values. By specifying `distinct`, therefore, the `count()` function examines the values of a column for each member of the group in order to find and remove duplicates, rather than simply counting the number of values in the group.

Using Expressions

Along with using columns as arguments to aggregate functions, you can use expressions as well. For example, you may want to find the maximum number of days between when a film was rented and subsequently returned. You can achieve this via the following query:

```
mysql> SELECT MAX(datediff(return_date,rental_date))
    -> FROM rental;
+-----+
| MAX(datediff(return_date,rental_date)) |
+-----+
|                               33 |
+-----+
1 row in set (0.01 sec)
```

The `datediff` function is used to compute the number of days between the return date and the rental date for every rental, and the `max` function returns the highest value, which in this case is 33 days.

While this example uses a fairly simple expression, expressions used as arguments to aggregate functions can be as complex as needed, as long as they return a number, string, or date. In [Chapter 11](#), I show you how you can use `case` expressions with aggregate functions to determine whether a particular row should or should not be included in an aggregation.

How Nulls Are Handled

When performing aggregations, or, indeed, any type of numeric calculation, you should always consider how `null` values might affect the outcome of your calculation. To illustrate, I will build a simple table to hold numeric data and populate it with the set {1, 3, 5}:

```
mysql> CREATE TABLE number_tbl
    -> (val SMALLINT);
Query OK, 0 rows affected (0.01 sec)
```

```
mysql> INSERT INTO number_tbl VALUES (1);
Query OK, 1 row affected (0.00 sec)

mysql> INSERT INTO number_tbl VALUES (3);
Query OK, 1 row affected (0.00 sec)

mysql> INSERT INTO number_tbl VALUES (5);
Query OK, 1 row affected (0.00 sec)
```

Consider the following query, which performs five aggregate functions on the set of numbers:

```
mysql> SELECT COUNT(*) num_rows,
->   COUNT(val) num_vals,
->   SUM(val) total,
->   MAX(val) max_val,
->   AVG(val) avg_val
-> FROM number_tbl;
+-----+-----+-----+-----+
| num_rows | num_vals | total | max_val | avg_val |
+-----+-----+-----+-----+
|      3 |       3 |     9 |       5 |   3.0000 |
+-----+-----+-----+-----+
1 row in set (0.08 sec)
```

The results are as you would expect: both `count(*)` and `count(val)` return the value 3, `sum(val)` returns the value 9, `max(val)` returns 5, and `avg(val)` returns 3. Next, I will add a `null` value to the `number_tbl` table and run the query again:

```
mysql> INSERT INTO number_tbl VALUES (NULL);
Query OK, 1 row affected (0.01 sec)

mysql> SELECT COUNT(*) num_rows,
->   COUNT(val) num_vals,
->   SUM(val) total,
->   MAX(val) max_val,
->   AVG(val) avg_val
-> FROM number_tbl;
+-----+-----+-----+-----+
| num_rows | num_vals | total | max_val | avg_val |
+-----+-----+-----+-----+
|      4 |       3 |     9 |       5 |   3.0000 |
+-----+-----+-----+-----+
1 row in set (0.00 sec)
```

Even with the addition of the `null` value to the table, the `sum()`, `max()`, and `avg()` functions all return the same values, indicating that they ignore any `null` values encountered. The `count(*)` function now returns the value 4, which is valid since the `number_tbl` table contains four rows, while the `count(val)` function still returns the value 3. The difference is that `count(*)` counts the number of rows, whereas

`count(val)` counts the number of *values* contained in the `val` column and ignores any `null` values encountered.

Generating Groups

People are rarely interested in looking at raw data; instead, people engaging in data analysis will want to manipulate the raw data to better suit their needs. Examples of common data manipulations include:

- Generating totals for a geographic region, such as total European sales
- Finding outliers, such as the top salesperson for 2020
- Determining frequencies, such as the number of films rented in each month

To answer these types of queries, you will need to ask the database server to group rows together by one or more columns or expressions. As you have seen already in several examples, the `group by` clause is the mechanism for grouping data within a query. In this section, you will see how to group data by one or more columns, how to group data using expressions, and how to generate rollups within groups.

Single-Column Grouping

Single-column groups are the simplest and most-often-used type of grouping. If you want to find the number of films associated with each actor, for example, you need only group on the `film_actor.actor_id` column, as in:

```
mysql> SELECT actor_id, count(*)
    -> FROM film_actor
    -> GROUP BY actor_id;
+-----+-----+
| actor_id | count(*) |
+-----+-----+
|      1 |      19 |
|      2 |      25 |
|      3 |      22 |
|      4 |      22 |
...
|   197 |      33 |
|   198 |      40 |
|   199 |      15 |
|   200 |      20 |
+-----+-----+
200 rows in set (0.11 sec)
```

This query generates 200 groups, one for each actor, and then sums the number of films for each member of the group.

Multicolumn Grouping

In some cases, you may want to generate groups that span *more* than one column. Expanding on the previous example, imagine that you want to find the total number of films for each film rating (G, PG, ...) for each actor. The following example shows how you can accomplish this:

```
mysql> SELECT fa.actor_id, f.rating, count(*)
-> FROM film_actor fa
-> INNER JOIN film f
-> ON fa.film_id = f.film_id
-> GROUP BY fa.actor_id, f.rating
-> ORDER BY 1,2;
+-----+-----+-----+
| actor_id | rating | count(*) |
+-----+-----+-----+
|      1 | G     |      4 |
|      1 | PG    |      6 |
|      1 | PG-13 |      1 |
|      1 | R     |      3 |
|      1 | NC-17 |      5 |
|      2 | G     |      7 |
|      2 | PG    |      6 |
|      2 | PG-13 |      2 |
|      2 | R     |      2 |
|      2 | NC-17 |      8 |
|
|      ... |        |        |
|      199 | G     |      3 |
|      199 | PG    |      4 |
|      199 | PG-13 |      4 |
|      199 | R     |      2 |
|      199 | NC-17 |      2 |
|      200 | G     |      5 |
|      200 | PG    |      3 |
|      200 | PG-13 |      2 |
|      200 | R     |      6 |
|      200 | NC-17 |      4 |
+-----+-----+-----+
996 rows in set (0.01 sec)
```

This version of the query generates 996 groups, one for each combination of actor and film rating found by joining the `film_actor` table with the `film` table. Along with adding the `rating` column to the `select` clause, I also added it to the `group by` clause, since `rating` is retrieved from a table and is not generated via an aggregate function such as `max` or `count`.

Grouping via Expressions

Along with using columns to group data, you can build groups based on the values generated by expressions. Consider the following query, which groups rentals by year:

```
mysql> SELECT extract(YEAR FROM rental_date) year,
->      COUNT(*) how_many
->   FROM rental
->  GROUP BY extract(YEAR FROM rental_date);
+-----+-----+
| year | how_many |
+-----+-----+
| 2005 |     15862 |
| 2006 |       182 |
+-----+-----+
2 rows in set (0.01 sec)
```

This query employs a fairly simple expression that uses the `extract()` function to return only the year portion of a date to group the rows in the `rental` table.

Generating Rollups

In “[Multicolumn Grouping](#)” on page 156, I showed an example that counts the number of films for each actor and film rating. Let’s say, however, that along with the total count for each actor/rating combination, you also want total counts for each distinct actor. You could run an additional query and merge the results, you could load the results of the query into a spreadsheet, or you could build a Python script, Java program, or some other mechanism to take that data and perform the additional calculations. Better yet, you could use the `WITH ROLLUP` option to have the database server do the work for you. Here’s the revised query using `WITH ROLLUP` in the `group by` clause:

```
mysql> SELECT fa.actor_id, f.rating, count(*)
->   FROM film_actor fa
->   INNER JOIN film f
->   ON fa.film_id = f.film_id
->  GROUP BY fa.actor_id, f.rating WITH ROLLUP
->  ORDER BY 1,2;
+-----+-----+-----+
| actor_id | rating | count(*) |
+-----+-----+-----+
|    NULL  |  NULL  |     5462 |
|      1   |  NULL  |       19 |
|      1   |    G    |        4 |
|      1   |    PG   |       6 |
|      1   |  PG-13 |       1 |
|      1   |      R  |       3 |
|      1   |  NC-17 |       5 |
|      2   |  NULL  |      25 |
|      2   |      G  |       7 |
```

```

|      2 | PG      |      6 |
|      2 | PG-13   |      2 |
|      2 | R       |      2 |
|      2 | NC-17   |      8 |
...
| 199 | NULL    |     15 |
| 199 | G       |      3 |
| 199 | PG      |      4 |
| 199 | PG-13   |      4 |
| 199 | R       |      2 |
| 199 | NC-17   |      2 |
| 200 | NULL    |     20 |
| 200 | G       |      5 |
| 200 | PG      |      3 |
| 200 | PG-13   |      2 |
| 200 | R       |      6 |
| 200 | NC-17   |      4 |
+-----+
1197 rows in set (0.07 sec)

```

There are now 201 additional rows in the result set, one for each of the 200 distinct actors and one for the grand total (all actors combined). For the 200 actor rollups, a `null` value is provided for the `rating` column, since the rollup is being performed across all ratings. Looking at the first line for `actor_id 200`, for example, you will see that a total of 20 films are associated with the actor; this equals the sum of the counts for each rating ($4 \text{ NC-17} + 6 \text{ R} + 2 \text{ PG-13} + 3 \text{ PG} + 5 \text{ G}$). For the grand total row in the first line of the output, a `null` value is provided for both the `actor_id` and `rating` columns; the total for the first line of output equals 5,462, which is equal to the number of rows in the `film_actor` table.



If you are using Oracle Database, you need to use a slightly different syntax to indicate that you want a rollup performed. The `group by` clause for the previous query would look as follows when using Oracle:

```
GROUP BY ROLLUP(fa.actor_id, f.rating)
```

The advantage of this syntax is that it allows you to perform rollups on a subset of the columns in the `group_by` clause. If you are grouping by columns `a`, `b`, and `c`, for example, you could indicate that the server should perform rollups on only columns `b` and `c` via the following:

```
GROUP BY a, ROLLUP(b, c)
```

If in addition to totals by actor you also want to calculate totals per rating, then you can use the `with cube` option, which will generate summary rows for *all* possible combinations of the grouping columns. Unfortunately, `with cube` is not available in version 8.0 of MySQL, but it is available with SQL Server and Oracle Database.

Group Filter Conditions

In [Chapter 4](#), I introduced you to various types of filter conditions and showed how you can use them in the `where` clause. When grouping data, you also can apply filter conditions to the data *after* the groups have been generated. The `having` clause is where you should place these types of filter conditions. Consider the following example:

```
mysql> SELECT fa.actor_id, f.rating, count(*)
-> FROM film_actor fa
-> INNER JOIN film f
-> ON fa.film_id = f.film_id
-> WHERE f.rating IN ('G','PG')
-> GROUP BY fa.actor_id, f.rating
-> HAVING count(*) > 9;
+-----+-----+-----+
| actor_id | rating | count(*) |
+-----+-----+-----+
|      137 | PG     |      10 |
|       37 | PG     |      12 |
|      180 | PG     |      12 |
|        7 | G      |      10 |
|      83 | G      |      14 |
|     129 | G      |      12 |
|     111 | PG     |      15 |
|      44 | PG     |      12 |
|      26 | PG     |      11 |
|      92 | PG     |      12 |
|       17 | G      |      12 |
|     158 | PG     |      10 |
|     147 | PG     |      10 |
|      14 | G      |      10 |
|     102 | PG     |      11 |
|     133 | PG     |      10 |
+-----+-----+-----+
16 rows in set (0.01 sec)
```

This query has two filter conditions: one in the `where` clause, which filters out any films rated something other than G or PG, and another in the `having` clause, which filters out any actors who appeared in less than 10 films. Thus, one of the filters acts on data *before* it is grouped, and the other filter acts on data *after* the groups have been created. If you mistakenly put both filters in the `where` clause, you will see the following error:

```
mysql> SELECT fa.actor_id, f.rating, count(*)
-> FROM film_actor fa
->   INNER JOIN film f
->     ON fa.film_id = f.film_id
-> WHERE f.rating IN ('G','PG')
->   AND count(*) > 9
-> GROUP BY fa.actor_id, f.rating;
ERROR 1111 (HY000): Invalid use of group function
```

This query fails because you cannot include an aggregate function in a query's `where` clause. This is because the filters in the `where` clause are evaluated *before* the grouping occurs, so the server can't yet perform any functions on groups.



When adding filters to a query that includes a `group by` clause, think carefully about whether the filter acts on raw data, in which case it belongs in the `where` clause, or on grouped data, in which case it belongs in the `having` clause.

Test Your Knowledge

Work through the following exercises to test your grasp of SQL's grouping and aggregating features. Check your work with the answers in [Appendix B](#).

Exercise 8-1

Construct a query that counts the number of rows in the `payment` table.

Exercise 8-2

Modify your query from Exercise 8-1 to count the number of payments made by each customer. Show the customer ID and the total amount paid for each customer.

Exercise 8-3

Modify your query from Exercise 8-2 to include only those customers who have made at least 40 payments.

CHAPTER 9

Subqueries

Subqueries are a powerful tool that you can use in all four SQL data statements. In this chapter, I'll show you how subqueries can be used to filter data, generate values, and construct temporary data sets. After a little experimentation, I think you'll agree that subqueries are one of the most powerful features of the SQL language.

What Is a Subquery?

A *subquery* is a query contained within another SQL statement (which I refer to as the *containing statement* for the rest of this discussion). A subquery is always enclosed within parentheses, and it is usually executed prior to the containing statement. Like any query, a subquery returns a result set that may consist of:

- A single row with a single column
- Multiple rows with a single column
- Multiple rows having multiple columns

The type of result set returned by the subquery determines how it may be used and which operators the containing statement may use to interact with the data the subquery returns. When the containing statement has finished executing, the data returned by any subqueries is discarded, making a subquery act like a temporary table with *statement scope* (meaning that the server frees up any memory allocated to the subquery results after the SQL statement has finished execution).

You already saw several examples of subqueries in earlier chapters, but here's a simple example to get started:

```
mysql> SELECT customer_id, first_name, last_name  
-> FROM customer  
-> WHERE customer_id = (SELECT MAX(customer_id) FROM customer);
```

```
+-----+-----+-----+
| customer_id | first_name | last_name |
+-----+-----+-----+
|      599 | AUSTIN     | CINTRON    |
+-----+-----+-----+
1 row in set (0.27 sec)
```

In this example, the subquery returns the maximum value found in the `customer_id` column in the `customer` table, and the containing statement then returns data about that customer. If you are ever confused about what a subquery is doing, you can run the subquery by itself (without the parentheses) to see what it returns. Here's the subquery from the previous example:

```
mysql> SELECT MAX(customer_id) FROM customer;
+-----+
| MAX(customer_id) |
+-----+
|      599 |
+-----+
1 row in set (0.00 sec)
```

The subquery returns a single row with a single column, which allows it to be used as one of the expressions in an equality condition (if the subquery returned two or more rows, it could be *compared* to something but could not be *equal* to anything, but more on this later). In this case, you can take the value the subquery returned and substitute it into the righthand expression of the filter condition in the containing query, as in the following:

```
mysql> SELECT customer_id, first_name, last_name
   -> FROM customer
   -> WHERE customer_id = 599;
+-----+-----+-----+
| customer_id | first_name | last_name |
+-----+-----+-----+
|      599 | AUSTIN     | CINTRON    |
+-----+-----+-----+
1 row in set (0.00 sec)
```

The subquery is useful in this case because it allows you to retrieve information about the customer with the highest ID in a single query, rather than retrieving the maximum `customer_id` using one query and then writing a second query to retrieve the desired data from the `customer` table. As you will see, subqueries are useful in many other situations as well and may become one of the most powerful tools in your SQL toolkit.

Subquery Types

Along with the differences noted previously regarding the type of result set returned by a subquery (single row/column, single row/multicolumn, or multiple columns), you can use another feature to differentiate subqueries; some subqueries are completely self-contained (called *noncorrelated subqueries*), while others reference columns from the containing statement (called *correlated subqueries*). The next several sections explore these two subquery types and show the different operators that you can employ to interact with them.

Noncorrelated Subqueries

The example from earlier in the chapter is a noncorrelated subquery; it may be executed alone and does not reference anything from the containing statement. Most subqueries that you encounter will be of this type unless you are writing update or delete statements, which frequently make use of correlated subqueries (more on this later). Along with being noncorrelated, the example from earlier in the chapter also returns a result set containing a single row and column. This type of subquery is known as a *scalar subquery* and can appear on either side of a condition using the usual operators (`=`, `<>`, `<`, `>`, `<=`, `>=`). The next example shows how you can use a scalar subquery in an inequality condition:

```
mysql> SELECT city_id, city
-> FROM city
-> WHERE country_id <>
-> (SELECT country_id FROM country WHERE country = 'India');
+-----+-----+
| city_id | city
+-----+-----+
|      1 | A Corua (La Corua)
|      2 | Abha
|      3 | Abu Dhabi
|      4 | Acua
|      5 | Adana
|      6 | Addis Abeba
...
|    595 | Zapopan
|    596 | Zaria
|    597 | Zeleznogorsk
|    598 | Zhezqazghan
|    599 | Zhoushan
|   600 | Ziguinchor
+-----+
540 rows in set (0.02 sec)
```

This query returns all cities that are not in India. The subquery, which is found on the last line of the statement, returns the country ID for India, and the containing query

returns all cities that do not have that country ID. While the subquery in this example is quite simple, subqueries may be as complex as you need them to be, and they may utilize any and all the available query clauses (`select`, `from`, `where`, `group by`, `having`, and `order by`).

If you use a subquery in an equality condition but the subquery returns more than one row, you will receive an error. For example, if you modify the previous query such that the subquery returns *all* countries *except for* India, you will receive the following error:

```
mysql> SELECT city_id, city
-> FROM city
-> WHERE country_id <>
-> (SELECT country_id FROM country WHERE country <> 'India');
ERROR 1242 (21000): Subquery returns more than 1 row
```

If you run the subquery by itself, you will see the following results:

```
mysql> SELECT country_id FROM country WHERE country <> 'India';
+-----+
| country_id |
+-----+
|      1 |
|      2 |
|      3 |
|      4 |
...
|    106 |
|    107 |
|    108 |
|    109 |
+-----+
108 rows in set (0.00 sec)
```

The containing query fails because an expression (`country_id`) cannot be equated to a set of expressions (`country_ids 1, 2, 3, ..., 109`). In other words, a single thing cannot be equated to a set of things. In the next section, you will see how to fix the problem by using a different operator.

Multiple-Row, Single-Column Subqueries

If your subquery returns more than one row, you will not be able to use it on one side of an equality condition, as the previous example demonstrated. However, there are four additional operators that you can use to build conditions with these types of subqueries.

The in and not in operators

While you can't *equate* a single value to a set of values, you can check to see whether a single value can be found *within* a set of values. The next example, while it doesn't use a subquery, demonstrates how to build a condition that uses the `in` operator to search for a value within a set of values:

```
mysql> SELECT country_id
-> FROM country
-> WHERE country IN ('Canada','Mexico');
+-----+
| country_id |
+-----+
|      20 |
|      60 |
+-----+
2 rows in set (0.00 sec)
```

The expression on the lefthand side of the condition is the `country` column, while the righthand side of the condition is a set of strings. The `in` operator checks to see whether either of the strings can be found in the `country` column; if so, the condition is met, and the row is added to the result set. You could achieve the same results using two equality conditions, as in:

```
mysql> SELECT country_id
-> FROM country
-> WHERE country = 'Canada' OR country = 'Mexico';
+-----+
| country_id |
+-----+
|      20 |
|      60 |
+-----+
2 rows in set (0.00 sec)
```

While this approach seems reasonable when the set contains only two expressions, it is easy to see why a single condition using the `in` operator would be preferable if the set contained dozens (or hundreds, thousands, etc.) of values.

Although you will occasionally create a set of strings, dates, or numbers to use on one side of a condition, you are more likely to generate the set using a subquery that returns one or more rows. The following query uses the `in` operator with a subquery on the righthand side of the filter condition to return all cities that are in Canada or Mexico:

```
mysql> SELECT city_id, city
-> FROM city
-> WHERE country_id IN
-> (SELECT country_id
->   FROM country
->   WHERE country IN ('Canada','Mexico'));
```

```

+-----+
| city_id | city
+-----+
| 179 | Gatineau
| 196 | Halifax
| 300 | Lethbridge
| 313 | London
| 383 | Oshawa
| 430 | Richmond Hill
| 565 | Vancouver
...
| 452 | San Juan Bautista Tuxtepec
| 541 | Torren
| 556 | Uruapan
| 563 | Valle de Santiago
| 595 | Zapopan
+-----+
37 rows in set (0.00 sec)

```

Along with seeing whether a value exists within a set of values, you can check the converse using the `not in` operator. Here's another version of the previous query using `not in` instead of `in`:

```

mysql> SELECT city_id, city
    -> FROM city
    -> WHERE country_id NOT IN
    -> (SELECT country_id
    ->     FROM country
    ->     WHERE country IN ('Canada','Mexico'));
+-----+
| city_id | city
+-----+
| 1 | A Corua (La Corua)
| 2 | Abha
| 3 | Abu Dhabi
| 5 | Adana
| 6 | Addis Abeba
...
| 596 | Zaria
| 597 | Zeleznogorsk
| 598 | Zhezqazghan
| 599 | Zhoushan
| 600 | Ziguinchor
+-----+
563 rows in set (0.00 sec)

```

This query finds all cities that are *not* in Canada or Mexico.

The all operator

While the `in` operator is used to see whether an expression can be found within a set of expressions, the `all` operator allows you to make comparisons between a single value and every value in a set. To build such a condition, you will need to use one of the comparison operators (`=`, `<>`, `<`, `>`, etc.) in conjunction with the `all` operator. For example, the next query finds all customers who have never gotten a free film rental:

```
mysql> SELECT first_name, last_name
-> FROM customer
-> WHERE customer_id <> ALL
-> (SELECT customer_id
-> FROM payment
-> WHERE amount = 0);
+-----+-----+
| first_name | last_name |
+-----+-----+
| MARY       | SMITH    |
| PATRICIA   | JOHNSON  |
| LINDA      | WILLIAMS |
| BARBARA    | JONES    |
...
| EDUARDO    | HIATT    |
| TERENCE    | GUNDERSON |
| ENRIQUE    | FORSYTHE |
| FREDDIE    | DUGGAN   |
| WADE       | DELVALLE |
| AUSTIN     | CINTRON  |
+-----+-----+
576 rows in set (0.01 sec)
```

The subquery returns the set of IDs for customers who have paid \$0 for a film rental, and the containing query returns the names of all customers whose ID is not in the set returned by the subquery. If this approach seems a bit clumsy to you, you are in good company; most people would prefer to phrase the query differently and avoid using the `all` operator. To illustrate, the previous query generates the same results as the next example, which uses the `not in` operator:

```
SELECT first_name, last_name
FROM customer
WHERE customer_id NOT IN
(SELECT customer_id
FROM payment
WHERE amount = 0)
```

It's a matter of preference, but I think that most people would find the version that uses `not in` to be easier to understand.



When using `not in` or `<> all` to compare a value to a set of values, you must be careful to ensure that the set of values does not contain a `null` value, because the server equates the value on the left-hand side of the expression to each member of the set, and any attempt to equate a value to `null` yields `unknown`. Thus, the following query returns an empty set:

```
mysql> SELECT first_name, last_name
-> FROM customer
-> WHERE customer_id NOT IN (122, 452, NULL);
Empty set (0.00 sec)
```

Here's another example using the `all` operator, but this time the subquery is in the `having` clause:

```
mysql> SELECT customer_id, count(*)
-> FROM rental
-> GROUP BY customer_id
-> HAVING count(*) > ALL
-> (SELECT count(*)
-> FROM rental r
-> INNER JOIN customer c
-> ON r.customer_id = c.customer_id
-> INNER JOIN address a
-> ON c.address_id = a.address_id
-> INNER JOIN city ct
-> ON a.city_id = ct.city_id
-> INNER JOIN country co
-> ON ct.country_id = co.country_id
-> WHERE co.country IN ('United States','Mexico','Canada')
-> GROUP BY r.customer_id
-> );
+-----+-----+
| customer_id | count(*) |
+-----+-----+
|          148 |        46 |
+-----+-----+
1 row in set (0.01 sec)
```

The subquery in this example returns the total number of film rentals for all customers in North America, and the containing query returns all customers whose total number of film rentals exceeds any of the North American customers.

The `any` operator

Like the `all` operator, the `any` operator allows a value to be compared to the members of a set of values; unlike `all`, however, a condition using the `any` operator evaluates to `true` as soon as a single comparison is favorable. This example will find all customers whose total film rental payments exceed the total payments for all customers in Bolivia, Paraguay, or Chile:

```

mysql> SELECT customer_id, sum(amount)
->   FROM payment
->   GROUP BY customer_id
->   HAVING sum(amount) > ANY
->     (SELECT sum(p.amount)
->       FROM payment p
->       INNER JOIN customer c
->         ON p.customer_id = c.customer_id
->       INNER JOIN address a
->         ON c.address_id = a.address_id
->       INNER JOIN city ct
->         ON a.city_id = ct.city_id
->       INNER JOIN country co
->         ON ct.country_id = co.country_id
->       WHERE co.country IN ('Bolivia','Paraguay','Chile')
->     GROUP BY co.country
->   );
+-----+-----+
| customer_id | sum(amount) |
+-----+-----+
|      137 |    194.61 |
|      144 |    195.58 |
|      148 |    216.54 |
|      178 |    194.61 |
|      459 |    186.62 |
|      526 |    221.55 |
+-----+-----+
6 rows in set (0.03 sec)

```

The subquery returns the total film rental fees for all customers in Bolivia, Paraguay, and Chile, and the containing query returns all customers who outspent at least one of these three countries (if you find yourself outspending an entire country, perhaps you need to cancel your Netflix subscription and book a trip to Bolivia, Paraguay, or Chile...).



Although most people prefer to use `in`, using `= any` is equivalent to using the `in` operator.

Multicolumn Subqueries

So far, the subquery examples in this chapter have returned a single column and one or more rows. In certain situations, however, you can use subqueries that return two or more columns. To show the utility of multicolumn subqueries, it might help to look first at an example that uses multiple, single-column subqueries:

```

mysql> SELECT fa.actor_id, fa.film_id
->   FROM film_actor fa

```

```

-> WHERE fa.actor_id IN
-> (SELECT actor_id FROM actor WHERE last_name = 'MONROE')
-> AND fa.film_id IN
-> (SELECT film_id FROM film WHERE rating = 'PG');

+-----+-----+
| actor_id | film_id |
+-----+-----+
|      120 |      63 |
|      120 |     144 |
|      120 |     414 |
|      120 |     590 |
|      120 |     715 |
|      120 |     894 |
|      178 |     164 |
|      178 |     194 |
|      178 |     273 |
|      178 |     311 |
|      178 |     983 |
+-----+-----+
11 rows in set (0.00 sec)

```

This query uses two subqueries to identify all actors with the last name Monroe and all films rated PG, and the containing query then uses this information to retrieve all cases where an actor named Monroe appeared in a PG film. However, you could merge the two single-column subqueries into one multicolumn subquery and compare the results to two columns in the `film_actor` table. To do so, your filter condition must name both columns from the `film_actor` table surrounded by parentheses and in the same order as returned by the subquery, as in:

```

mysql> SELECT actor_id, film_id
-> FROM film_actor
-> WHERE (actor_id, film_id) IN
-> (SELECT a.actor_id, f.film_id
->   FROM actor a
->       CROSS JOIN film f
->   WHERE a.last_name = 'MONROE'
->   AND f.rating = 'PG');

+-----+-----+
| actor_id | film_id |
+-----+-----+
|      120 |      63 |
|      120 |     144 |
|      120 |     414 |
|      120 |     590 |
|      120 |     715 |
|      120 |     894 |
|      178 |     164 |
|      178 |     194 |
|      178 |     273 |
|      178 |     311 |
|      178 |     983 |
+-----+-----+

```

```
+-----+-----+
11 rows in set (0.00 sec)
```

This version of the query performs the same function as the previous example, but with a single subquery that returns two columns instead of two subqueries that each return a single column. The subquery in this version uses a type of join called a *cross join*, which will be explored in the next chapter. The basic idea is to return all combinations of actors named Monroe (2) and all films rated PG (194) for a total of 388 rows, 11 of which can be found in the `film_actor` table.

Correlated Subqueries

All of the subqueries shown thus far have been independent of their containing statements, meaning that you can execute them by themselves and inspect the results. A *correlated subquery*, on the other hand, is *dependent* on its containing statement from which it references one or more columns. Unlike a noncorrelated subquery, a correlated subquery is not executed once prior to execution of the containing statement; instead, the correlated subquery is executed once for each candidate row (rows that might be included in the final results). For example, the following query uses a correlated subquery to count the number of film rentals for each customer, and the containing query then retrieves those customers who have rented exactly 20 films:

```
mysql> SELECT c.first_name, c.last_name
   -> FROM customer c
   -> WHERE 20 =
   -> (SELECT count(*) FROM rental r
   -> WHERE r.customer_id = c.customer_id);
+-----+-----+
| first_name | last_name  |
+-----+-----+
| LAUREN     | HUDSON    |
| JEANETTE   | GREENE    |
| TARA        | RYAN      |
| WILMA      | RICHARDS  |
| JO          | FOWLER    |
| KAY         | CALDWELL  |
| DANIEL     | CABRAL    |
| ANTHONY    | SCHWAB   |
| TERRY       | GRISSOM   |
| LUIS        | YANEZ    |
| HERBERT    | KRUGER    |
| OSCAR       | AQUINO   |
| RAUL        | FORTIER  |
| NELSON     | CHRISTENSON |
| ALFREDO    | MCADAMS  |
+-----+-----+
15 rows in set (0.01 sec)
```

The reference to `c.customer_id` at the very end of the subquery is what makes the subquery correlated; the containing query must supply values for `c.customer_id` for the subquery to execute. In this case, the containing query retrieves all 599 rows from the `customer` table and executes the subquery once for each customer, passing in the appropriate customer ID for each execution. If the subquery returns the value 20, then the filter condition is met, and the row is added to the result set.



One word of caution: since the correlated subquery will be executed once for each row of the containing query, the use of correlated subqueries can cause performance issues if the containing query returns a large number of rows.

Along with equality conditions, you can use correlated subqueries in other types of conditions, such as the range condition illustrated here:

```
mysql> SELECT c.first_name, c.last_name
-> FROM customer c
-> WHERE
-> (SELECT sum(p.amount) FROM payment p
-> WHERE p.customer_id = c.customer_id)
-> BETWEEN 180 AND 240;
+-----+-----+
| first_name | last_name |
+-----+-----+
| RHONDA    | KENNEDY   |
| CLARA     | SHAW       |
| ELEANOR   | HUNT       |
| MARION    | SNYDER    |
| TOMMY     | COLLAZO   |
| KARL      | SEAL       |
+-----+-----+
6 rows in set (0.03 sec)
```

This variation on the previous query finds all customers whose total payments for all film rentals are between \$180 and \$240. Once again, the correlated subquery is executed 599 times (once for each customer row), and each execution of the subquery returns the total account balance for the given customer.



Another subtle difference in the previous query is that the subquery is on the lefthand side of the condition, which may look a bit odd but is perfectly valid.